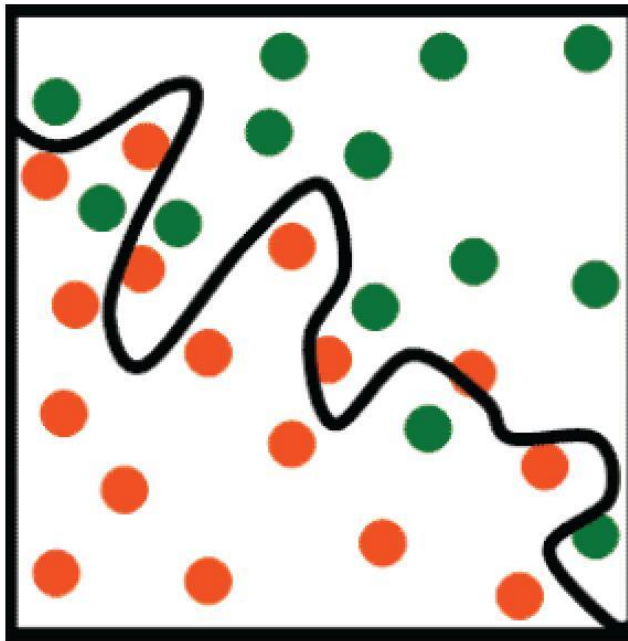


# Supervised Learning (II)

---



# This Lecture: Supervised Learning Methods

---

- I. Linear Model (detail)
- II. Decision Tree (detail)
- III. Neural Network (brief)
- IV. k-Nearest Neighbours (brief)
- V. Support Vector Machine (brief)

# I. Linear Model (线性模型)

**I.1 Univariate Linear Regression (ULR) 单元/一元线性回归**

I.2 Multivariate Linear Regression (MLR) 多元线性回归

I.3 Multivariate Linear Classification (MLC) 多元线性分类

# Example: House Price

---

- **[Question]** How to predict the price of a new house?
- **[Answer]** Fit a straight line using the training data  
⇒ **linear regression**.
- One variable (house size)  
⇒ **univariate**.

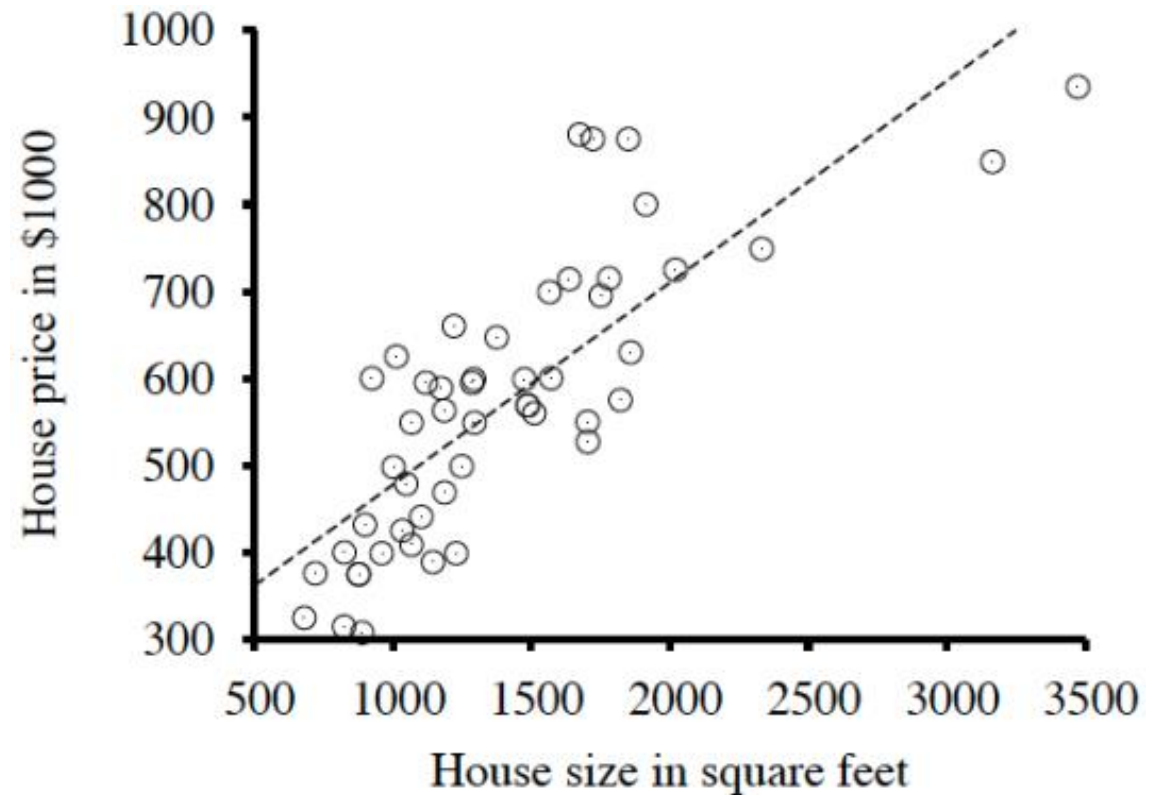


Image source: Figure 18.13.a of the AI book by S. Russell & P. Novig.

# Model Formulation

---

- **Linear model:**  $h_{\mathbf{w}}(x) = w_0 + w_1x$ ,
  - $\mathbf{w} = [w_0, w_1]^T \in \mathbb{R}^{2 \times 1}$  : model parameters,
  - $x \in \mathbb{R}^1$ : input feature (特征), e.g. house size.
- **Training data:**  $\mathcal{D} = \{(x^{(n)}, y^{(n)}) \mid \mathbf{y}^{(n)} \in \mathbb{R}^1\}_{n=1}^N$ .
- **Aim:** find the optimal  $\mathbf{w}$  fitting the observations in  $\mathcal{D}$ .
- **Optimization:** minimize empirical square loss regarding  $\mathbf{w}$  as

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - (w_1 x^{(n)} + w_0)]^2.$$

# Model Formulation

---

- **Aim:** find the optimal  $\mathbf{w}$  fitting the observations in  $\mathcal{D}$ .
- **Optimization:** minimize empirical square loss regarding  $\mathbf{w}$  as


$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - (w_1 x^{(n)} + w_0)]^2.$$

- **Remarks:**
  - Square loss -> Euclidean distance (欧式距离) -> least square method (最小二乘法)
  - Finding optimal  $\mathbf{w}$ : parameter estimation (参数估计)

# Model Parameter Space

---

- Plot 3D graph for  $\mathcal{L}(w_0, w_1) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - (w_1 x^{(n)} + w_0)]^2$



**Convex!**  
**Global optima assured!**

*Image source: Figure 18.13.b of the AI book by S. Russell & P. Novig.*

# 1. Closed-form Solution (闭式解)

---

- **Optimization:**  $\mathcal{L}(w_0, w_1) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - (w_1 x^{(n)} + w_0)]^2$ .
- **First-order equations:**

$$\frac{\partial}{\partial w_0} \mathcal{L}(w_0, w_1) = 0, \quad \frac{\partial}{\partial w_1} \mathcal{L}(w_0, w_1) = 0.$$

- **Solution:**

$$w_1 = \frac{N(\sum_n x^{(n)} y^{(n)}) - (\sum_n x^{(n)})(\sum_n y^{(n)})}{N(\sum_n (x^{(n)})^2) - (\sum_n x^{(n)})^2}$$

$$w_0 = \frac{\sum_n y^{(n)} - w_1 \sum_n x^{(n)}}{N}.$$



## 2. Iterative Solution

---

- Sometimes there is no closed-form solution.

- **Gradient Descent (GD, 梯度下降法):**

1.  $\mathbf{w} \leftarrow$  any point in the parameter space

2. **LOOP** until convergence **DO**

3. **FOR** each  $w_i$  in  $\mathbf{w}$  **DO**

4.  $w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} \mathcal{L}(\mathbf{w}),$

- $\alpha$ : learning rate, positive.






Image source: [https://commons.wikimedia.org/wiki/File:Gradient\\_descent.svg](https://commons.wikimedia.org/wiki/File:Gradient_descent.svg)

## 2. Iterative Solution: Advanced

---

- **Batch GD**: update  $\mathbf{w}$  once with all training samples.
  - Guarantee global optimum but slow.
- **Stochastic GD**: update  $\mathbf{w}$   $N$  times with one training data for one update.
  - Fast but do not guarantee global optimum with a fixed  $\alpha$ .
  - Online/offline settings
- **Mini-batch SGD**: update  $\mathbf{w}$  several times with a subset of  $\mathcal{D}$  for one update.



*Zigzag problem of SGD. Image source: Figure 4.10 of "Fundamentals of Deep Learning" by Nikhil Buduma.*

# I. Linear Model (线性模型)

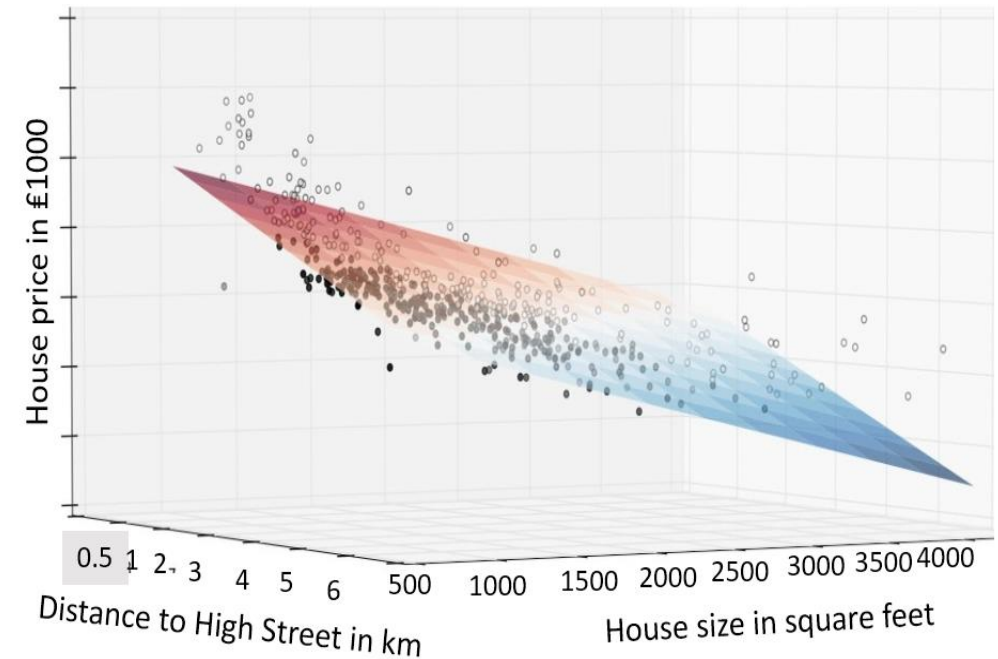
I.1 Univariate Linear Regression (ULR) 单元/一元线性回归

**I.2 Multivariate Linear Regression (MLR) 多元线性回归**

I.3 Multivariate Linear Classification (MLC) 多元线性分类

# Example: House Price Revisit

- One more feature: distance to High Street.
- **[Question]** How to predict the price of a new house with the two features?
- **[Answer]** Fit a plane using  $\mathcal{D}$   
⇒ linear regression.
- Multiple variables (distance + size)  
⇒ multivariate.



# Model Formulation

---

- **Linear model:**  $h_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1x + \cdots + w_mx_m = \mathbf{w}^T \mathbf{x} \in \mathbb{R}^1$ ,
  - $\mathbf{w} = [w_0, w_1, \cdots, w_m]^T \in \mathbb{R}^{(m+1)}$ : model parameters.
  - $\mathbf{x} = [1, x_1, \cdots, x_m]^T \in \mathbb{R}^{(m+1)}$ : input features; e.g. size, distance to High Street, etc.
- **Training data:**  $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)}) \mid y^{(n)} \in \mathbb{R}^1\}_{n=1}^N$ .
- **Aim:** find the optimal  $\mathbf{w}$  fitting the observations in  $\mathcal{D}$ .
- **Optimization:** minimize empirical loss regarding  $\mathbf{w}$  as

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - \mathbf{w}^T \mathbf{x}^{(n)}]^2.$$

# 1. Closed-form Solution

---

- **Optimization:**  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - \mathbf{w}^T \mathbf{x}^{(n)}]^2.$
- **First-order equations:**  $\frac{\partial}{\partial w_i} \mathcal{L}(\mathbf{w}) = 0, \forall i = 0, 1, \dots, m$
- **Solution:** solve the system of linear equations to have

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y},$$

- $\mathbf{X} = [x'^{(1)}, \dots, x'^{(N)}]^T \in \mathbb{R}^{N \times (m+1)},$
- $\mathbf{y} = [y^{(1)}, \dots, y^{(N)}]^T \in \mathbb{R}^{N \times 1}.$

## 2. Iterative Solution

---

- **Optimization:**  $\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2} \sum_{n=1}^N [y^{(n)} - \mathbf{w}^T \mathbf{x}^{(n)}]^2$ .
- **Gradient descent (GD):**

$$w_i \leftarrow w_i + \alpha \sum_n \left( y_i^{(n)} - \mathbf{w}^T \mathbf{x}^{(n)} \right) \cdot x_i^{(n)}$$

- $\alpha$ : learning rate, positive.

# Overfitting for MLR

---

- MLR in a high-dimensional space may encounter **overfitting**.
- MLR: common to use **regularization**.
- ULR does not have this problem – only 1 feature.



# Regularized Objective for MLR

---

- Regularized Objective:

$$\min_{\mathbf{w}} \mathcal{L}_{tr}(\mathbf{w}) + \lambda \cdot \Omega(\mathbf{w}).$$

- $\mathcal{L}_{tr}(\mathbf{w})$  : training loss; measures how well the model fits the training data.
  - Square loss:  $l(y^{(n)}, \widehat{y}^{(n)}) = (y^{(n)} - \widehat{y}^{(n)})^2 = (y^{(n)} - \mathbf{w}^T \mathbf{x}^{(n)})^2$ .
  - Logistic loss:  $l(y^{(n)}, \widehat{y}^{(n)}) = y^{(n)} \ln(1 + e^{-\widehat{y}^{(n)}}) + (1 - y^{(n)}) \ln(1 + e^{\widehat{y}^{(n)}})$
- $\lambda$ : trade-off & manually tuning parameter.


# Regularization

---

- $\Omega(\mathbf{w})$ : regularization; how complex the model is?
- $\Omega(\mathbf{w}) \triangleq \ell_p(\mathbf{w}) = \sum_i |w_i|^p$ , in particular:
  - $\ell_0$  regularization:  $p = 0$ , penalize #(non-zero parameters);
  - $\ell_1$  regularization:  $p = 1$ , penalize the sum of the absolute parameters;
  - $\ell_2$  regularization:  $p = 2$ , penalize the sum of square parameters.
- **[Question]** Which  $\ell_p$  (p范数) should we use?
- **[Answer]** Depend on the specific problem.

# Illustration: $\ell_1$ vs $\ell_2$

- Let  $w = [w_1, w_2]^T$ , we have:
  - $\ell_1 = |w_1| + |w_2|$ ,
  - $\ell_2 = w_1^2 + w_2^2$ .
- Plot contours for  $\ell_1 = \ell_2 = c$ .
- Goodness of  $\ell_1$ : **sparse** model.



*Image source: Figure 18.14 of the AI book by S. Russell & P. Novig.*

# Exercise: Closed-form Solution of MLR

---

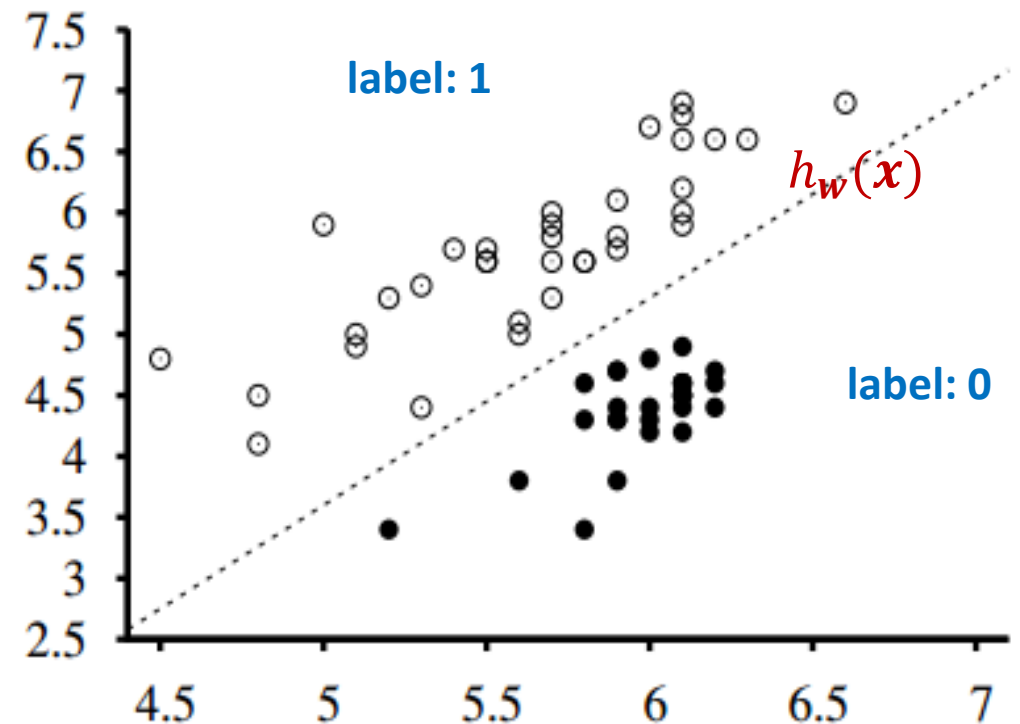
- **[Question]** Derive the closed-form solution of MLR?
- **[Hint]** At **variable-level**:
  - 1) Compute and write the  $i^{th}$  equation  $\frac{\partial}{\partial w_i} \mathcal{L}(\mathbf{w}) = 0$ .
  - 2) Re-write the  $i^{th}$  equation into 'vector-vector' form.
  - 3) Align the  $m$  equations about  $x_i^{(n)}$  into a matrix-vector form. Note to check the match of dimensionality.
  - 4) Obtain the representation of  $\mathbf{w}$ .
- **[Hint]** The solution is much easier to characterize in **matrix notation**.

# I. Linear Model (线性模型)

1. Univariate Linear Regression (ULR) 单元/一元线性回归
2. Multivariate Linear Regression (MLR) 多元线性回归
3. **Multivariate Linear Classification (MLC) 多元线性分类**

# Example: Earthquakes or Explosions

- [Question] How to distinguish earthquakes ( $\circ$ ) from explosions ( $\bullet$ ) using two features?
- [Answer] Learn a linear decision boundary that can separate the two-class points.
- [Denote] the classifier as  $h_w(x)$ .
  - Classifier: linear decision boundary.



Linear separable data points. Image source: Figure 18.15.a of the AI book by S. Russell & P. Novig.

# Classification Problem Formulation

---

- **Training data**:  $\mathcal{D} = \{(x^{(n)}, y^{(n)}) \mid y^{(n)} \in \{0, 1\}\}_{n=1}^N$ .
- **Aim**: find the optimal  $\mathbf{w}$  fitting the observations in  $\mathcal{D}$ .
- **Optimization**: minimize square-error regarding  $\mathbf{w}$  as

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N [y^{(n)} - h_{\mathbf{w}}(\mathbf{x}^{(n)})]^2.$$

- **Remaining Question**: How to formulate  $h_{\mathbf{w}}(\mathbf{x})$ ?

# Problem of MLR Model in Classification

---

- **MLR model:**  $h_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^T \mathbf{x} \in \mathbb{R}^1$ .
- **Problem:**  $h_{\mathbf{w}}(\mathbf{x}) \in \mathbb{R}^1$  cannot constrain 0/1 output.

=> Hard-threshold Linear Classifier (帶硬閾值的线性分类器)



# Hard-threshold Classifier


- Hard-threshold function:

$$\sigma(z) = \mathbb{1}_{z \geq 0} = \begin{cases} 1, & z \geq 0 \\ 0, & z < 0 \end{cases}.$$

- Hard-threshold classification model:

$$h_{\mathbf{w}}(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x}) \in \{0, 1\},$$

- $h_{\mathbf{w}}(\mathbf{x})$  has 0/1 output.



*Image source: Figure 18.17.a of the AI book by S. Russell & P. Novig.*

# Problem in Learning Method

---

- **Problem:** The below derivative does **NOT exist** (either 0 or undefined)

$$\frac{\partial}{\partial w_i} h_{\mathbf{w}}(\mathbf{x}) = \frac{\partial}{\partial z} \boxed{\sigma(z)} \frac{\partial (\mathbf{w}^T \mathbf{x})}{\partial w_i}.$$

- **Closed-form solution:** Cannot proceed.
- **Iterative solution:**  $w_i \leftarrow w_i - \alpha \frac{\partial}{\partial w_i} \mathcal{L}(\mathbf{w})$  needs the above derivative.

**Classification optimization:**

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N [y^{(n)} - h_{\mathbf{w}}(\mathbf{x}^{(n)})]^2.$$

# Proposed Learning Method

---

- **Perceptron learning rule:** For a training point  $(\mathbf{x}, y)$ , update as

$$w_i \leftarrow w_i + \alpha(y - h_{\mathbf{w}}(\mathbf{x})) \cdot x_i, \text{ for } i \in \{1, \dots, m\}.$$

- Identical to the MLR case in form.

- **Implementation:**

- $y \cdot h_{\mathbf{w}}(\mathbf{x}) = 1$ : keep  $w_i$  unchanged;
- $y = 1 \ \& \ h_{\mathbf{w}}(\mathbf{x}) = 0$ : increase  $w_i$  if  $x_i > 0$  and decrease  $w_i$  if  $x_i < 0$ ;
- $y = 0 \ \& \ h_{\mathbf{w}}(\mathbf{x}) = 1$ : decrease  $w_i$  if  $x_i > 0$  and increase  $w_i$  if  $x_i < 0$ .

# From Hard-threshold to Soft-threshold

**Hard threshold:**  $\sigma(z) = \mathbb{1}_{z \geq 0}$

- $\mathbb{1}_{z \geq 0}$ : indicator function.
- **Not differentiable.**




Image source: Figure 18.17.a of the AI book by S. Russell & P. Novig.

**Soft threshold function:**  $\sigma(z) = s(z) = \frac{1}{1 + e^{-z}}$

- $s(z)$ : sigmoid function.
- **Differentiable:**  $s'(z) = s(z)[1 - s(z)]$ .




Image source: Figure 18.17.b of the AI book by S. Russell & P. Novig.

# Logistic Regression (对数几率回归)

---

- **Logistic regression model:**  $h_{\mathbf{w}}(\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x})$ , where  $\sigma(z) = \frac{1}{1+e^{-z}}$ .
- **Closed-form solution:** Does not exist.
- **Iterative solution:**  $w_i \leftarrow w_i - \alpha \frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_i}$ 
  - $\frac{\partial \mathcal{L}(\mathbf{w})}{\partial w_i} = -\frac{1}{N} \sum_{n=1}^N [y^{(n)} - h_{\mathbf{w}}(\mathbf{x}^{(n)})] \cdot h_{\mathbf{w}}(\mathbf{x}^{(n)}) \cdot [1 - h_{\mathbf{w}}(\mathbf{x}^{(n)})] \cdot x_i^{(n)}$
  - $\alpha$ : learning rate, positive.

**Classification optimization:**

$$\min_{\mathbf{w}} \mathcal{L}(\mathbf{w}) = \frac{1}{2N} \sum_{n=1}^N [y^{(n)} - h_{\mathbf{w}}(\mathbf{x}^{(n)})]^2.$$

## II. Decision Tree (决策树)

1. **Tree Representation**
2. Decision Tree Construction with Heuristics
  - ❖ Information Gain: Good Feature Heuristics
  - ❖ Information Gain: Continuous Feature
  - ❖ Overall: Decision Tree Construction
3. Tree Overfitting
4. Decision Tree for Regression

# Example: Tree-shape Model

- **[Question]** How to judge an animal to be a mammal by two features?
- **[Answer]** Learn a **decision tree** that can separate the training samples.
- **Goodness:** Natural for humans, easy to interpret the results.

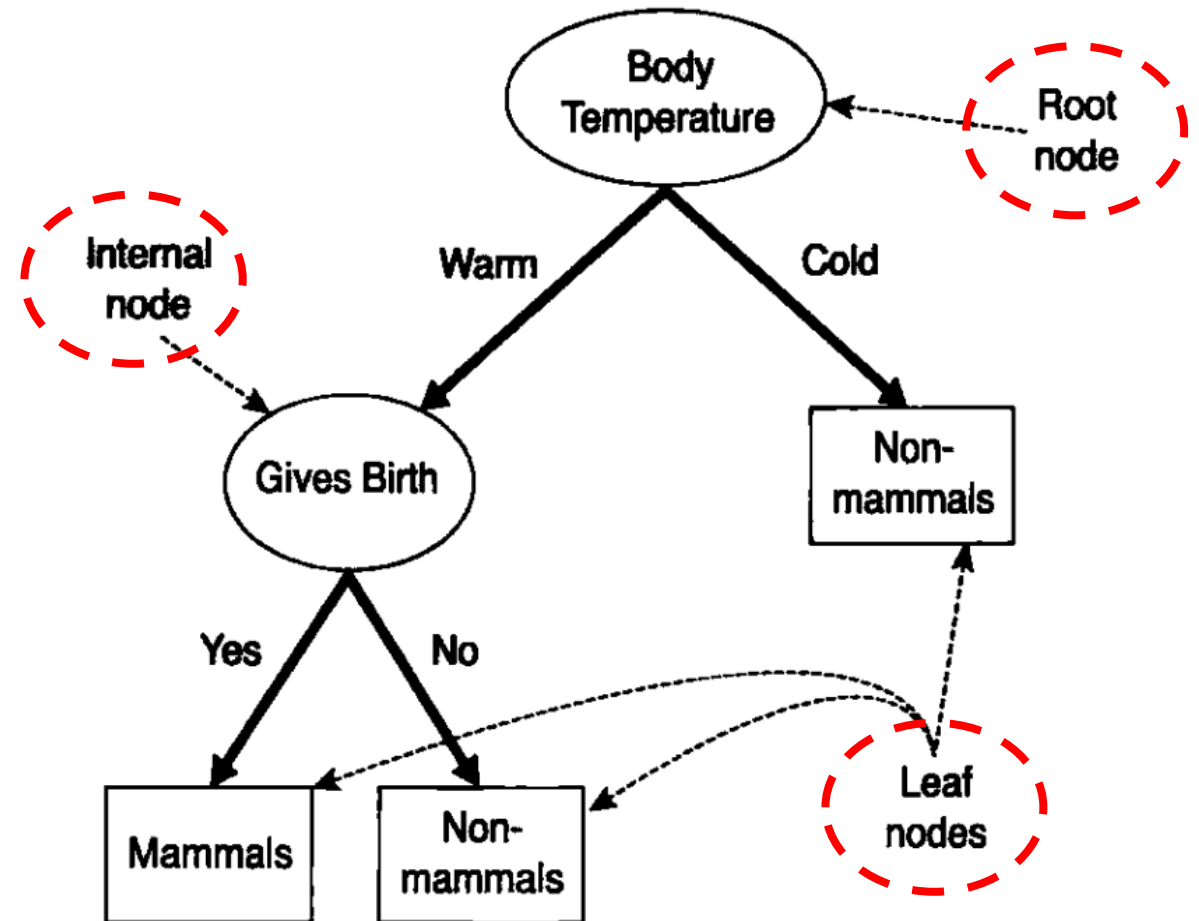


Image source: Figure 5.6 of "Introduction to Data Mining" by P. Tan, M. Steinbach and V. Kumar.

# Tree Representation

- **Tree model**: a function mapping feature vector  $x$  to a decision  $y$  via a **sequence of tests**.
  - Discrete  $y$ : a decision tree for classification.
  - Continuous  $y$ : a decision tree for regression.
- **Two types of nodes**:
  - Decision nodes: a test on some feature.
  - Leaf nodes: a decision of the tree model.

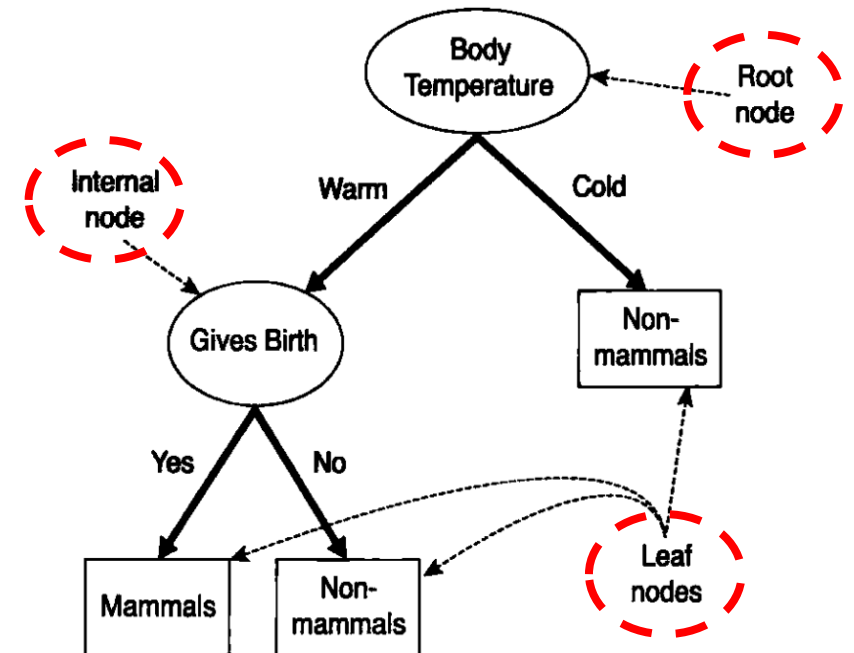


Image source: Figure 5.6 of "Introduction to Data Mining" by P. Tan, M. Steinbach and V. Kumar.



# [Example] Waiting at a Restaurant

---

- Prediction: Should we wait for a table?
- # Input features: 10
  - Alternate: is there an alternative restaurant nearby?
  - Bar: is there a comfortable bar area to wait in?
  - Fri/Sat: is today Friday or Saturday?
  - Hungry: are we hungry?
  - Patrons: # people in the restaurant (None, Some, Full)
  - Price: price range.
  - Raining: is it raining outside?
  - Reservation: have we made a reservation?
  - Type: kind of restaurant (French, Italian, Thai, Burger)
  - Wait-Estimate: estimated waiting time (0-10, 10-30, 30-60, >60)

# True Decision Tree $f$

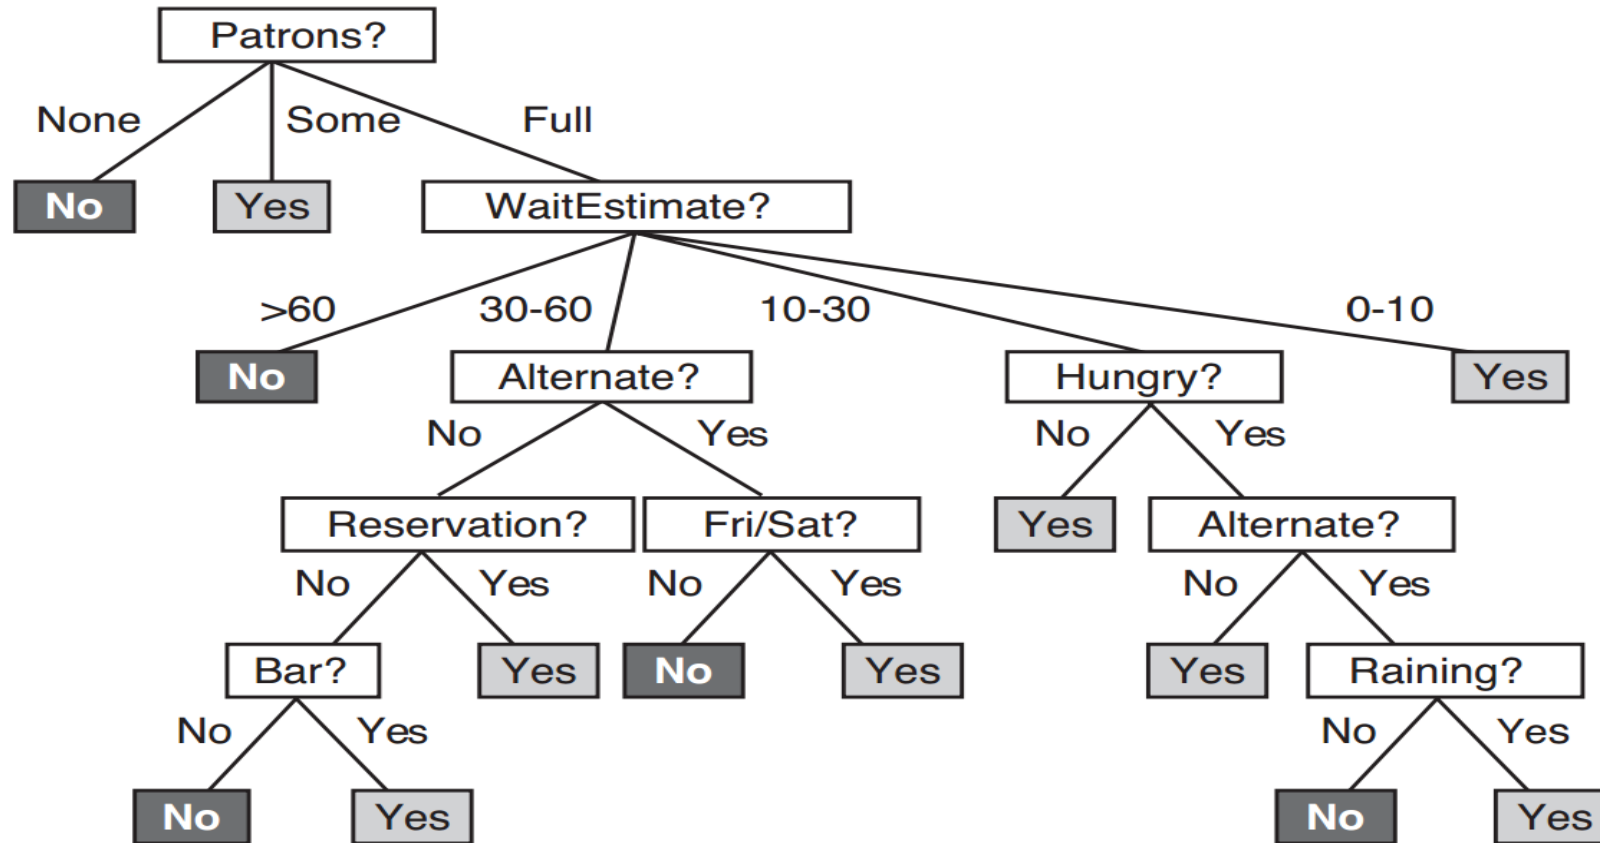


Fig. The decision tree  $f$  for deciding whether to wait for a table. Image source: Figure 18.4 of the AI book by S. Russell & P. Novig.

# 12 Training Examples (6+ 6-)

- Generate 12 training examples from the true decision tree.

Example	Input Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
<b>x<sub>1</sub></b>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>0–10</i>	<i>y<sub>1</sub> = Yes</i>
<b>x<sub>2</sub></b>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>30–60</i>	<i>y<sub>2</sub> = No</i>
<b>x<sub>3</sub></b>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Some</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	<i>y<sub>3</sub> = Yes</i>
<b>x<sub>4</sub></b>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Thai</i>	<i>10–30</i>	<i>y<sub>4</sub> = Yes</i>
<b>x<sub>5</sub></b>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>&gt;60</i>	<i>y<sub>5</sub> = No</i>
<b>x<sub>6</sub></b>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Italian</i>	<i>0–10</i>	<i>y<sub>6</sub> = Yes</i>
<b>x<sub>7</sub></b>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	<i>y<sub>7</sub> = No</i>
<b>x<sub>8</sub></b>	<i>No</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Thai</i>	<i>0–10</i>	<i>y<sub>8</sub> = Yes</i>
<b>x<sub>9</sub></b>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>&gt;60</i>	<i>y<sub>9</sub> = No</i>
<b>x<sub>10</sub></b>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>Italian</i>	<i>10–30</i>	<i>y<sub>10</sub> = No</i>
<b>x<sub>11</sub></b>	<i>No</i>	<i>No</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>0–10</i>	<i>y<sub>11</sub> = No</i>
<b>x<sub>12</sub></b>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>30–60</i>	<i>y<sub>12</sub> = Yes</i>

Image source: Figure 18.3 of the AI book by S. Russell & P. Novig.

# True Decision Tree $f$

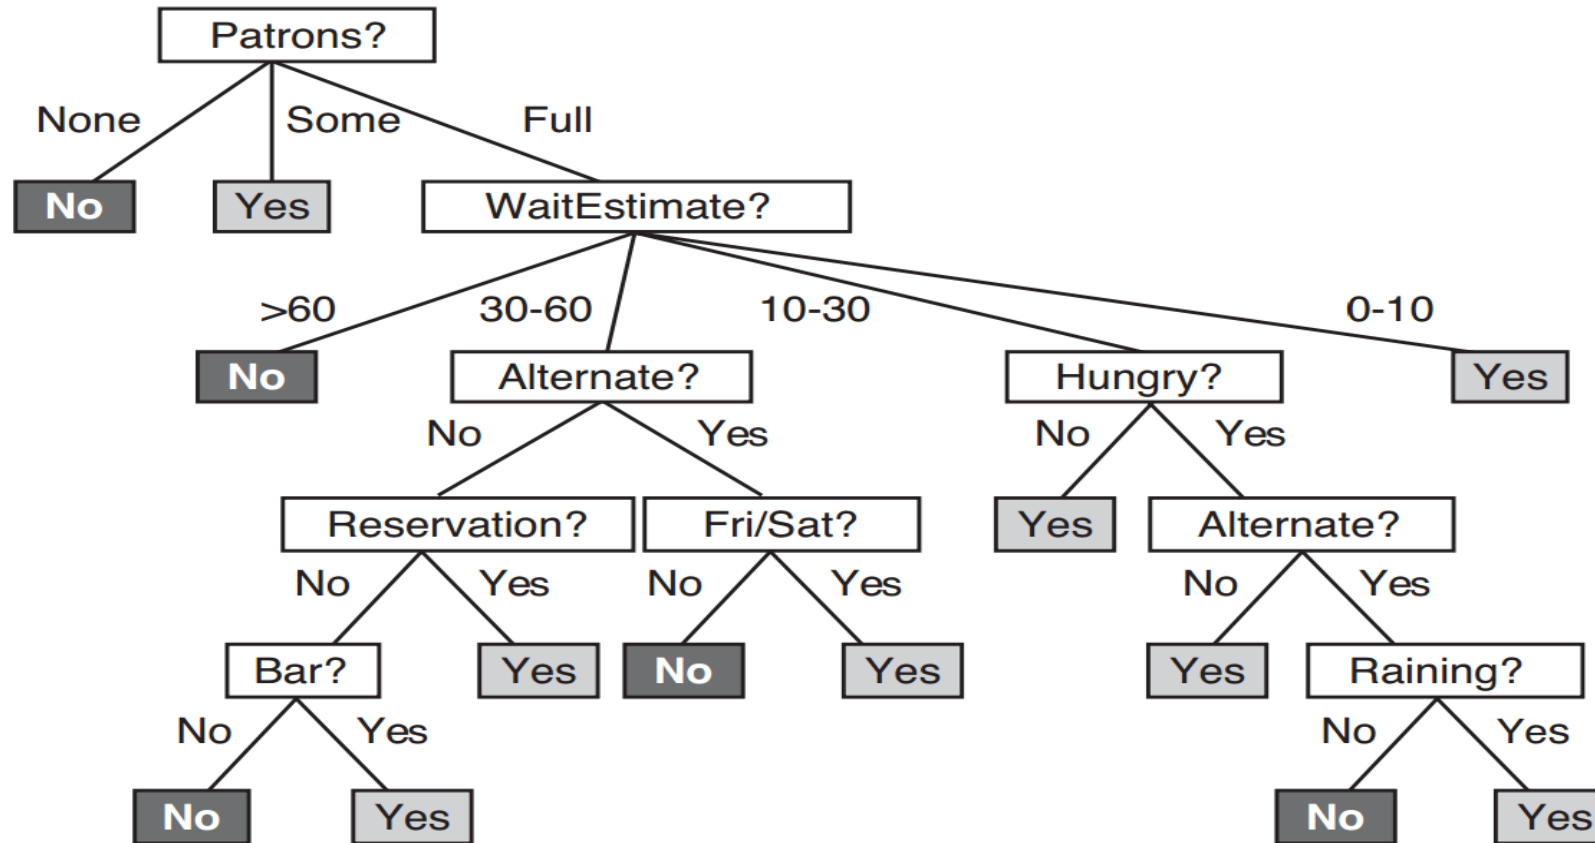


Fig. The decision tree  $f$  for deciding whether to wait for a table. Image source: Figure 18.4 of the AI book by S. Russell & P. Novig.

# Question

---

How to induce the decision tree from the training examples?

## II. Decision Tree (决策树)

1. Tree Representation
2. **Decision Tree Construction with Heuristics**
  - ❖ **Information Gain: Good Feature Heuristics**
  - ❖ **Information Gain: Continuous Feature**
  - ❖ **Overall: Decision Tree Construction**
3. Tree Overfitting
4. Decision Tree for Regression

# Learning Decision Tree is Hard

---

- **Resources:** The 12 training examples.
- **Aim:** Build the **smallest** tree that classifies the training data correctly.
  - Ockham's razor.
- **Challenge:** Finding the smallest tree is **NP-hard** [Hyafil & Rivest'76].

# Learning Decision Tree is Hard

---

- **[Question]** How many decision trees can be expressed (at least)?

Input space	Output
$\left\{ \begin{array}{l} 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0 \\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 1 \\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 0\ 1\ 0 \\ \dots \\ 0\ 0\ 0\ 0\ 1\ 0\ 0\ 0\ 1\ 0 \\ \dots \\ 1\ 1\ 1\ 1\ 1\ 1\ 1\ 1\ 1\ 1 \end{array} \right.$	$\left. \begin{array}{l} 0\ or\ 1 \\ 0\ or\ 1 \\ 0\ or\ 1 \\ \dots \\ 0\ or\ 1 \\ \dots \\ 0\ or\ 1 \end{array} \right\}$

- **[Answer]**  $2^{2^{10}}$ , super huge search space!
- Need talented **heuristics** to **guide the search** in such a huge space!



# Greedy Divide-and-conquer Strategy

---

- **Approach**: Greedy divide-and-conquer strategy - **heuristic search**.
  - (1) Start from empty tree.
  - (2) Decide the **best feature** based on **heuristics**.
  - (3) Divide the problem into smaller subproblems;
  - (4) Repeat (2)~(3) until stopping criteria.
- **Heuristics**: Pick the feature that maximizes **information gain** (信息增益).
  - The most informative feature.

# Good Feature: Type vs Patrons?

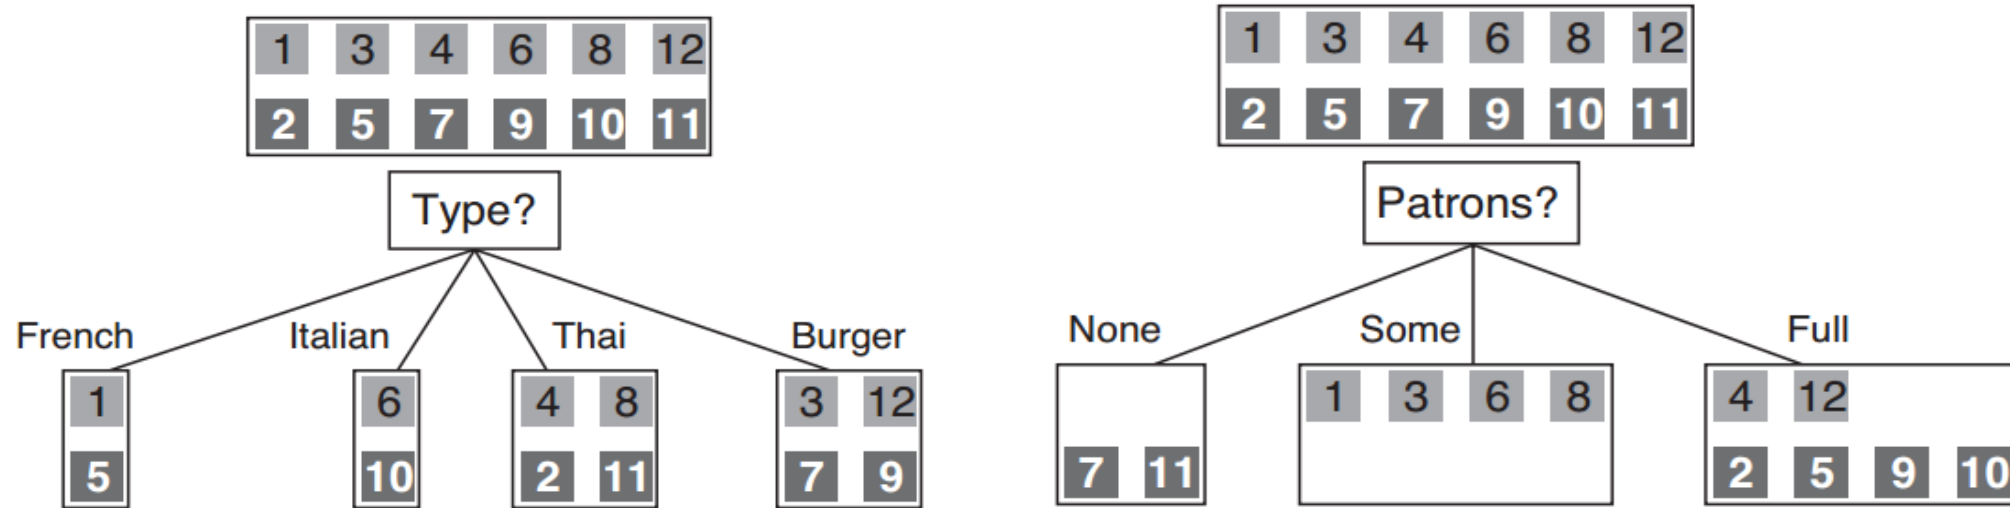


Image source: Figure 18.4 of the AI book by S. Russell & P. Novig.

- **[Question]** Split the tree based on 'type' or 'patrons'?
- **[Answer]** 'patrons'.
- **Reason:** Divides the 12 data into **more distinguishable** sub-sets.

# Heuristics for Good Feature

- **Intuition**: More certain about the classification after split regarding this feature.
  - Deterministic (all true or false): perfect
  - Uniform distribution: bad
  - What about in between?
- **[Question]** How to measure the **goodness** of a feature formally?
- **[Answer]** *Information gain*.

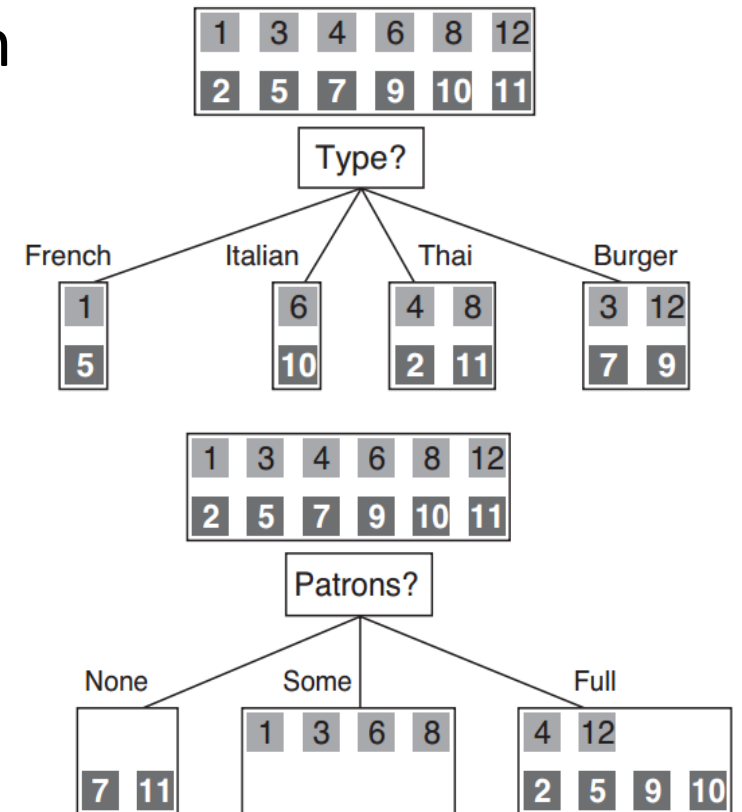


Image source: Figure 18.13 of the AI book by S. Russell & P. Novig.

# Preliminary: Entropy (熵)

- **Entropy**:  $\mathcal{H}(Y) \triangleq -\sum_k p(y_k) \log_2 p(y_k)$ .
- Larger entropy, more uncertainty.
  - High entropy:  $Y \sim$  uniform or flat distribution  $\rightarrow$  less predictable
  - Low entropy:  $Y \sim$  peak/valley distribution  $\rightarrow$  more predictable
- **Example 1**:  $\mathcal{H}(Y = \text{'label'}) = -(0.5 \log_2 0.5 + 0.5 \log_2 0.5) = 1$ .

Goal
<i>WillWait</i>
$y_1 = \text{Yes}$
$y_2 = \text{No}$
$y_3 = \text{Yes}$
$y_4 = \text{Yes}$
$y_5 = \text{No}$
$y_6 = \text{Yes}$
$y_7 = \text{No}$
$y_8 = \text{Yes}$
$y_9 = \text{No}$
$y_{10} = \text{No}$
$y_{11} = \text{No}$
$y_{12} = \text{Yes}$

Image source: Figure 18.3 of the AI book by S. Russell & P. Novig.

# Preliminary: Conditional Entropy

- Conditional entropy:

$$\mathcal{H}(Y|X) \triangleq \sum_j p(X = x_j) \cdot \mathcal{H}(Y|X = x_j).$$

- Example 2:  $Y \sim \text{label} \ \& \ X \sim \text{Type}$

- $p(X = \text{French}) = p(X = \text{Italian}) = \frac{2}{12};$
- $p(X = \text{Thai}) = p(X = \text{Burger}) = \frac{4}{12};$
- $\mathcal{H}(Y|X = \text{French or Italian}): \quad -\frac{1}{2} \log\left(\frac{1}{2}\right) - \frac{1}{2} \log\left(\frac{1}{2}\right) = -\log\left(\frac{1}{2}\right);$
- $\mathcal{H}(Y|X = \text{Thai or Burger}): \quad -\frac{2}{4} \log\left(\frac{2}{4}\right) - \frac{2}{4} \log\left(\frac{2}{4}\right) = -\log\left(\frac{1}{2}\right);$
- $\mathcal{H}(Y|X) = -\left[\frac{2}{12} \cdot \log\left(\frac{1}{2}\right) \cdot 2 + \frac{4}{12} \cdot \log\left(\frac{1}{2}\right) \cdot 2\right] = -\log\left(\frac{1}{2}\right) = 1.$

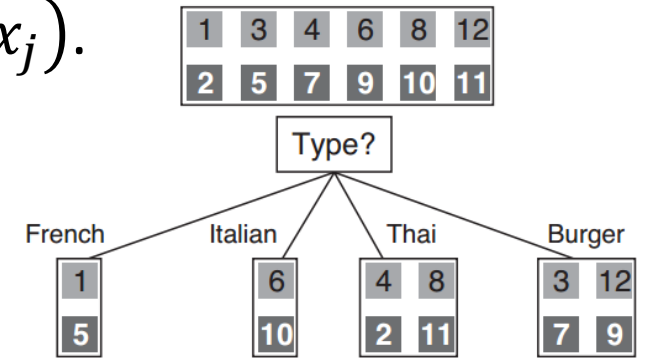


Image source: Figure 18.13 of the AI book by S. Russell & P. Novig.

# Preliminary: Conditional Entropy

- Conditional entropy:

$$\mathcal{H}(Y|X) \triangleq \sum_j p(X = x_j) \cdot \mathcal{H}(Y|X = x_j).$$

- Example 3:  $Y \sim \text{label}$  &  $X \sim \text{Patrons}$

- $p(X = \text{None}) = \frac{2}{12}$ ;  $p(X = \text{Some}) = \frac{4}{12}$ ;  $p(X = \text{Full}) = \frac{6}{12}$ ;
- $\mathcal{H}(Y|X = \text{None})$ :  $\frac{2}{2} \log \left( \frac{2}{2} \right) = 0$ ;
- $\mathcal{H}(Y|X = \text{Some})$ :  $\frac{4}{4} \log \left( \frac{4}{4} \right) = 0$ ;
- $\mathcal{H}(Y|X = \text{Full})$ :  $\frac{2}{6} \log \left( \frac{2}{6} \right) + \frac{4}{6} \log \left( \frac{4}{6} \right) = -0.9183$ ;
- $\mathcal{H}(Y|X) = - \left[ \frac{2}{12} \cdot 0 + \frac{4}{12} \cdot 0 + \frac{6}{12} \cdot (-0.9183) \right] = 0.4591$ .




Image source: Figure 18.4(b) of the AI book by S. Russell & P. Novig.

# Information Gain (信息增益)

---

- **Information gain**: Decrease in entropy after splitting

$$IG(X) = \mathcal{H}(Y) - \mathcal{H}(Y|X)$$

- $X$ : input feature,
- $Y$ : classification label.

- **Example 4**:  $Y \sim \text{label} \ \& \ X \sim \text{type/patrons}$

- $IG(\text{Type}) = \mathcal{H}(Y) - \mathcal{H}(\text{label}|\text{Type}) = 1 - 1 = 0.$
- $IG(\text{Patrons}) = \mathcal{H}(Y) - \mathcal{H}(\text{label}|\text{Patrons}) = 1 - 0.4591 = 0.541.$

# Information Gain: Type vs Patrons?

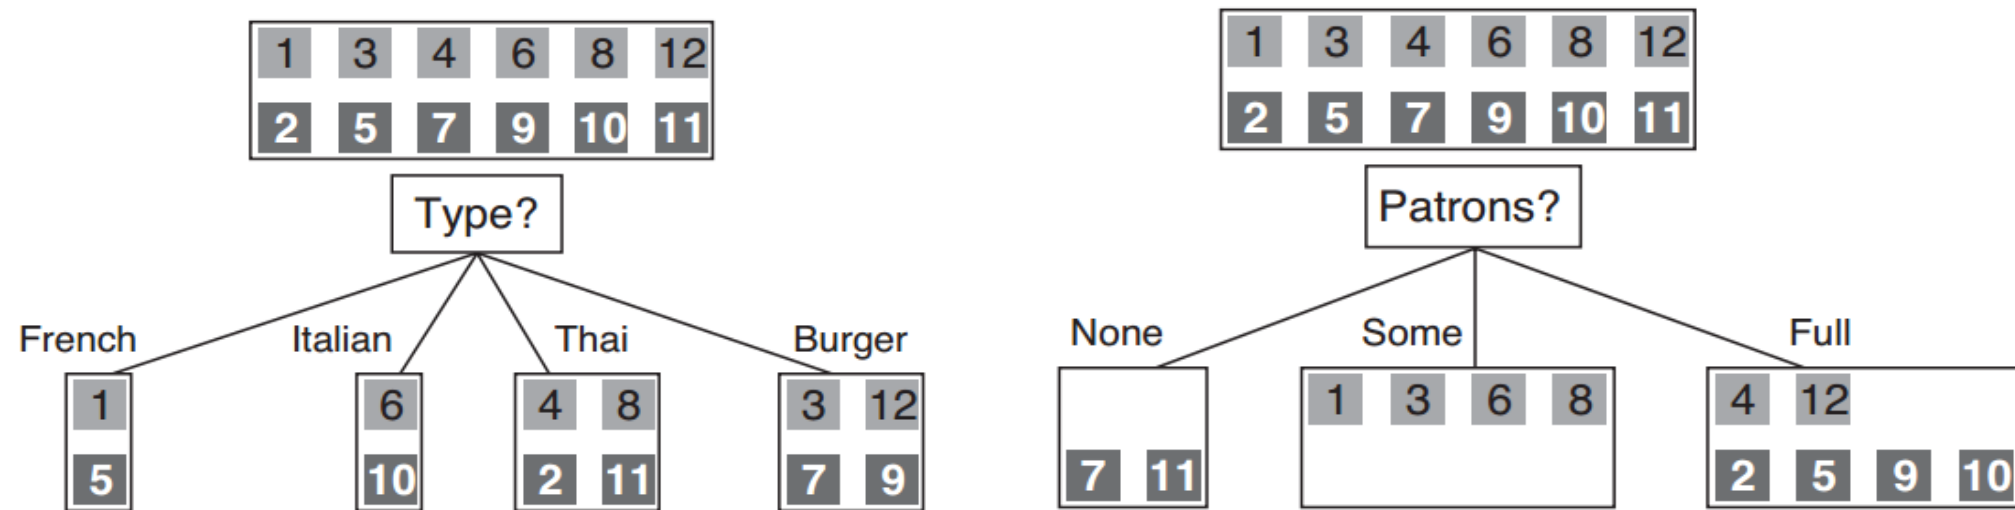


Image source: Figure 18.4 of the AI book by S. Russell & P. Novig.

- **[Answer]**  $IG(Patrons) > IG(Type) \Rightarrow$  Patrons is better.



# Continuous Feature $Est$

- **[Question]** What should we do for  $Est \in \mathbb{R}^1$ ?
  - $Est$ : estimated waiting time.
- **Binary tree**: Split on  $Est$  at value  $t$ ,
  - One branch:  $Est < t$ ,
  - Other branch:  $Est \geq t$ .
- **Note**: Allow repeated features along a path.

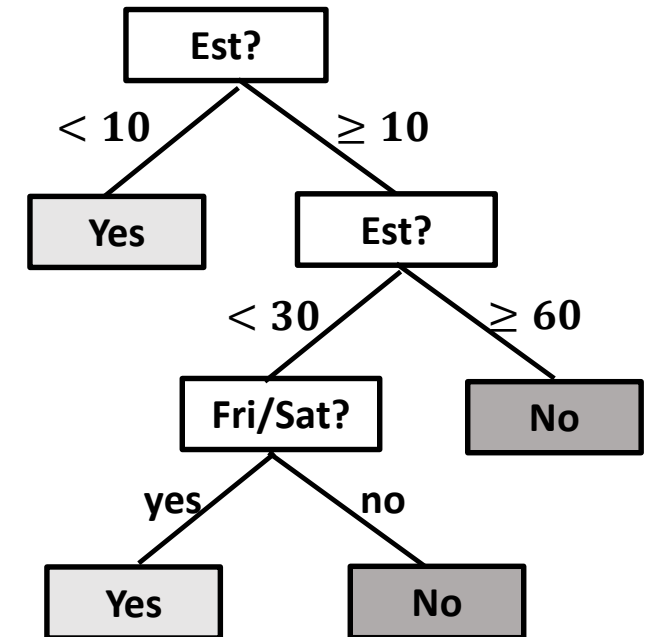


Figure generated by Liyan Song.

## Possible $\{t\}$ for $Est$

---

- [Question] How to decide the possible  $\{t\}$  for  $Est$ ?
- [Concern] Search through  $\mathbb{R}^1 \Rightarrow$  Too hard!
- [Answer] Only a finite number of values are useful:
  - Sort values of  $Est$  into  $\{x_1, \dots, x_m\}$  with non-duplicated values;
  - Consider candidates  $\left\{t_i = x_i + \frac{x_{i+1} - x_i}{2} \mid i = 1, \dots, m - 1\right\}$ .

## Best $t^*$ for $Est$ and its Information Gain

---

- Take the best  $t$  from  $\{t\}$ : Denote  $X \sim Est$ ,
  - (1) Define  $\mathcal{H}(Y|X:t) = p(X < t) \cdot \mathcal{H}(Y|X < t) + p(X \geq t) \cdot \mathcal{H}(Y|X \geq t)$ ;
  - (2) Compute  $IG(Y|X:t_i) = \mathcal{H}(Y) - \mathcal{H}(Y|X:t_i) \forall t_i$ ;
  - (3) Choose  $t^* = \arg \max_{t_i} IG(Y|X:t_i)$
- Use:  $IG^*(Est) = IG(Y|X:t^*) = \max_{t_i} IG(Y|X:t_i)$ .

# When to Stop?

---

- Criterion 1: all records in current subset have the same label.
- Criterion 2: there are no remaining features to help partitioning.
- Criterion 3: The associated dataset is empty.




Fig. 1 Criterion 1

*Image source: Figure 18.4 of the AI book by S. Russell & P. Novig.*

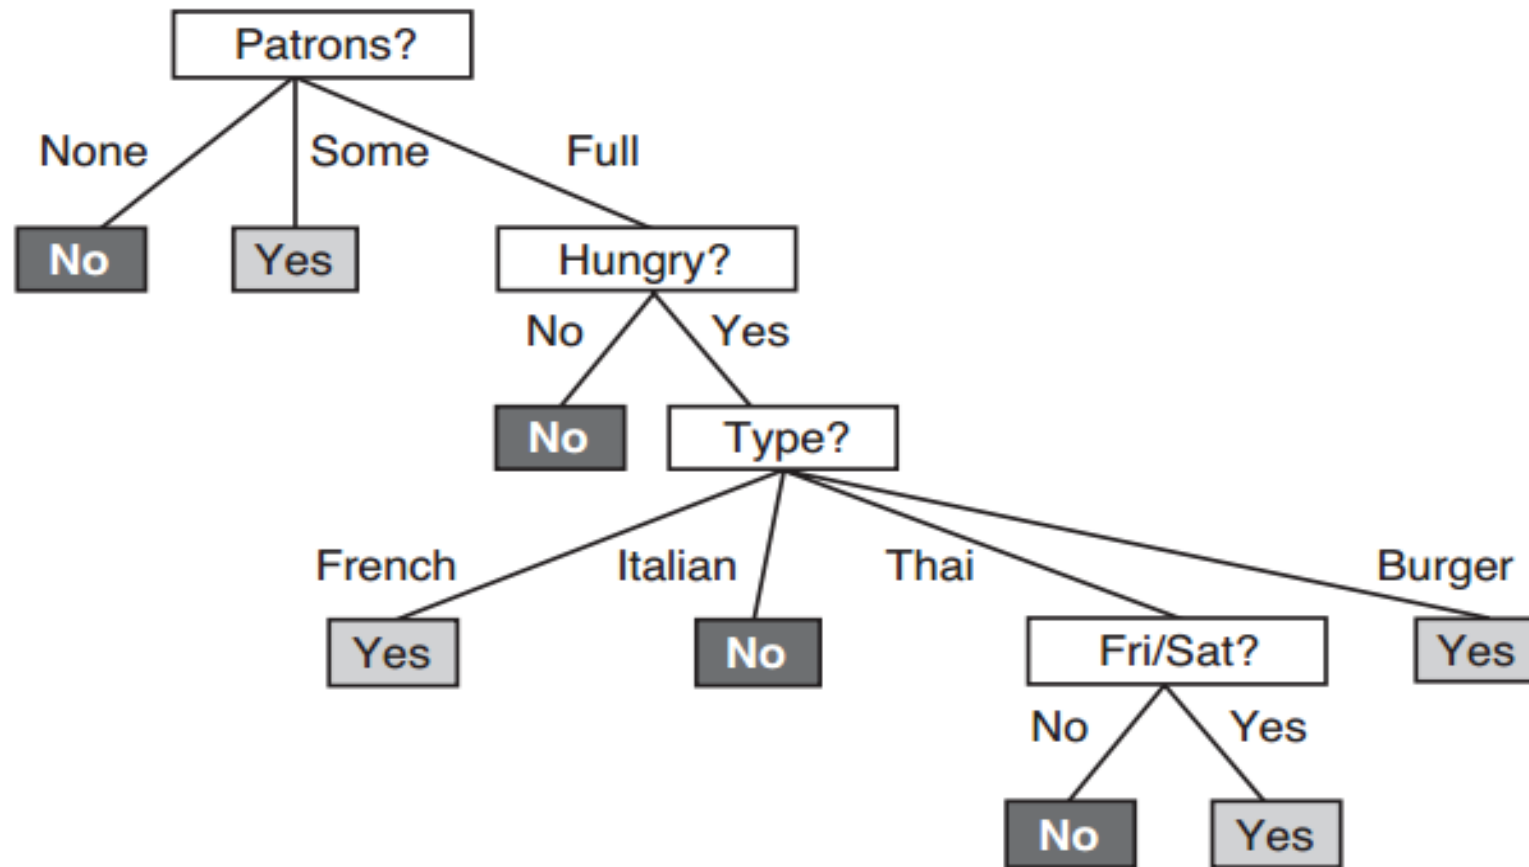
# Learning Decision Trees

---

- Start from empty tree.
- Split on next best feature based on *information gain*.
- Repeat

# An Example of Induced Decision Tree

---



*Fig. An estimated decision tree  $g$  induced from 12 examples. Image source: Figure 18.6 of the AI book by S. Russell & P. Novig.*


## II. Decision Tree (决策树)

1. Tree Representation
2. Decision Tree Construction with Heuristics
  - ❖ Information Gain: Good Feature Heuristics
  - ❖ Information Gain: Continuous Feature
  - ❖ Overall: Decision Tree Construction
- 3. Tree Overfitting**
4. Decision Tree for Regression

# Decision Tree May Overfit

---

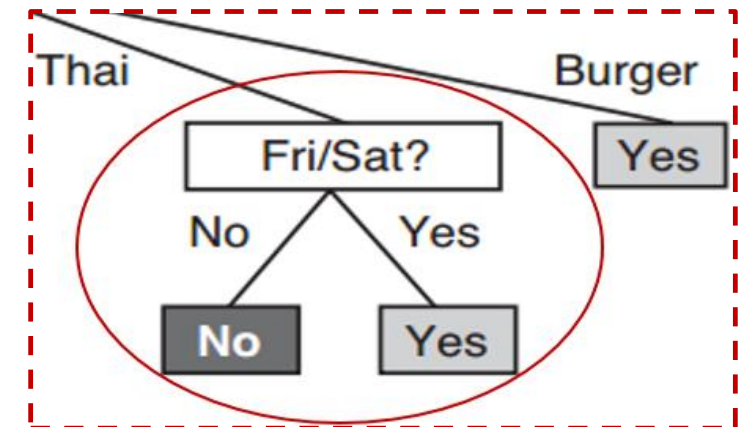
- More #feature, more likely overfitting; more #(train data), less likely overfitting.





# Dealing with Overfitting

- Decision tree pruning:
  - (1) Build a fully grown tree.
  - (2) Choose a node that has only leaf nodes as children.
  - (3) Testing the feature 'relevance' for this node:
    - (a) relevant→reserve this node.
    - (b) irrelevant: replace it based on its leaf nodes.
  - Repeat (2)~(3) until no such irrelevant nodes.
- Other strategies:
  - Fixed depth
  - Fixed #leaves



A testing node at step (2). Edited from Figure 18.6 of the AI book by S. Russell & P. Novig.

# Measure Feature Relevance

---

- [Question] How to detect that a node is testing an irrelevant feature?
- [Answer] the node splits the examples evenly & information gain is close to 0  $\rightarrow$  irrelevant feature.
- [Question] How large a gain should be required to split on the feature?
- [Answer] Using  $\chi^2$  statistical test, namely  $\chi^2$  pruning.

## II. Decision Tree (决策树)

1. Tree Representation
2. Decision Tree Construction with Heuristics
  - ❖ Information Gain: Good Feature Heuristics
  - ❖ Information Gain: Continuous Feature
  - ❖ Overall: Decision Tree Construction
3. Tree Overfitting
4. **Decision Tree for Regression**

# Example: Predict Car Price


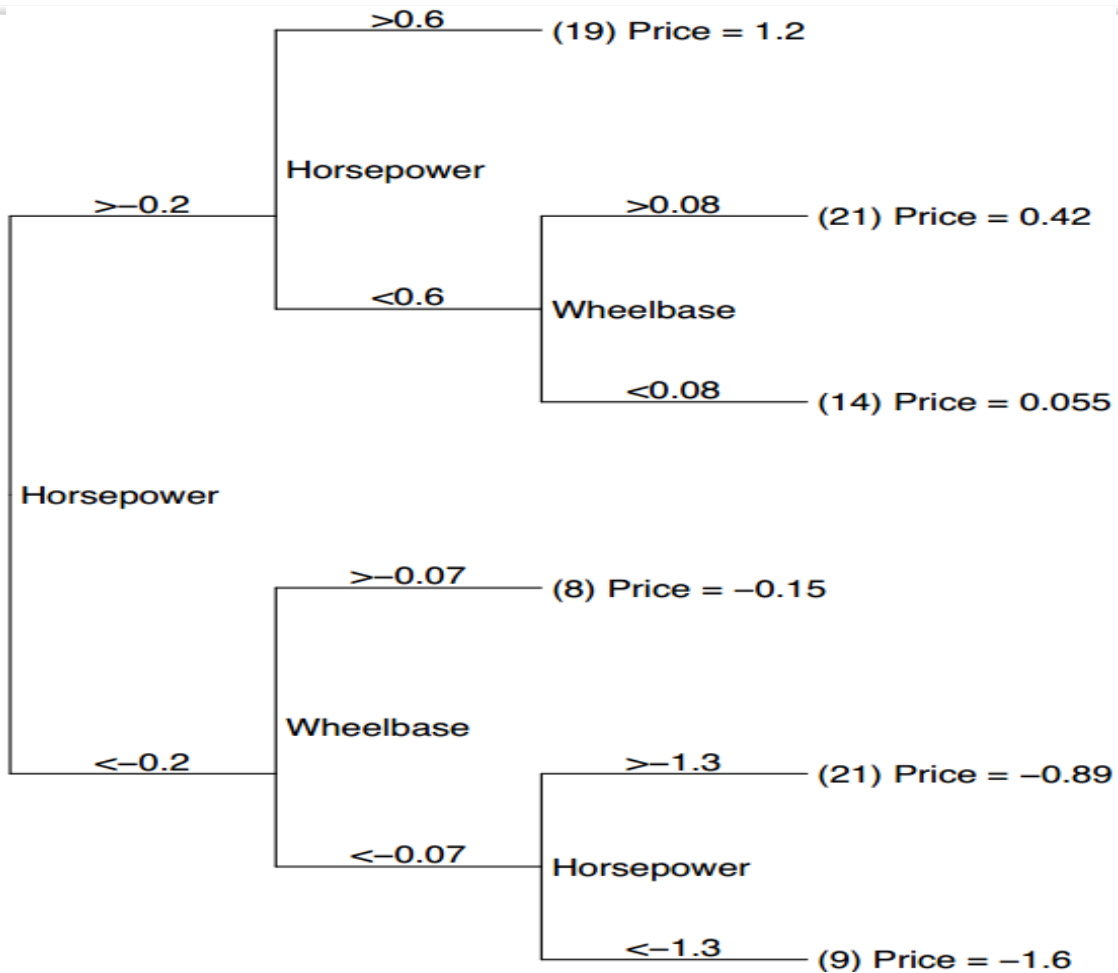


Image source: <http://www.stat.cmu.edu/~cshalizi/350-2006/lecture-10.pdf>

# Regression Tree (RT)

---

- **Construction**: Similar to classification trees except:
  - **Output**: continuous value.
  - **Leaf node**: (a) mean/median: piecewise constant RT; or (b) a regression model: piecewise linear RT.

# Piecewise Constant RT

---

- **Leaf node**: Mean of the examples in leaf node  $C$

$$\widehat{y}_C = \frac{1}{||C||} \sum_{i=1}^{||C||} y_i$$

- **Algorithms examples**: AID, CART.

# Piecewise Linear RT

---

- **Leaf node**: Linear regression model on the examples in each leaf node.
- **Algorithms examples**:
  - M5': (1) Construct a constant regression tree. (2) Fit a linear regression model for each leaf node.
  - GUIDE: (1) Fit a regression model (linear or nonlinear) and compute the residuals. (2) Label the examples with 1 for positive and 2 for negative residuals. (3) Apply the GUIDE for classification tree to split the node.

# RT Construction

---

- **Problem:** No labels to split features by  $IG(X_i)$ .

- **Feature splitting criteria:** Sum of squared error

$$SSE = \sum_{C \in L} \sum_{i \in C} (y_i - \widehat{y}_C)^2,$$

- $L$ : a set of leaf nodes.
- $\widehat{y}_C$ : the estimation on leaf node  $C$  from its examples.
- $y_i$  for  $i \in C$ : output of the  $i^{th}$  example of leaf node  $C$ .
- **Learning:** Search all **binary splits** that reduce  $SSE$  to the full.
- **When to stop:**  $SSE \leq \delta$  or fixed #leaves.



# III. Neural Network (神经网络)

# Artificial Neural Network: Formulation




Fig.(1) A brain neuron. Image source: Figure 1.2 of the AI book by S. Russell & P. Novig.




Fig.(2) An artificial neuron. Image source: Figure 18.9 of the AI book by S. Russell & P. Novig.

- $in_j = \sum_{i=0}^n w_{i,j} a_i$
- $a_j = g(in_j)$

# Activation Function (激活函数)

- **Physics**: Simulate the activation process of real neuron.
- **Math**: Nonlinear activation functions encodes the ability to estimate a nonlinear function from inputs to outputs.
- Popular activation functions:
  - hard threshold,
  - logistic function,
  - sigmoid
  - Tanh
  - ReLU, Leaky ReLU
  - .....





Image source: Figure 18.17 of the AI book by S. Russell & P. Novig.

# Single-layer Neural Network

---

- **Single-layer neural network:** All input neurons connect directly with the output neurons.



*Image source: Figure 18.20.a of the AI book by S. Russell & P. Novig.*

# Multilayer Neural Network

---


Output units  $O_i$

$W_{j,i}$

Hidden units  $a_j$

$W_{k,j}$


Input units  $I_k$



# Multilayer Neural Network

- Notation system:
  - Output index:  $1 \dots i \dots n$
  - Hidden index:  $1 \dots j \dots h$
  - Input index:  $1 \dots k \dots l$
  - Weights between output and hidden:  $w_{j,i}$
  - Weight between hidden and input:  $w_{k,j}$
  - Activation function of output units:  $\sigma()$
  - Activation function of hidden units:  $g()$
  - $h_w$ : the (non-linear) function the NN represents.
- Exercise: express  $o_i$  with the above notation system.

Output units  $O_i$   
 $w_{j,i}$   
Hidden units  $a_j$   
 $w_{k,j}$   
Input units  $I_k$



# Learning Multilayer Networks

---

- **Loss function** for an example:  $\ell_2(\mathbf{w}) = \frac{1}{2} ||y - o(\mathbf{x})||^2$

- $(\mathbf{x}, y)$ : a training example;
- $o(\mathbf{x})$ : estimated output for input  $\mathbf{x}$ .

- **Partial derivative for any  $w$** : ‘chain rule’

$$\frac{\partial}{\partial w} \ell_2(\mathbf{w}) = \frac{\partial}{\partial w} \frac{1}{2} \sum_i (y_i - o_i)^2 = - \sum_i (y_i - o_i) \frac{\partial o_i}{\partial w} = \dots$$

- **Back propagation (反向传播/逆传播)** to train ANN:

- Gradient descent for  $w_{j,i}$  from hidden to output:  $w_{j,i} \leftarrow w_{j,i} - \alpha \cdot \frac{\partial}{\partial w_{j,i}} \ell_2(\mathbf{w})$
- Gradient descent for  $w_{k,j}$  from input to hidden:  $w_{k,j} \leftarrow w_{k,j} - \alpha \cdot \frac{\partial}{\partial w_{k,j}} \ell_2(\mathbf{w})$

# Learning ANN Structures

---

- Fully connected networks: decide #hidden layers and their sizes using **cross-validation**.
- Not fully connected networks: **optimal brain damage** algorithm begins with a fully connected network and removes connections from it.
- Grow a larger network from a smaller one: Subsequent units are added to cater for the examples that the first unit got wrong in the **tiling** algorithm.




# IV. $k$ -Nearest Neighbor ( $k$ -近邻)

# $k$ -NN

---

- $k$ -Nearest neighbor method:
  - For classification: find  $k$  nearest neighbors of the testing point and take a vote.
  - For regression: take mean or median of the  $k$  nearest neighbors, or do a local regression on them.



Example of  $k$ -NN classification. Image source: <https://en.wikipedia.org/wiki/File:KnnClassification.svg#filelinks>.

# $k$ -NN Issues

---

- Distance metric: e.g.,  $\ell_p(\mathbf{x}_j, \mathbf{x}_q) = \left( \sum_i |x_{j,i} - x_{q,i}|^p \right)^{1/p}$ 
  - $\ell_1$ : Manhattan distance,
  - $\ell_2$ , Euclidean distance.
- Model parameter  $k$ : increasing  $k$  reduces variance and increases bias.
- Memory-based method: must store all training samples.

# $k$ -NN Issues

---

- Advantage:
  - Training is very fast.
  - Learn complex target functions.
  - Do not lose information.
- Disadvantage:
  - Slow at query time.
  - Easily fooled by irrelevant attributes.

# V. Support Vector Machine (支持向量机)

# Geometry and SVM Formulation

- Input:  $\mathbf{x}$
- Output:  $y$  ( $-1$  or  $1$ )
- Model:  $\{\mathbf{w}^T \mathbf{x} + b = 0\}$
- Two Boundaries:

$$\{\mathbf{w} \cdot \mathbf{x} + b = \pm 1\}$$

- Margin:  $d = \frac{2}{\|\mathbf{w}\|^2}$




Image source: Figure 5.25 of "Introduction to Data Mining" by P. Tan, M. Steinbach and V. Kumar.

# Maximize the Margin




---

- **Training Data:**  $\mathcal{D} = \{(\mathbf{x}^{(n)}, y^{(n)})\}_{n=1}^N$ .
- **Optimization:** maximize the margin with the constraints as

$$\begin{aligned} \max_{\mathbf{w}} \quad & \frac{2}{\|\mathbf{w}\|^2}, \\ \text{s.t.} \quad & [\mathbf{w} \cdot \mathbf{x}^{(n)} + b] \cdot y_i^{(n)} \geq 1 \end{aligned}$$

- **Learning algorithm [3]:**
  - Lagrange multiplier with KKT condition  $\Rightarrow$  Dual representation.
  - Gradient descent.

# Kernel Trick



*Image source: Figure 18.31 of the AI book by S. Russell & P. Novig.*

**Kernel trick:** nonlinear-separable feature space  $\Rightarrow$  linear-separable one.



# Reading Materials for This Lecture

---

- [1] AI book (P693-748).
- [2] P. Tan, M. Steinbach and V. Kumar. *Introduction to Data Mining* (pages 223-225, 256-276).
- [3] Laurent H, Rivest R L. *Constructing optimal binary decision trees is NP-complete*. Information processing letters, 1976, 5(1): 15-17.
- [4] Gradient Descent: <http://runder.io/optimizing-gradient-descent/>
- [5] Classification and Regression Trees: <http://www.stat.wisc.edu/~loh/treeprogs/guide/wires11.pdf>
- [6] <http://scikit-learn.org/stable/modules/ensemble.html>