

Advanced Routing



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Xuetao Wei

weixt@sustech.edu.cn

Contents



- Routing Basics
- Intra-AS Routing
- Inter-AS Routing
- Multicast Routing
- MPLS

router : 1.control plane 2. data plane

quiz考前三个

查看

traceroute

目标 IP : 64.238.147.121

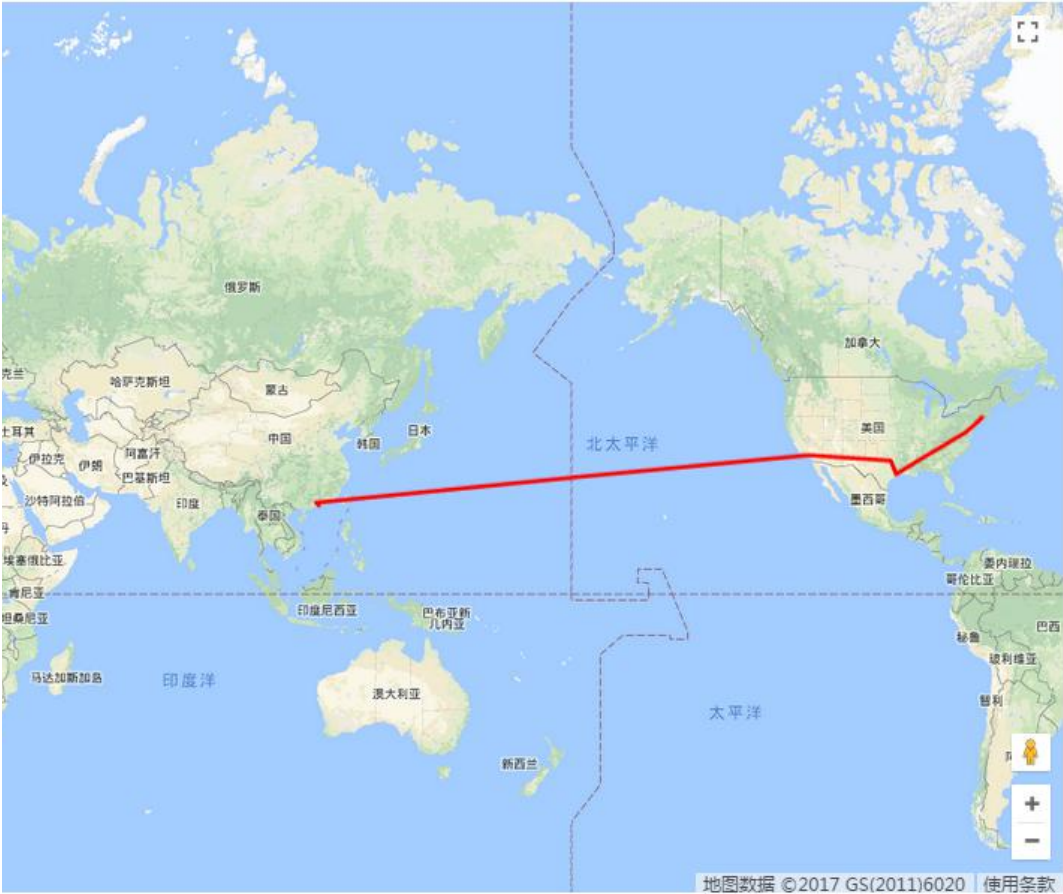
监测点赞助商 : 速云

跳数	IP	主机名	地区 (仅供参考)	AS号 (仅供参考)	时间 (毫秒)
1	172.17.11.254	172.17.11.254	局域网		1.2 / 1.1 / 1
2	172.17.2.1	172.17.2.1	局域网		0.4 / 0.8 / 1.2
3	61.144.170.169	61.144.170.169	广东深圳 电信	AS4134	1.3 / 1 / 1.8
4	219.133.30.82	219.133.30.82	广东深圳 电信	AS4134	2.1 / 2.1 / 1.8
5	58.60.24.41	58.60.24.41	广东深圳 电信	AS4134	2 / 2 / 2.1
6	183.56.65.70	183.56.65.70	广东深圳 电信	AS4134	7.6 / 5.3 / 7.1
7	*	*	N/A	*	*
8	202.97.94.114 * *	202.97.94.114 * *	广东广州 电信 N/A N/A	AS4134 * *	7.6 * *
9	202.97.51.106	202.97.51.106	美国加利福尼亚州洛杉矶 电信	AS4134	249.1 / 247.1 / 248.8
10	202.97.49.154	202.97.49.154	美国加利福尼亚州洛杉矶 电信	AS4134	282.1 / 280.7 / 277.7
11	218.30.54.174	218.30.54.174	美国加利福尼亚州洛杉矶 CHINATELECOM	AS4134	209 / 210.2 / 197.5
12	64.125.28.230	ae13.cs1.lax112.us.eth.zayo.com	美国加利福尼亚州洛杉矶 zayo.com	AS6461	284.9 / 277.2 / 284.6
13	64.125.29.52	ae3.cs1.dfw2.us.eth.zayo.com	美国德克萨斯州达拉斯 zayo.com	AS6461	271.1 / 281.3 / 248.1
14	64.125.28.98 64.125.28.98 *	ae5.cs1.iah1.us.eth.zayo.com ae5.cs1.iah1.us.eth.zayo.com *	美国德克萨斯州休斯顿 zayo.com 美国德克萨斯州休斯顿 zayo.com N/A	AS6461 AS6461 *	320.5 320.3 *
15	64.125.29.48	ae3.cs1.dca2.us.eth.zayo.com	美国华盛顿 zayo.com	AS6461	261.4 / 257.2 / 269.4
16	64.125.29.202 204.16.61.21 204.16.61.21	ae4.cs1.lga5.us.eth.zayo.com 204.16.61.21 204.16.61.21	美国纽约州纽约 zayo.com 美国康涅狄格州 cyrusone.com 美国康涅狄格州 cyrusone.com	AS6461 AS62 AS62	343.5 315.5 306.1
17	64.125.22.9	ae12.mpr1.bdl4.us.zip.zayo.com	美国康涅狄格州哈特福德 zayo.com	AS6461	320.4 / 320.4 / 320.4
18	64.238.144.18	64.238.144.18	美国 cyrusone.com	AS19479	375.4 / 374.8 / 372.9
19	204.16.61.21	204.16.61.21	美国康涅狄格州 cyrusone.com	AS62	262 / 270.5 / 266.9
20	64.238.147.121	64.238.147.121	美国 cyrusone.com	AS19479	375.4 / 380.4 / 377




南方科技大学

SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



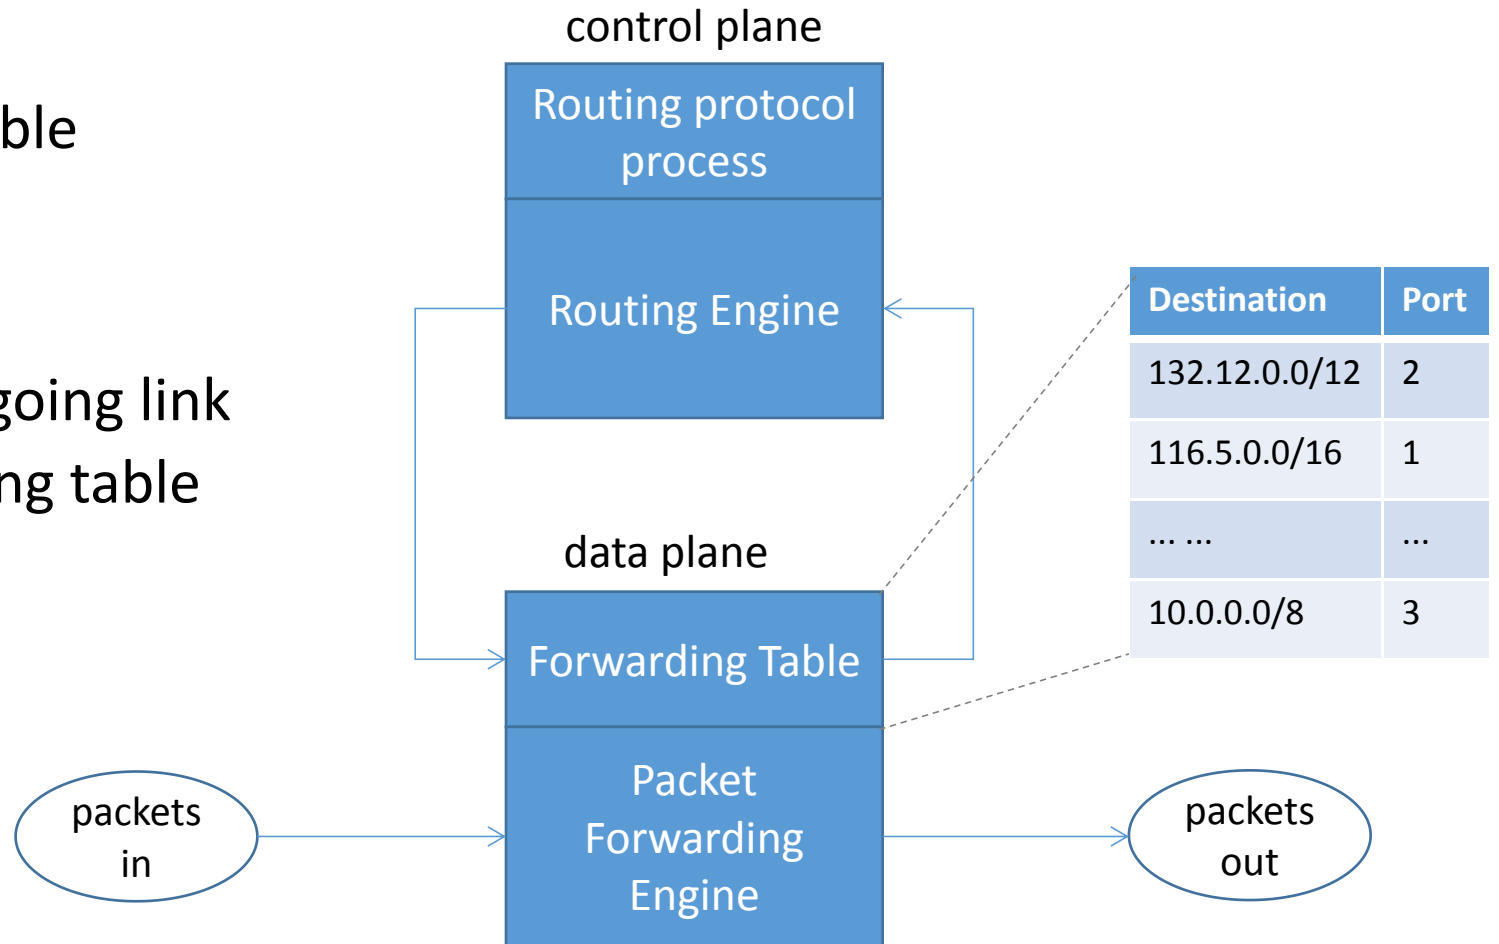
Routing: The Big Picture



Router



- **Routing:** control plane
 - Creating a forwarding table
 - Talking to other routers
- **Forwarding:** data plane
 - Send a packet to an outgoing link
 - Looking up the forwarding table



IP Addressing

- IP address: hierarchial or flat?
 - IPv4: dotted decimal format (a.b.c.d)
 - Internet is inherently hierarchical: Multiple ASes -> AS -> subnet
 - IP address: network + host
- Class-based addressing: class + network + host

Class A

0	Network	Host
---	---------	------

 127 nets, 16M hosts

Class B

		14			16
1	0	Network			Host

 16K nets, 64K hosts

Class C

			21				8
1	1	0	Network				Host

 2M nets, 254 hosts

Class D (1110) for multicast

Class E (1111) for experimental

CIDR Addressing



- CIDR: Classless Inter-Domain Routing
 - Allows arbitrary length of network prefix




- e.g. 132.12.1.8/12 (**10000100 0000**1100 00000001 00001000)
 - Network**
 - Host

- Pros: flexible, better utilization of IP resources
- Cons: more complicated lookup: **longest prefix match**

Autonomous System

- Autonomous System (AS)
 - Unit of routing policy
 - ~58k ASes in use
 - E.g., AT&T has AS#144, Princeton has AS#88, CHINANET(backbone) has AS#4134
- AS relationships
 - Peering: jointly pay for costs
 - Transit: customer pay for upstream and downstream traffic



Routing Protocols in Scenarios



- Intra-AS routing
 - Distance Vector Algorithms (RIP)
 - Link State Algorithms (OSPF)
- Inter-AS routing
 - Border Gateway Protocol (BGP)


Contents



- Routing Basics
- Intra-AS Routing
- Inter-AS Routing
- Multicast Routing
- MPLS

Intra-AS Routing

- Routing
 - Essentially a graph theory problem
 - Find “best” path between every pair of vertices (routers)



Global vs Decentralized Routing



- **Global routing algorithm**

- Global view and complete knowledge of the network topology
- Complete path for every pair of nodes
- Also known as **Link State routing** (e.g., OSPF)

- **Decentralized algorithm**

- Local view on the network
- No complete path at each node, only the next hop to destination
- Also known as **Distance Vector routing** (e.g., RIP)


1. Distance Vector Routing

- A distance vector $\{d_x(y) \mid y \in N\}$ at each node x
 - $d_x(y)$ = cost of least-cost path from X to Y
- Rippling vs flooding
 - Only notify neighbors on its updates, no explicit flooding to all nodes
 - But if a neighbor y gets updated as well, the update is rippled to y 's neighbors
- Bellman-Ford Algorithm
 - $d_x(y) = \min \{c(x,v) + d_v(y)\}$ over all neighbors V

Example: Initial State

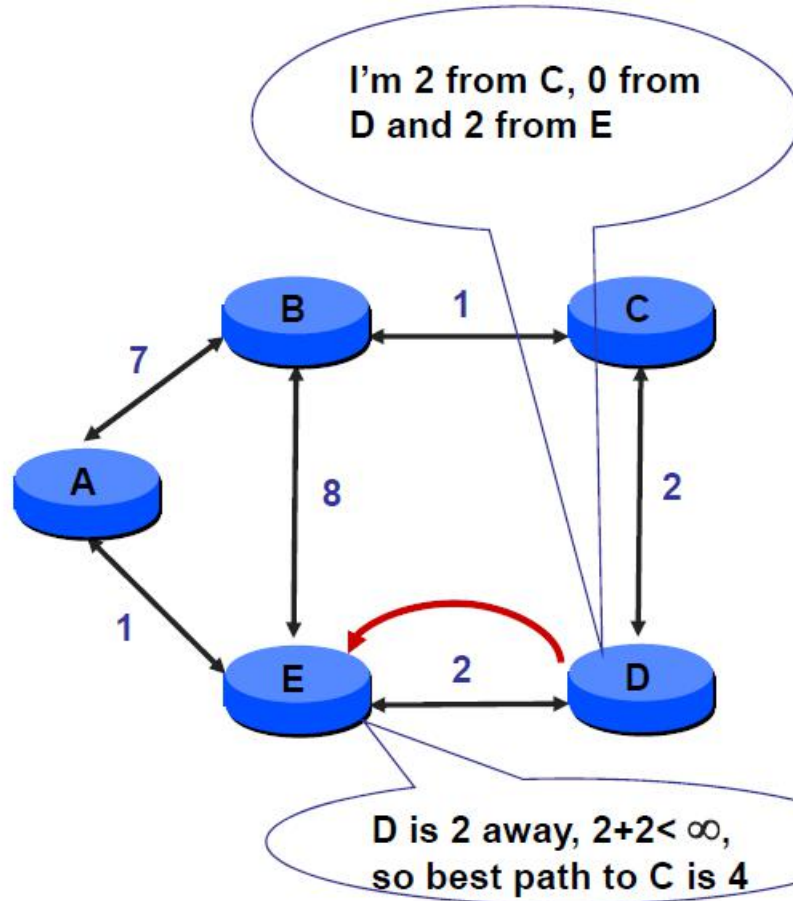


可能考



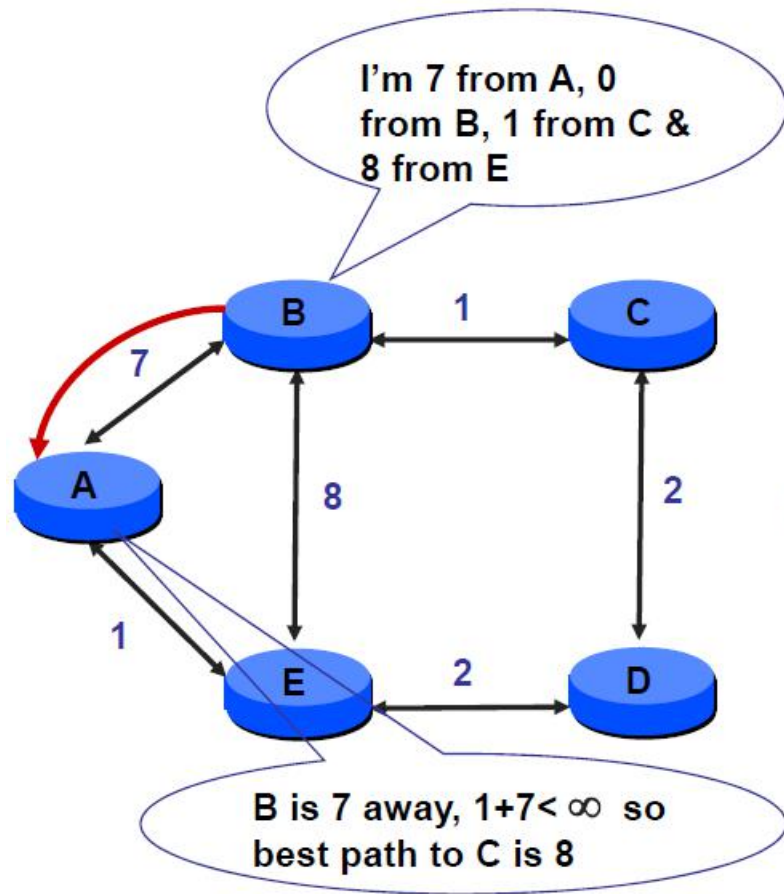
Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	∞	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	∞	2	0

D Sends Vector to E




Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	∞	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0

B Sends Vector to A



Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	8	∞	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0


E Sends Vector to A



Info at node	Distance to Node				
	A	B	C	D	E
A	0	7	5	3	1
B	7	0	1	∞	8
C	∞	1	0	2	∞
D	∞	∞	2	0	2
E	1	8	4	2	0


I'm 1 from A, 8 from B, 4 from C, 2 from D & 0 from E

... Until Convergence




Info at node	Distance to Node				
	A	B	C	D	E
A	0	6	5	3	1
B	6	0	1	3	5
C	5	1	0	2	4
D	3	3	2	0	2
E	1	5	4	2	0

Node B's Distance Vectors



Dest	Next hop		
	A	E	C
A	7	9	6
C	12	12	1
D	10	10	3
E	8	8	5


Worse: Count to Infinity



Why So and How to Solve?

- Why So?
 - Updates don't contain enough information
 - B accepts A's path to C that is implicitly through B!
- How to Solve?
 - **Split horizon**: never advertise a destination through its next hop
(Only works for node pairs)

Split Horizon




Why So and How to Solve?

- Why So?
 - Updates don't contain enough information
 - B accepts A's path to C that is implicitly through B!
- How to Solve?
 - **Split horizon**: never advertise a destination through its next hop
(Only works for node pairs)
 - **Split horizon + poison reverse**: send negative information (infinity cost) when advertising a destination through its next hop
(Only works for a loop of size 2)

Only Works for Loop of Size 2



- A tells B & C that D is unreachable
- B computes new route through C
 - Tells C that D is unreachable (poison reverse)
 - Tells A it has path of cost 3 (split horizon doesn't apply)
- A computes new route through B
- A tells C that D is now reachable



Why So and How to Solve?



- Why So?
 - Updates don't contain enough information
 - B accepts A's path to C that is implicitly through B!
- How to Solve?
 - **Split horizon**: never advertise a destination through its next hop
(Only works for node pairs)
 - **Split horizon + poison reverse**: send negative information (infinity cost) when advertising a destination through its next hop
(Only works for a loop of size 2)
 - **Hold down**: when a route is removed, no update of it accepted for some period of time (hold-down time) - to give everyone a chance to remove the route

Real DV Protocols

- RIP: Routing Information Protocol
 - DV protocol with hop count as metric
 - Infinity value is 16 hops; limits network size
 - Routers send vectors every 30 seconds
 - With triggered updates for link failures

2. Link State Routing

- Tell all routers the topology and each computes best paths for source-destination pair of nodes
- Two phases
 1. Topology dissemination (by flooding)
 2. Shortest-path calculation (Dijkstra's algorithm)

1. Flooding




- **Link State Packet (LSP)**

- ID of the node that created the LSP
- Cost of link to each directly connected neighbor
- Sequence number (SEQNO)
- Time-to-live (TTL) for this packet

一定要知道区别

- **Relible flooding**

- Store most recent LSP from each node
- Forward LSP to all nodes but one that sent it
- Generate new LSP periodically (increment SEQNO)
- Decrement TTL of each stored LSP (discard when TTL = 0)



Why “Reliable”?




- **Complications**

- Packet loss
- Out-of-order arrival

- **Solutions**

- Acknowledgments and retransmissions
- Sequence numbers
- Time-to-live for each packet



2. Shortest-path Calculation



Dijkstra's Algorithm:

s source node.


D_n cost of the least-cost path from node s to node n

```
M = {s};  
for each  $n \notin M$   
     $D_n = d_{sn}$ ;  
while (M  $\neq$  all nodes) do  
    Find  $w \notin M$  for which  $D_w = \min\{D_j ; j \notin M\}$ ;  
    Add  $w$  to  $M$ ;  
    for each  $n \notin M$   
         $D_n = \min_w [ D_n, D_w + d_{wn} ]$ ;  
        Update route;  
enddo
```

Example



Construct the routing table at node 1 using Dijkstra's algorithm




Iteration	M	D ₁	D ₂	D ₃	D ₄	D ₅	D ₆
Init							

Example





Construct the routing table at node 1 using Dijkstra's algorithm



	M	D1	D2	D3	D4	D5	D6
0	{1}	0	2	5	1	inf	inf
1	{1,4}	0	2	4	1	2	inf
2	{1,4,2,5}	0	2	3	1	2	4
3	{1,4,2,5,3}	0	2	3	1	2	4
4	{1,4,2,5,3,6}	0	2	3	1	2	4

Result Routing Table



The tree is translated into a routing table at node 1:

<u>Destination</u>	<u>Next Hop</u>
2	2
3	4
4	4
5	4
6	4

Real Link State Protocols



- **OSPF (Open Shortest Path First) and IS-IS**
 - Most widely used intra-domain routing protocols
 - Run by almost all ISPs and many large organizations
- **Implementational add-ons of OSPF**
 - Adds notion of areas for scalability
 - Area 0 is “backbone” area
 - Traffic between two areas must always go through area 0
 - Only need to know how to route within area
 - Load balancing: multiple equal cost routes

LS Routing vs DV Routing



- **DV: Tell your neighbors about the world**

- Global, consistent information
- Easy to get confused
- Slow convergence due to ripples and hold down
- Simpler sometimes (only update neighbors on some DV changes)

各自有什么优点和缺点

- **LS: Tell the world about your neighbors**


- Iterative, asynchronous and distributed information
- Harder to get confused
- Faster convergence (instantaneous update of link state changes)
- More expensive sometimes (flooding to every one on link changes)

Contents



- Routing Basics
- Intra-AS Routing
- Inter-AS Routing
- Multicast Routing
- MPLS

Inter-AS Routing: Scalability Problem



Issues with Link State




- Topology information is flooded
 - High bandwidth and storage overhead
- Entire path computed locally per node
 - High processing overhead in a large network
- Can't express policy
- What about Distance Vector?
 - Hides details of the network topology
 - Nodes determine only “next hop” toward the destination
 - But, slow convergence due to the counting-to-infinity problem
 - Idea: extend the notion of a distance vector

Path Vector Routing




- Key idea: advertise the entire path (sequence of ASes)
 - Distance vector: send distance metric per destination
 - Path vector: send the entire path for each destination



Loop Detection



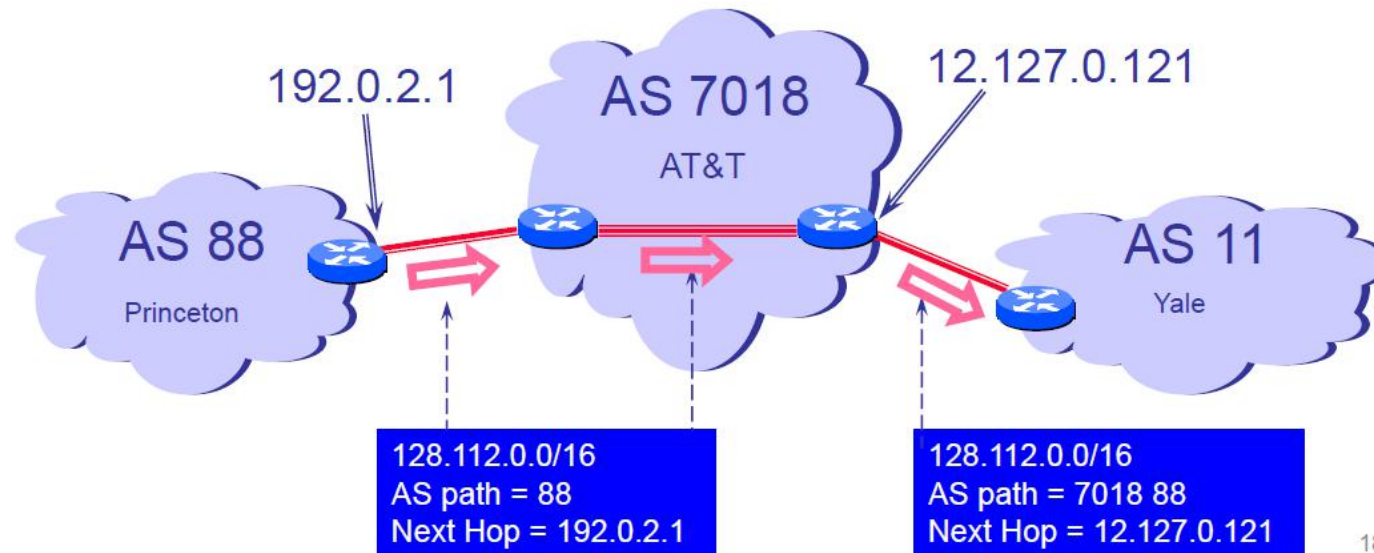
- Node can easily detect a loop
 - Look for its own node identifier in the path
 - E.g., node 1 sees itself in the path “3, 2, 1”
- Node can simply discard paths with loops
 - E.g., node 1 simply discards the advertisement



Border Gateway Protocol (BGP)



- BGP uses Path Vector Routing
- A simple BGP route with destination prefix: e.g., 128.112.0.0/16
 - Route attributes, including AS path (e.g., “7018 88”) and next-hop IP address (e.g., 12.127.0.121)



Contents




- Routing Basics
- Intra-AS Routing
- Inter-AS Routing
- Multicast Routing quiz不要求
- MPLS


Multicasting: Motivation



- Efficient delivery to multiple destinations (e.g. video broadcast)
- One-to-many addressing is needed for reducing traffic!



VS



IP Multicast service model



- Communications based on groups
 - Special IP addresses (Class D in IPv4) represent “multicast groups”
 - Anyone can join group to receive packets
 - Anyone can send to group (senders need not be part of group)
- Unreliable datagram service
 - Extension to unicast IP
 - Group membership not visible to hosts
 - No synchronization


Elements of IP Multicast

- Host interface
 - Application visible multicast API
 - Multicast addressing
 - Link-layer mapping
- Host-Router interface
 - IGMP
- Router-Router interface
 - Multicast routing protocols

Internet Group Management Protocol (IGMP)



- Goal: Communicate group membership between hosts and routers
 - Hosts explicitly inform their router about membership
 - Must periodically refresh membership report



Router broadcasts membership query to 224.0.0.1

Host sends membership report to group G when its timer expires


Router-Router Interface



- Protocols for multicast packet forwarding
- A simple solution
 - Flood packets from a source to entire network
 - If a router has not seen a packet before, forward it to all interfaces except the incoming one

Problems


1. Routers receive duplicate packets
2. Detecting that a packet is a duplicate requires storage, which can be quite expensive



A Clever Solution




- Reverse path forwarding
 - Forward packet from S to all interfaces if and only if packet arrives on the interface that corresponds to the shortest path *to* S
 - No need to remember past packets (C needs not forward packet received from D)



Cleverer

- Don't send a packet downstream if you are not on the shortest path from the downstream router to the source
- C needs not forward packet from A to E



Summary



- Routing Basics
- Intra-AS Routing
- Inter-AS Routing
- Multicast Routing
- MPLS