

# Transport Layer Issues



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

**Xuetao Wei**

weixt@sustech.edu.cn

# Contents



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- Basics of Transport Layer
- Reliable Transportation
- Flow Control
- Congestion Control
- Multipath TCP

quiz考前四个

# Network Layer vs Transport Layer



- **Network layer**

- Host-to-host communication
- Host addressing
- Routing


- **Transport layer**

- App-to-app communication
- Reliable communication
- Flow & Congestion control

# Logical Communication




- Socket = <IP address, port number>
- A pair of sockets **<client IP, server port>** and **<server IP, server port>** identify a transport layer connection (app to app)




# Packet vs Circuit Switching



**Mechanism at Transport Layer  
for Reliable Transmission!**



Unreliable Transmission



Reliable Transmission

# TCP vs UDP



- TCP: Transport Control Protocol
  - Reliable bidirectional stream of bytes

重点考tcp

- UDP: User Datagram Protocol
  - Simple (unreliable) message delivery

**TCP Segment Header Format**

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Sequence Number							
64	Acknowledgment Number							
96	Data Offset	Res	Flags		Window Size			
128	Header and Data Checksum				Urgent Pointer			
160...	Options							

**UDP Datagram Header Format**

Bit #	0	7	8	15	16	23	24	31
0	Source Port				Destination Port			
32	Length				Header and Data Checksum			

# UDP Application

- Features
  - Simple: no connection establishment (which can add delay)
  - Small packet header overhead (eight bytes long)
  - No congestion control (packets are sent as soon as possible)
- Applications
  - Multimedia streaming
  - Simple query protocols like DNS

But we often need **reliable & disciplined** data transmission => **TCP**

# TCP at Glance

TCP的特点，这四个要记住



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- **Connection oriented**
  - Explicit set-up and tear-down of TCP session
- **Reliable, in-order delivery**
  - Checksums to detect corrupted data
  - Sequence numbers to detect losses and reorder data
  - Acknowledgments & retransmissions for reliable delivery
- **Flow control** 重点
  - Prevent overflow of the receiver's buffer space
- **Congestion control**
  - Adapt to network congestion for the greater good



# Contents

- Basics of Transport Layer
- Reliable Transportation
- Flow Control
- Congestion Control
- Multipath TCP

# In-order Delivery of Packets



- **Option 1**

- Packets are required to arrive at the receiver **in-order**
- Almost impossible to achieve (for packet switching network)

- **Option 2**

- Packets may arrive **out-of-order**, but get identified by the receiver **in-order**
- More relaxed requirement and is possible to achieve, but HOW?

Answer => **sequence number** for packets

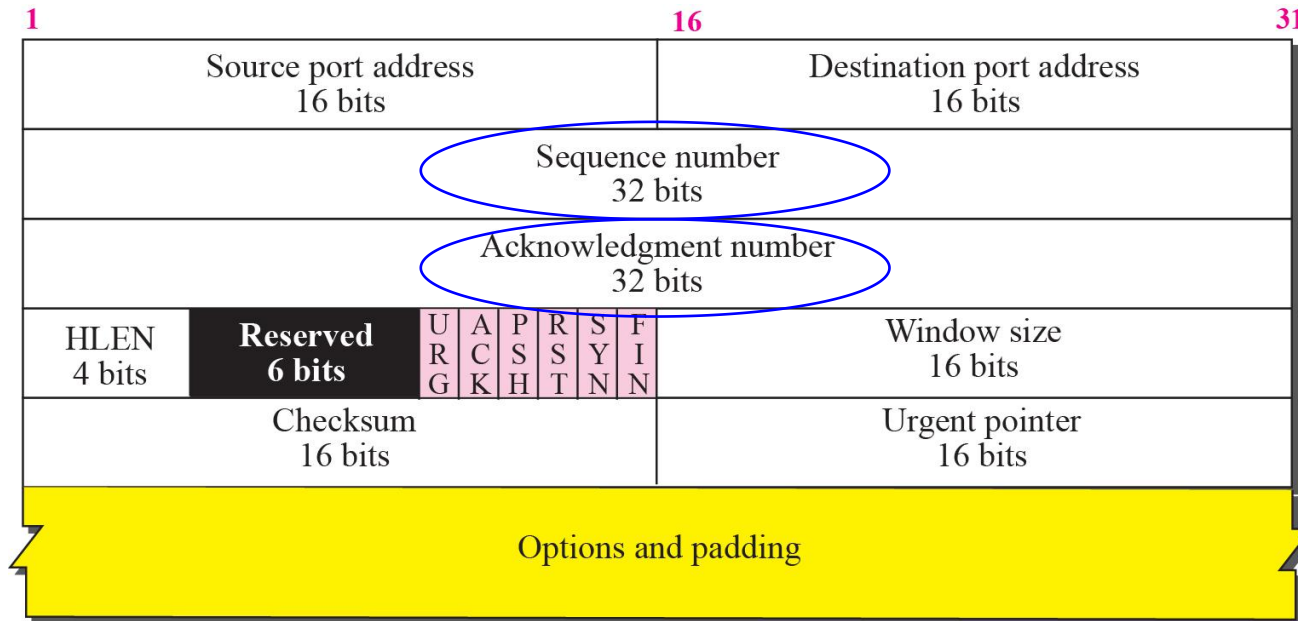
# TCP Segment



- IP Packet = IP Header + IP Data (TCP Segment)



a. Segment



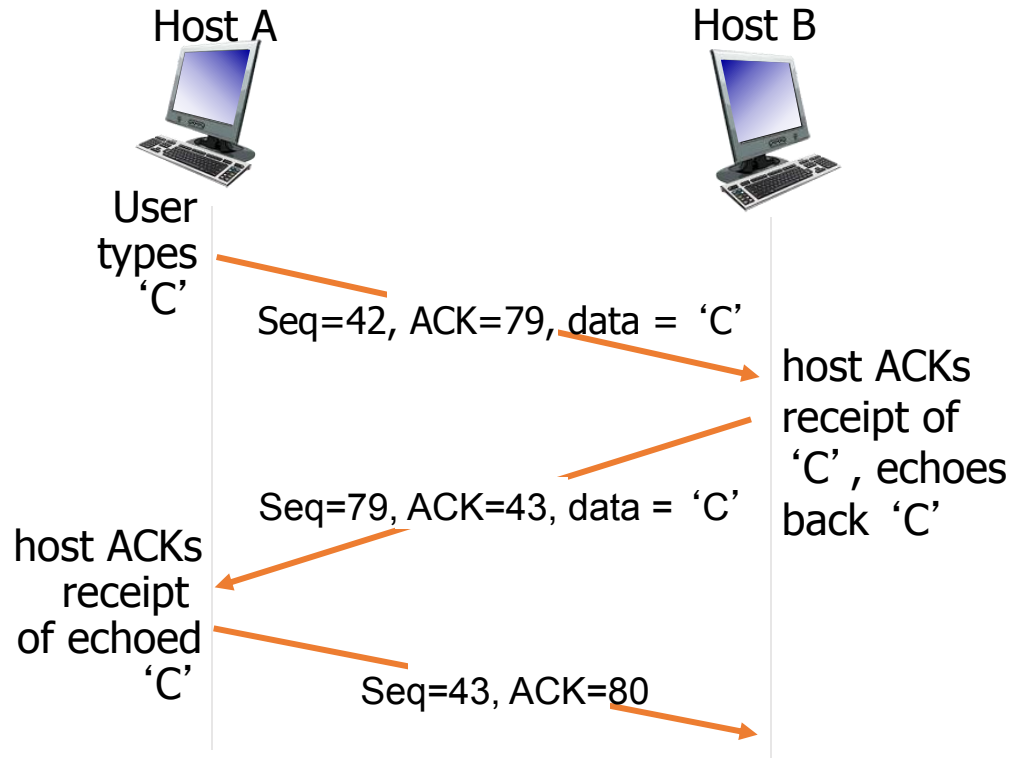
b. Header

# Sequence Number



- **Sequence number (SeqNo)**
    - 32 bits long ( $0 \leq \text{SeqNo} \leq 2^{32} - 1$ )
    - Each SeqNo identifies a byte (not a segment) in the byte stream
  - **Acknowledgement Number (AckNo)**
    - An A -> B segment can contain an acknowledgement for a B -> A segment earlier
    - The AckNo contains the next SeqNo that a host wants to receive
- Example: The acknowledgement for byte stream 0-1500 is AckNo=1501

# TCP: Sequence Numbers and Acks



simple telnet scenario

# Connection Management

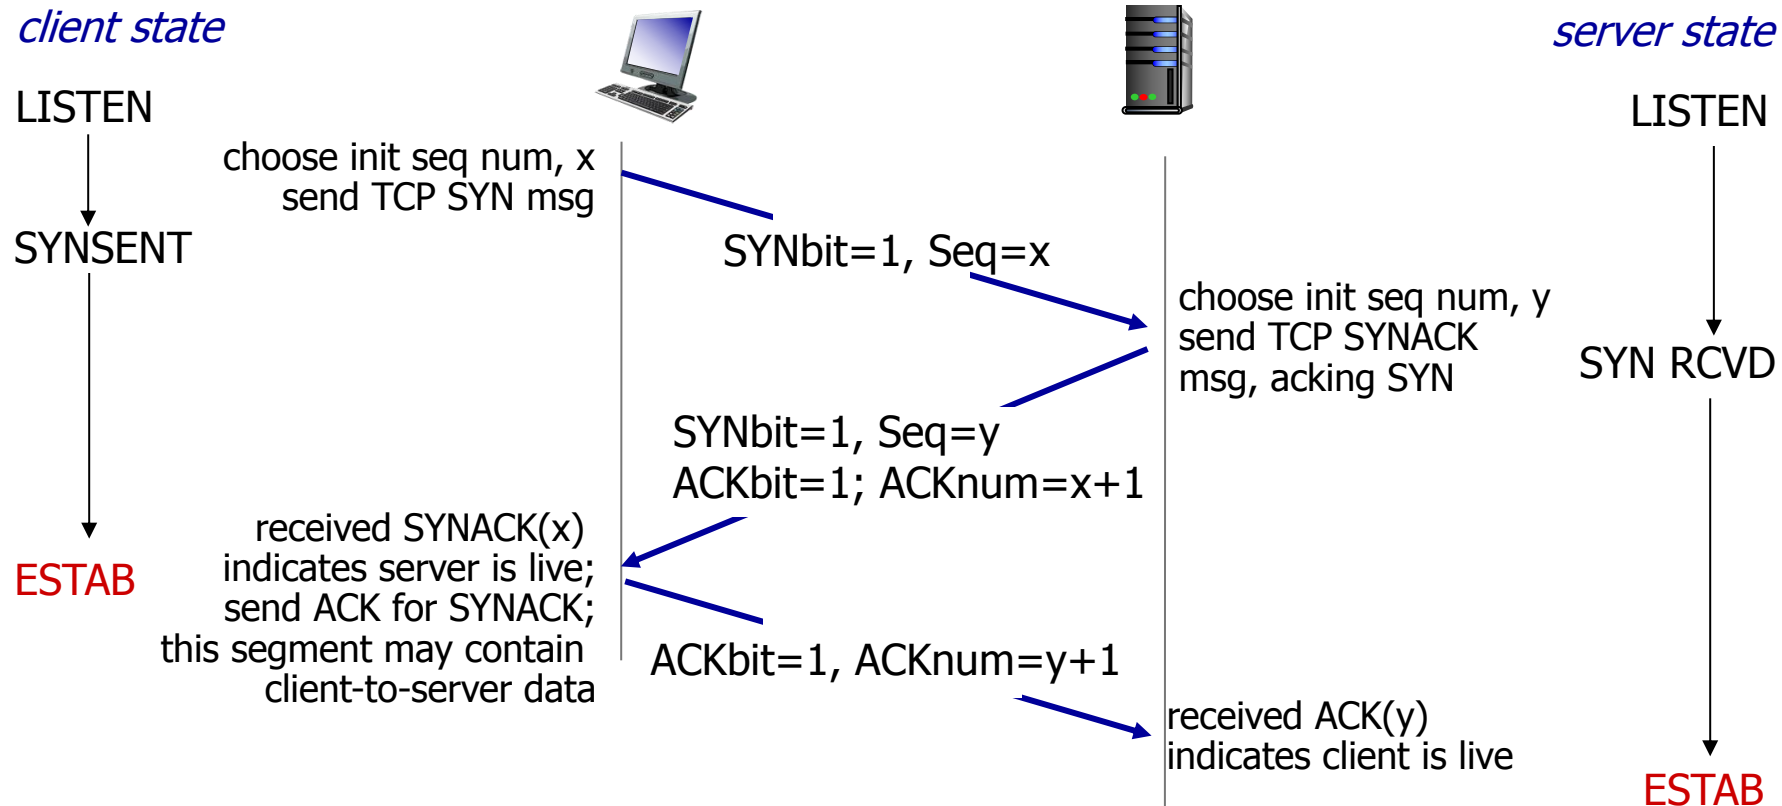
- **Connection establishment**
- **Connection termination**
- **State diagram**

# Connection Establishment



重点

## Three-way handshake

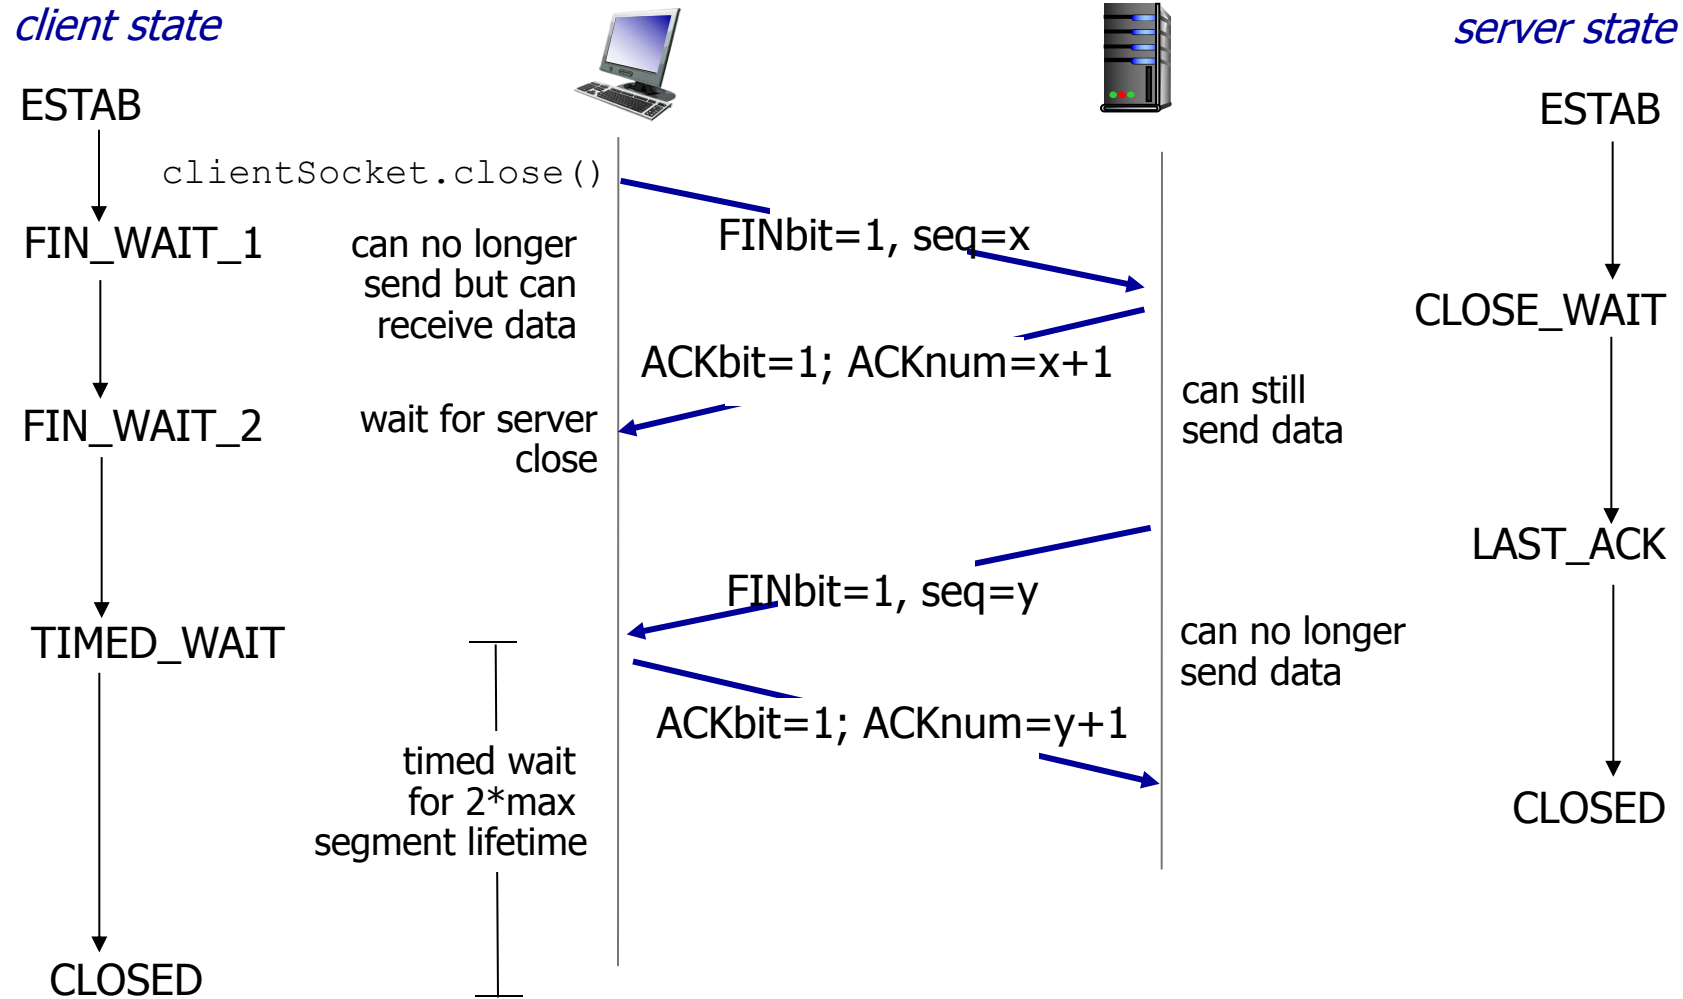


# Connection Termination

## Four-way handshake



重点



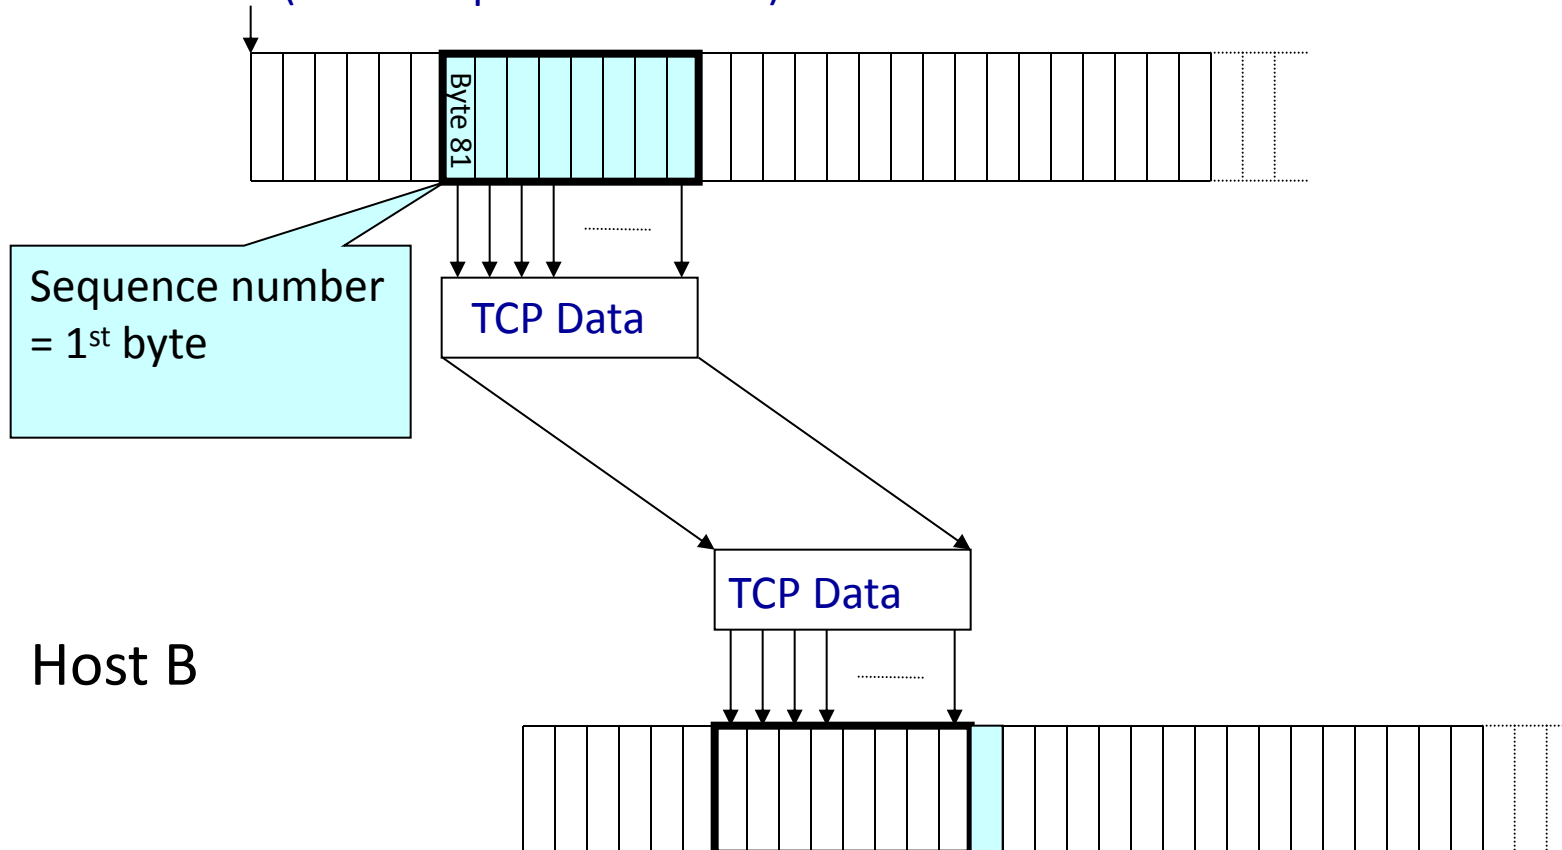


# Initial Sequence Number (ISN)



Host A

ISN (initial sequence number)



Host B

# Initial Sequence Number (ISN)

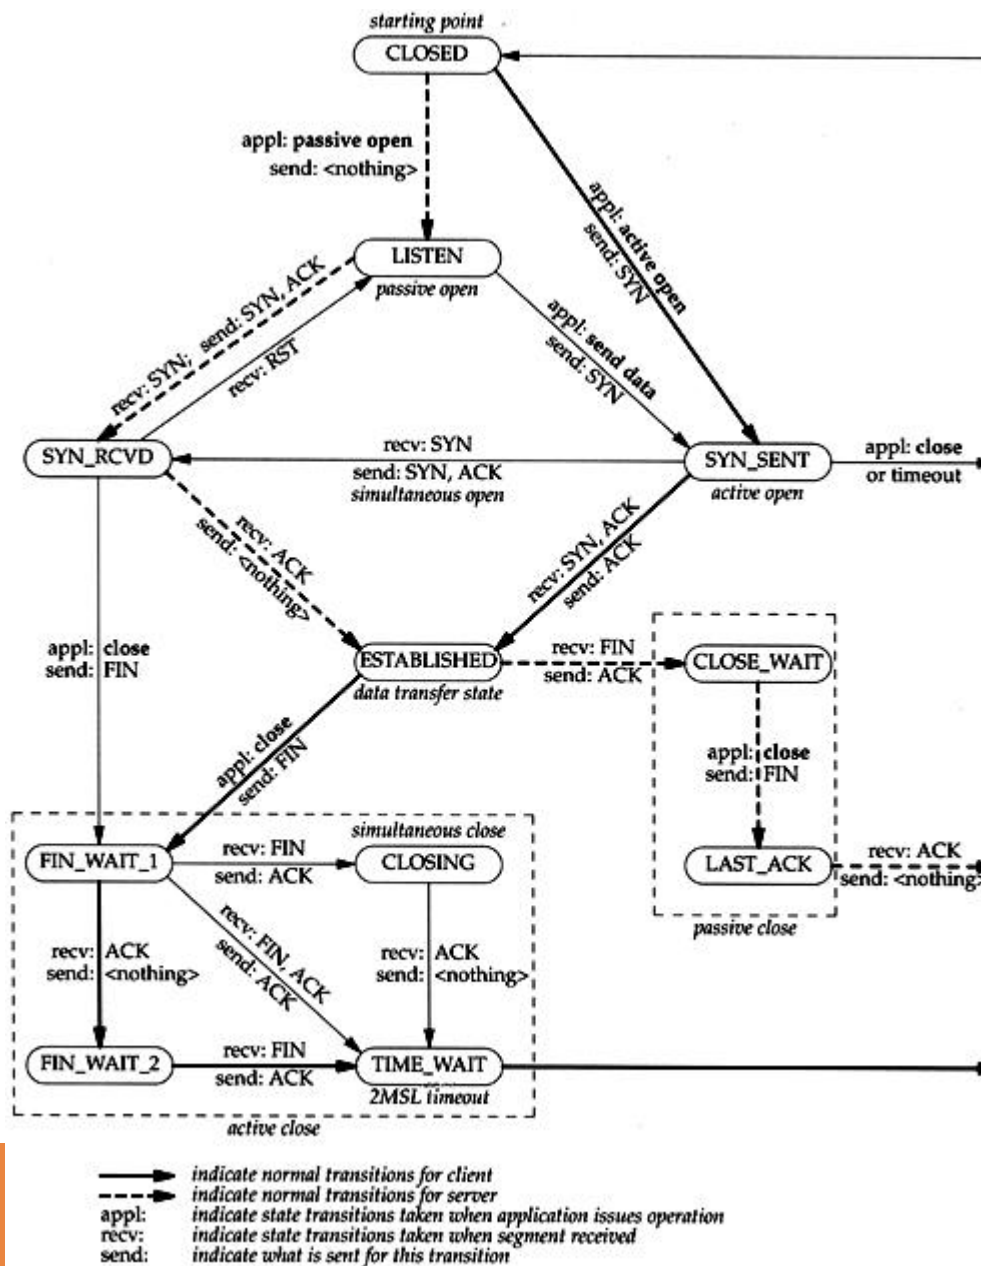


- Sequence number for the very first byte: why not 0 for ISN?
- Practical issue
  - IP addresses and ports uniquely identify a connection, but ports may get reused
  - ... and there is a chance an old packet is still in flight
  - ... and might be associated with the new connection
- So, TCP requires changing the ISN over time
  - Set from a 32-bit clock that ticks every 4 microseconds
  - ... which only wraps around once every 4.55 hours!
- But, this means the hosts need to exchange ISNs

# TCP State Diagram



要理解这个



# Reliable Transmission

重要





南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- **Detect bit errors:** checksum
  - Used to detect corrupted data at the receiver
  - ...leading the receiver to drop the packet
- **Detect missing data:** sequence number
  - Used to detect a gap in the stream of bytes
  - ... and for putting the data back in order
- **Recover from lost data:** retransmission
  - Sender retransmits lost or corrupted data
  - Sender retransmits lost or corrupted data

这三种实现可靠传输的方式要记住

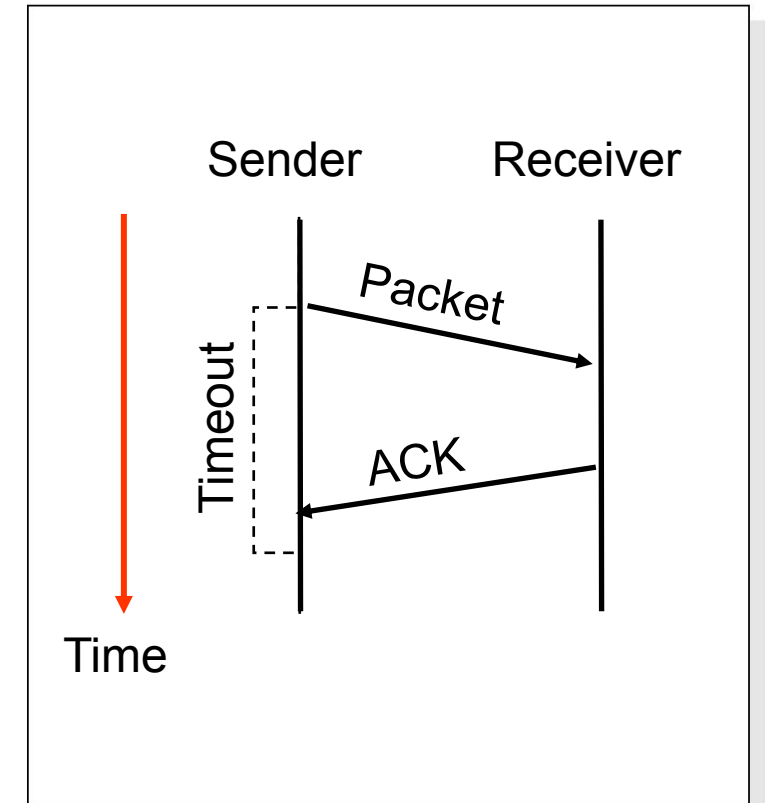
# Sender and Receiver Buffering




# Timeout for Data Loss Detection




- Automatic Repeat reQuest (ARQ) ARQ重要
  - Receiver sends acknowledgment (ACK) when it receives packet
  - Sender waits for ACK and timeouts if it does not arrive within some time period
- Simplest ARQ protocol
  - Stop and wait: send a packet, stop and wait until ACK arrives




# Reasons for Retransmission



Packet lost



ACK lost  
DUPLICATE  
PACKET

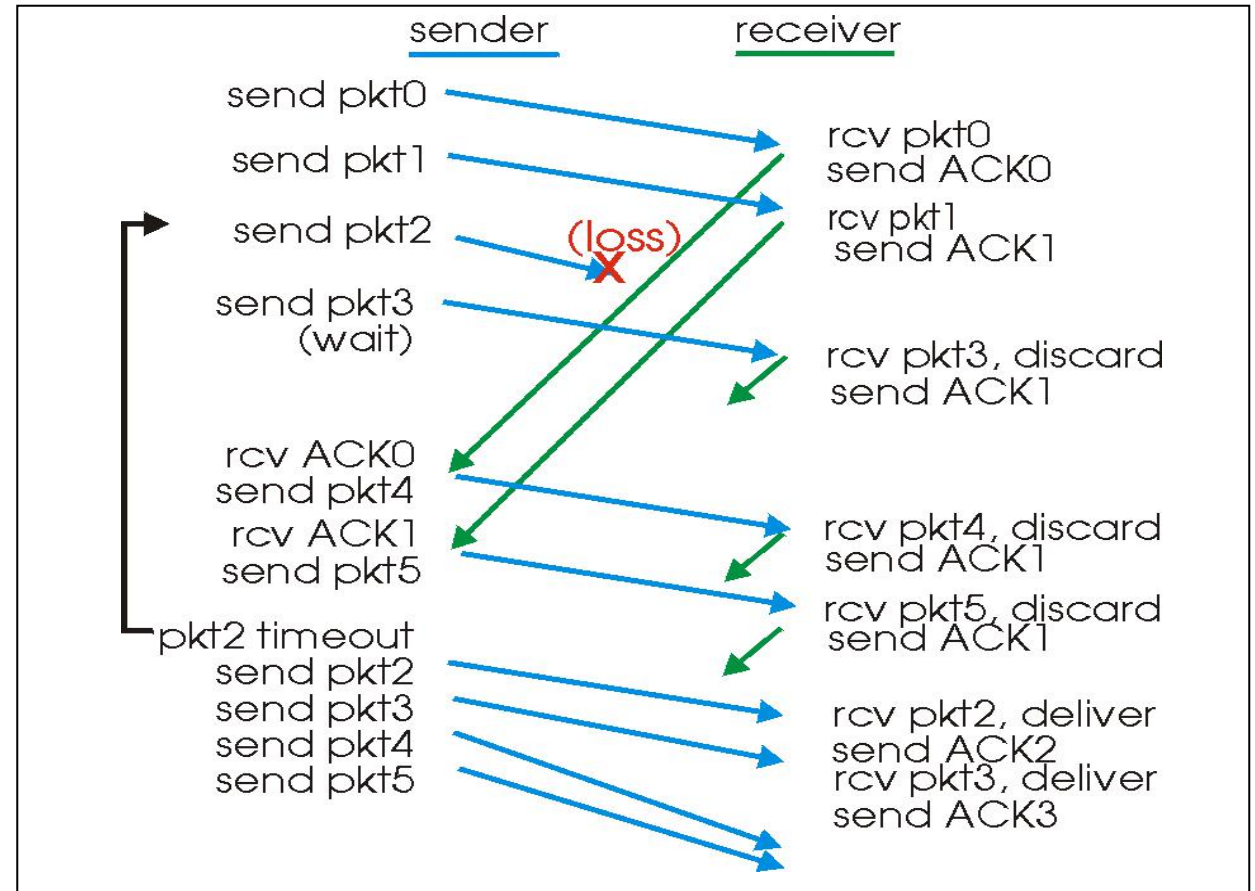


Early timeout  
DUPLICATE  
PACKETS

# Timeout-based Retransmission



- Sender transmits a packet and waits until timer expires
- ... and then retransmits from **the lost packet onward**





# Timeout Setting



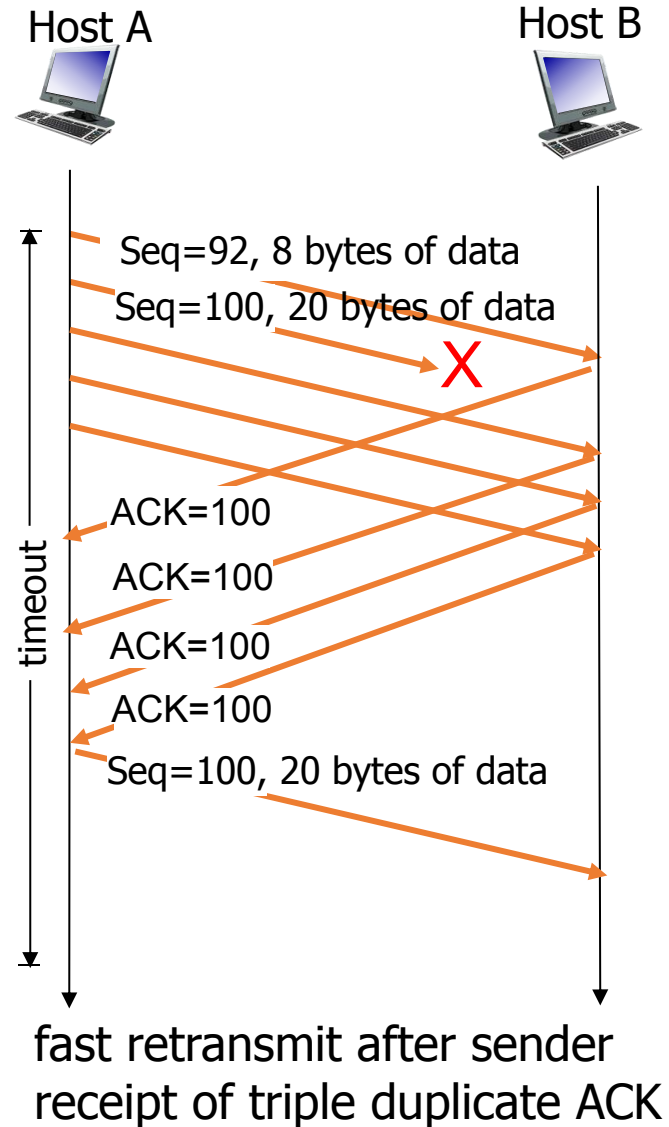
- Sender sets a timeout to wait for an ACK
  - Too short: wasted retransmissions
  - Too long: excessive delays when packet lost
- TCP sets timeout as a function of the Round-Trip-Time (RTT)
  - Expect ACK to arrive after an “round-trip time”
  - ... plus additional time spent on queuing
- But, how does the sender know the RTT?
  - Can estimate the RTT by watching the ACKs
  - Smooth estimate: keep a running average of the RTT  
$$\text{EstimatedRTT} = a * \text{EstimatedRTT} + (1 - a) * \text{SampleRTT}$$
  - Compute timeout:  $\text{TimeOut} = 2 * \text{EstimatedRTT}$

# Fast Retransmission



- ARQ/Timeout is slow and inefficient
- Better solution possible under sliding window
  - Although packet  $n$  might have been lost
  - ... packets  $n+1$ ,  $n+2$ , and so on might get through
- Idea: have the receiver send ACK packets
  - ACK says that receiver is still awaiting  $n^{\text{th}}$  packet
  - And *repeated* ACKs suggest later packets have arrived
  - Sender can view the “duplicate ACKs” as an early hint for lost of the  $n^{\text{th}}$  packet
- Fast retransmission
  - Sender retransmits data after the triple duplicate ACK

# Tripple Duplicate ACK



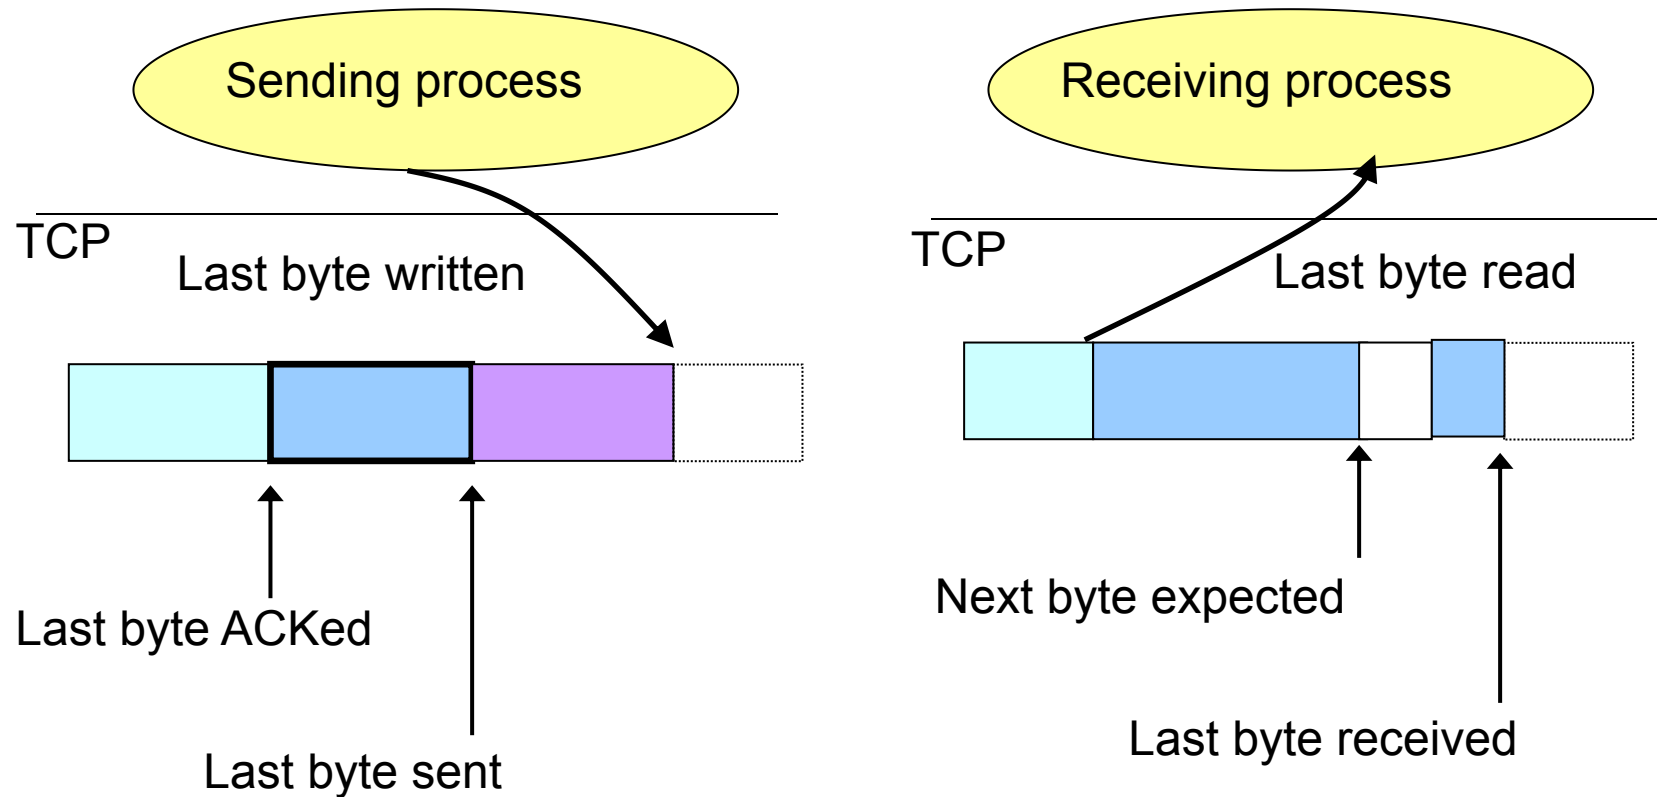
# Contents

- Basics of Transport Layer
- Reliable Transportation
- Flow Control
- Congestion Control
- Multipath TCP

# Flow Control



- The receiver controls the sender, so sender won't overflow receiver's buffer by transmitting too much or too fast




# Sliding Window

- Stop-and-wait is inefficient
  - Only one TCP segment is “in flight” at a time
  - Especially bad when delay-bandwidth product is high
- Sliding window
  - Allow a larger amount of data “in flight”
  - Allow sender to get ahead of the receiver

# Receiver Buffering



- Window size
  - Amount that can be sent without acknowledgment
  - Receiver needs to be able to store this amount of data
- Receiver advertises the window to the sender




# TCP Header Field for Receiver Buffering



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Source port		Destination port	
Sequence number			
Acknowledgment			
HdrLen	0	Flags	Advertised window
Checksum		Urgent pointer	
Options (variable)			
Data			

## Window-Size Example






# Contents

- Basics of Transport Layer
- Reliable Transportation
- Flow Control
- Congestion Control
- Multipath TCP

# Problem of Congestion



- Problem
  - Too many sources sending too much data too fast for *network* to handle
- Manifestations
  - Lost packets (buffer overflow at routers)
  - Long delays (queueing in router buffers)



# How to Learn Congestion?

- **Delay**

重要

- Round-trip time (RTT) estimate


- **Loss**

- Timeout
- Duplicate acknowledgments

- **Mark**

- Packets marked by routers with large queues

# Drop-tail Queuing



# Congestion Collapse




- **Rough definition:** “When an increase in network load produces a decrease in useful work”
- **Why does it happen?**
  - Sender sends faster than bottleneck link speed
  - Packets queue until dropped
  - In response to packets being dropped, sender retransmits
  - Retransmissions further congest link
  - All hosts repeat in steady state...

# Proactive vs Reactive



- **Congestion avoidance:** try to stay to the left of the “knee”
- **Congestion control:** try to stay to the left of the “cliff”



重要

# Congestion Control



- **Flow control**

- To prevent resource exhaustion **at end node**
- Basic idea: receiver advertises sliding window **awnd** with each ACK

- **Congestion control**

- To prevent resource exhaustion **within network**
- Basic idea: source calculates congestion window **cwnd** from indication of network congestion (losses, delay, mark)

- **Sender TCP window** =  $\min \{ \text{awnd}, \text{cwnd} \}$

# Congestion Control Algorithms




- Tahoe (Jacobson 1988)
  - Slow Start
  - Congestion Avoidance
  - Fast Retransmit
- Reno (Jacobson 1990)
  - Fast Recovery
- Vegas (Brakmo & Peterson 1994)
  - New Congestion Avoidance
- RED, REM...



# 1. Slow Start

- Goal: quickly find the equilibrium sending rate
- When connection begins, increase rate exponentially
  - initially **cwnd** = 1 MSS
  - double **cwnd** every RTT
- When to stop exponential increase?
  - First packet loss detected
  - Threshold encountered (ssthresh)




## 2. Congestion Avoidance



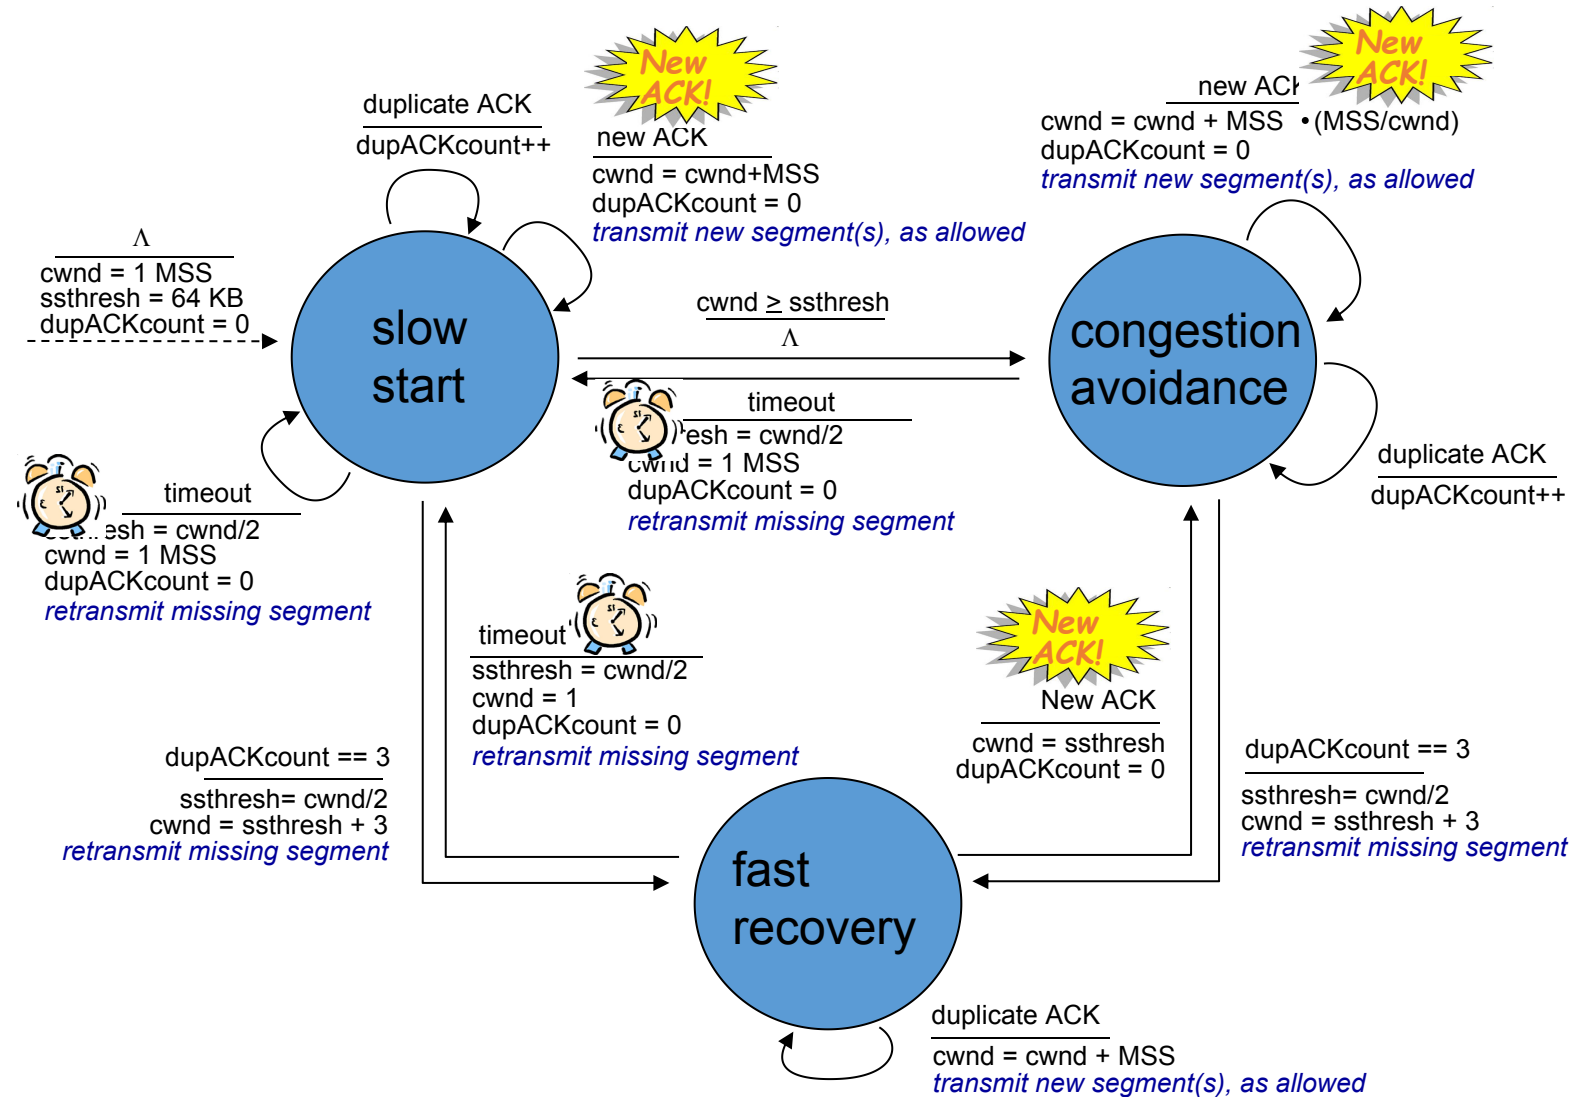
- Goal: detecting and reacting to loss
- Loss indicated by timeout
  - **cwnd** set to 1 MSS
  - Window grows exponentially (as in slow start) to threshold, then grows linearly
- Loss indicated by 3 duplicate ACKs: TCP RENO
  - Dup ACKs indicate network capable of delivering some segments
  - **cwnd** is cut in half window then grows linearly
- TCP Tahoe always sets **cwnd** to 1 (timeout or 3 duplicate acks)

# Evolution of Cwnd




on loss event, **sssthresh** is set to  $1/2$  of **cwnd** just before loss event


# State Transition Diagram



# TCP Sawtooth Behavior



- AIMD: Additive Increase, Multiplicative Decrease
  - Mechanism working in congestion avoidance and fast transmission



# Discussion

- TCP flows can be
  - Elephant flows
  - Mice flows
- Which type of flows are dominant
  - In terms of number of flows?
  - In terms of bandwidth consumption?
- How does Tahoe/Reno work on mice flows?

# Contents

- Basics of Transport Layer
- Reliable Transportation
- Flow Control
- Congestion Control
- Multipath TCP



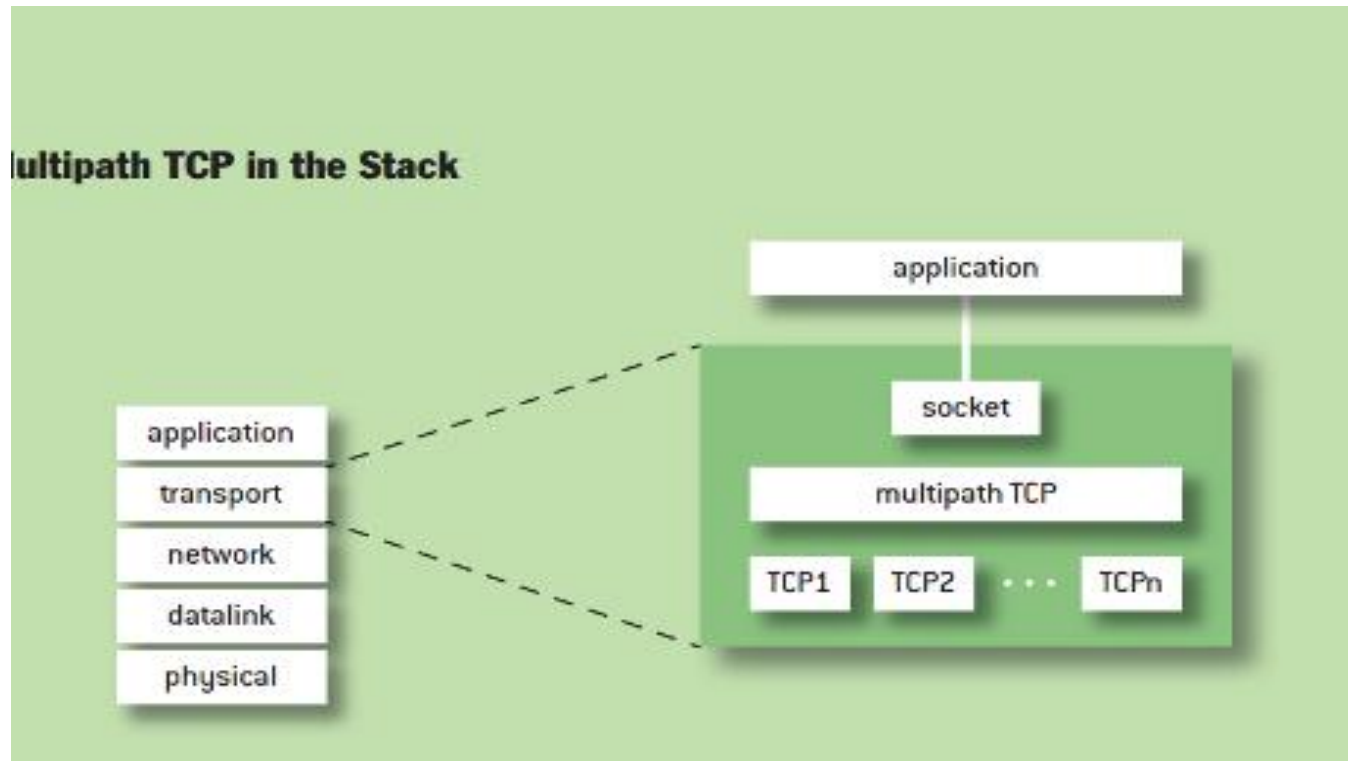
# Multipath TCP

- Regular TCP
  - Single IP connection
- Modern infrastructures
  - **High-end server**: multiple Ethernet cards
  - **Data centers**: rich topologies with many paths
  - **Mobile users**: WiFi and cellular at the same time
- Benefits of multipath
  - Higher throughput
  - Failover from one path to another
  - Seamless mobility

# Working with Unmodified Apps



- Present the same socket API and expectations (IP, port, protocol)




# Working with Unmodified Apps



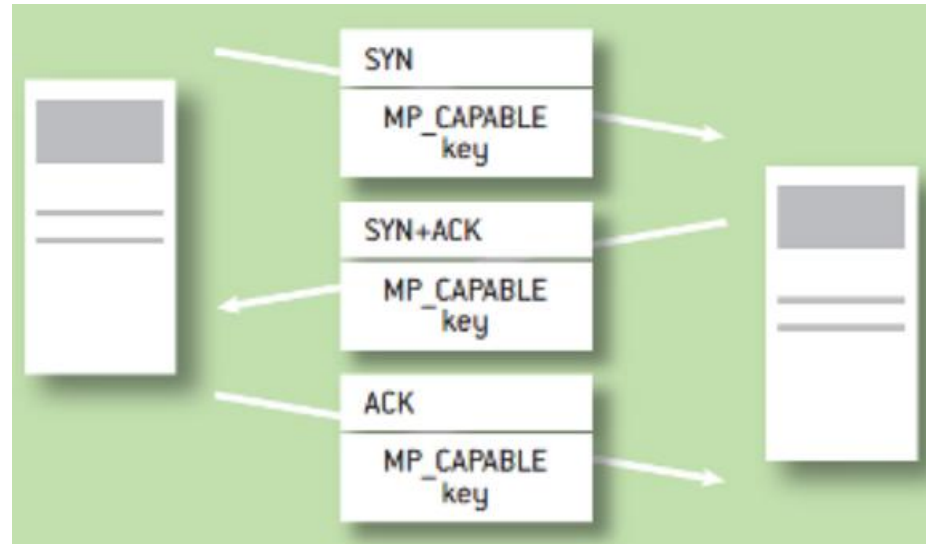
- Establish the TCP connection in the normal way
  - Create a socket to a single remote IP address/port
  - And then add more subflows, if possible

Each host tells its *Initial Sequence Number (ISN)* to the other host.



# Negotiating MPTCP Capability

- How do hosts know they both speak MPTCP?
- During the 3-way SYN/SYN-ACK/ACK handshake
  - If SYN-ACK doesn't contain MP\_CAPABLE, don't try to add any subflows!



# Adding Subflows, Idealized



- How to associate a new subflow with the connection?
  - Use a token generated from original subflow set-up
- How to start using the new subflow?
  - Simply start sending packets with new IP/port pairs
  - ... and associate them with the existing connection
- How could two end-points learn about extra IP addresses for establishing new subflows?
  - Implicitly: one end-point establishes a new subflow, to already-known address(es) at the other end-point

# Sequence Numbers



- Challenges across subflows
  - More out-of-order packets due to RTT differences
  - Access networks that rewrite sequence numbers
  - Middleboxes upset by discontinuous TCP byte stream
  - Need to retransmit lost packets on a different subflow
- Solution: two levels of sequence numbers
  - Sequence numbers per subflow
  - Sequence numbers for the entire connection
- Enables
  - Efficient detection of loss on each subflow
  - Retransmission of lost packet on a different subflow

# Read More on MPTCP



- Improving Datacenter Performance and Robustness with Multipath TCP  
[\[SIGCOMM 2011\]](#)
- Design, Implementation and Evaluation of Congestion Control for Multipath TCP [\[NSDI 2011\]](#)

# Summary

- What is transportation layer for?
- Major issues
  - Reliable transmission
  - Flow control
  - Congestion control
- New infrastructures and applications drive new extensions