

# Software Defined Networking (1/2)



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

**Xuetao Wei**

weixt@sustech.edu.cn

# Outline

- **Motivation for SDN**
- SDN Overview
- OpenFlow
- Case Study

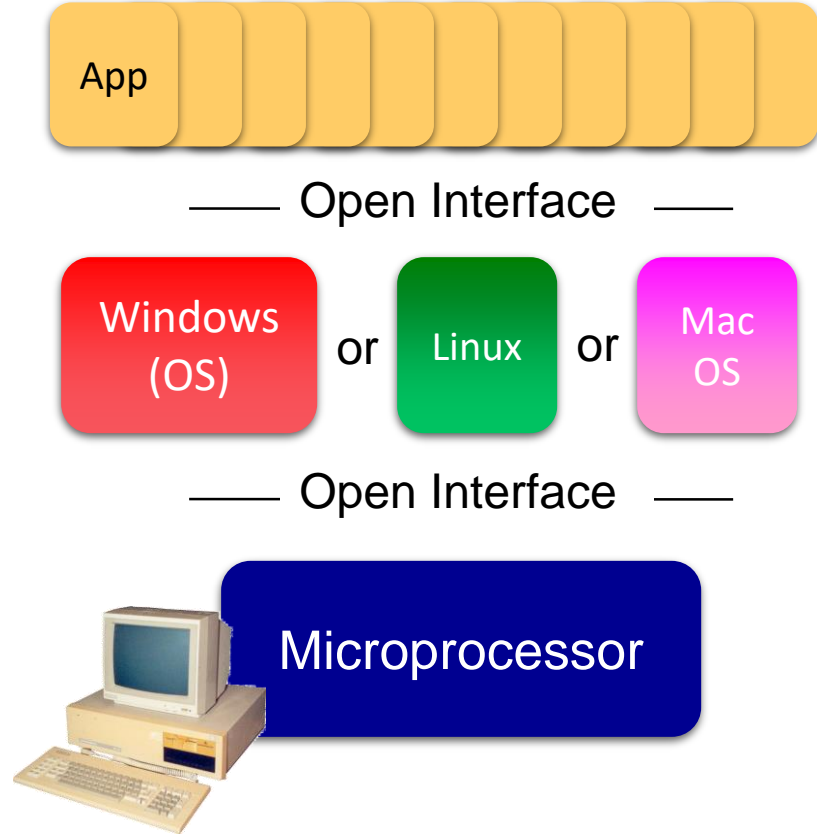


# Computer Industry



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



**Traditional Computer Industry:**  
Vertically integrated, Closed  
Slow innovation, Small Industry



**PC era Industry:**  
Horizontal, Open interfaces  
Rapid innovation, Huge industry

# Networking Industry




南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



Vertically integrated, with “mainframe” mind-set

Routing, management, mobility management, access control, VPNs, ...



Million of lines of source code      6,000 RFCs

Billions of gates      Bloated      Power Hungry

# Software Defined Network

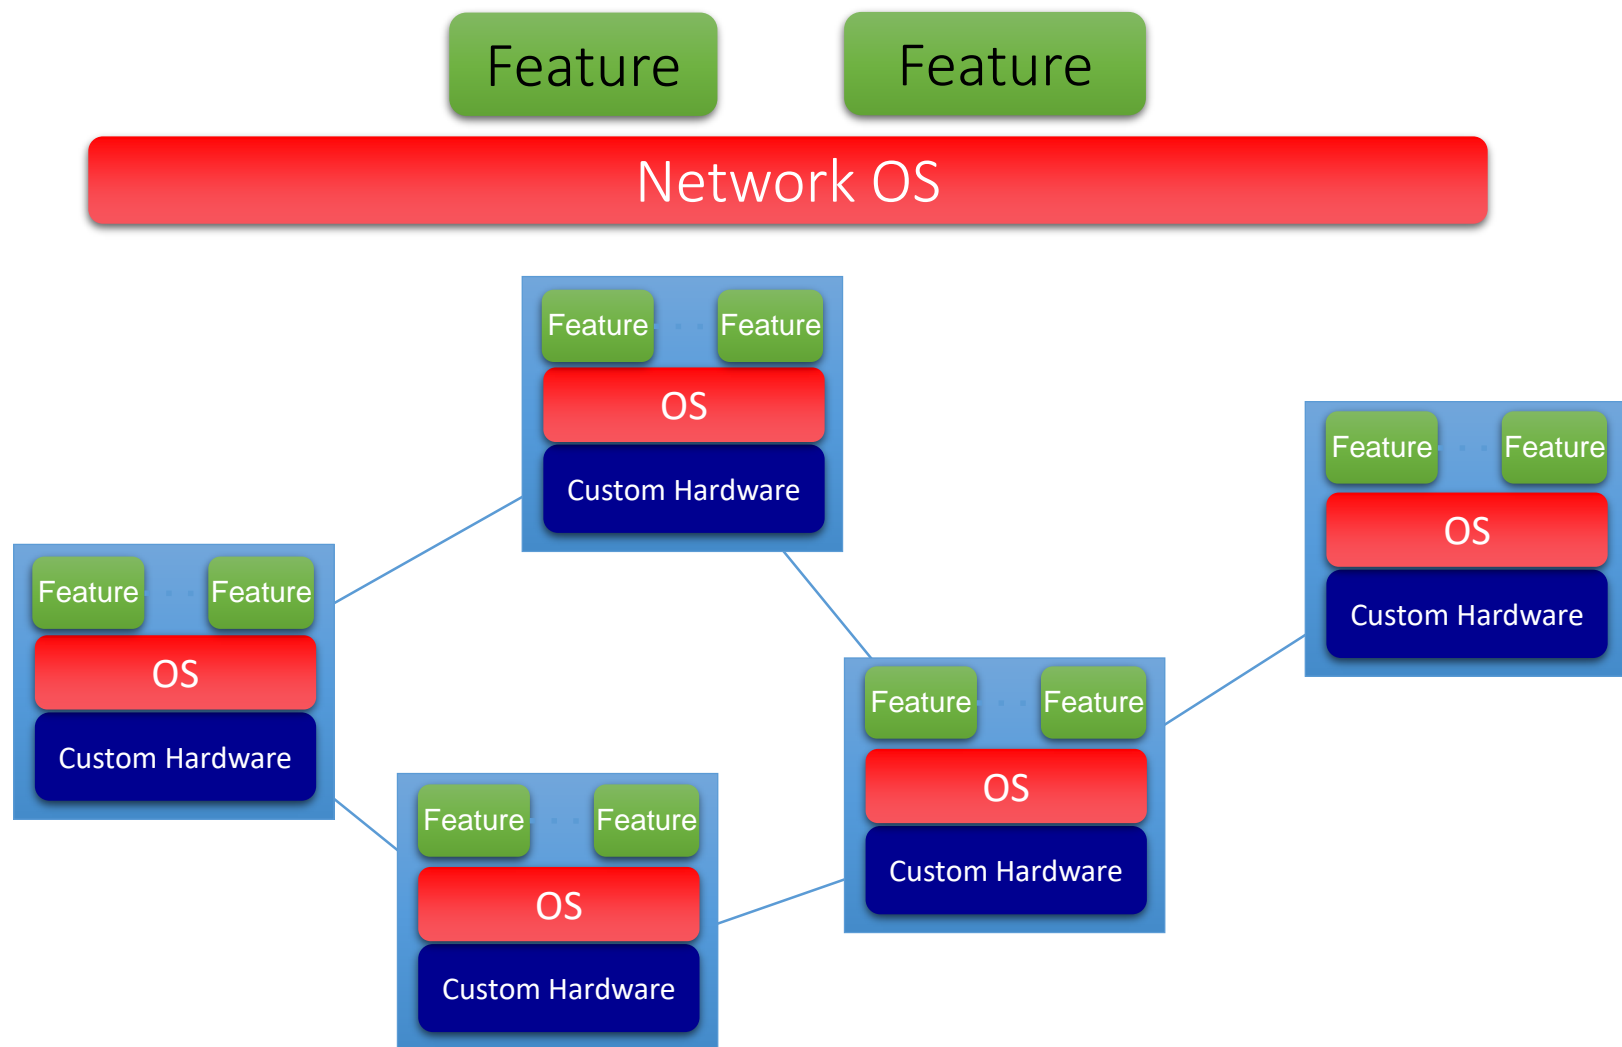


南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

A network in which the **control plane** is physically separate from the **forwarding plane**.

*and*

A single **control plane** controls several **forwarding devices**.




# Network Industry



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



Vertically integrated  
Closed, Slow  
innovation




— Open Interface —



— Open Interface —



Horizontal  
Open interfaces, Rapid innovation



# Immediate Consequences

- Networks cost less: hardware is simple and streamlined again
- Networks cost less to operate: customized
- Networks are more reliable and more secure
- Puts network owners and operators in control of their destiny

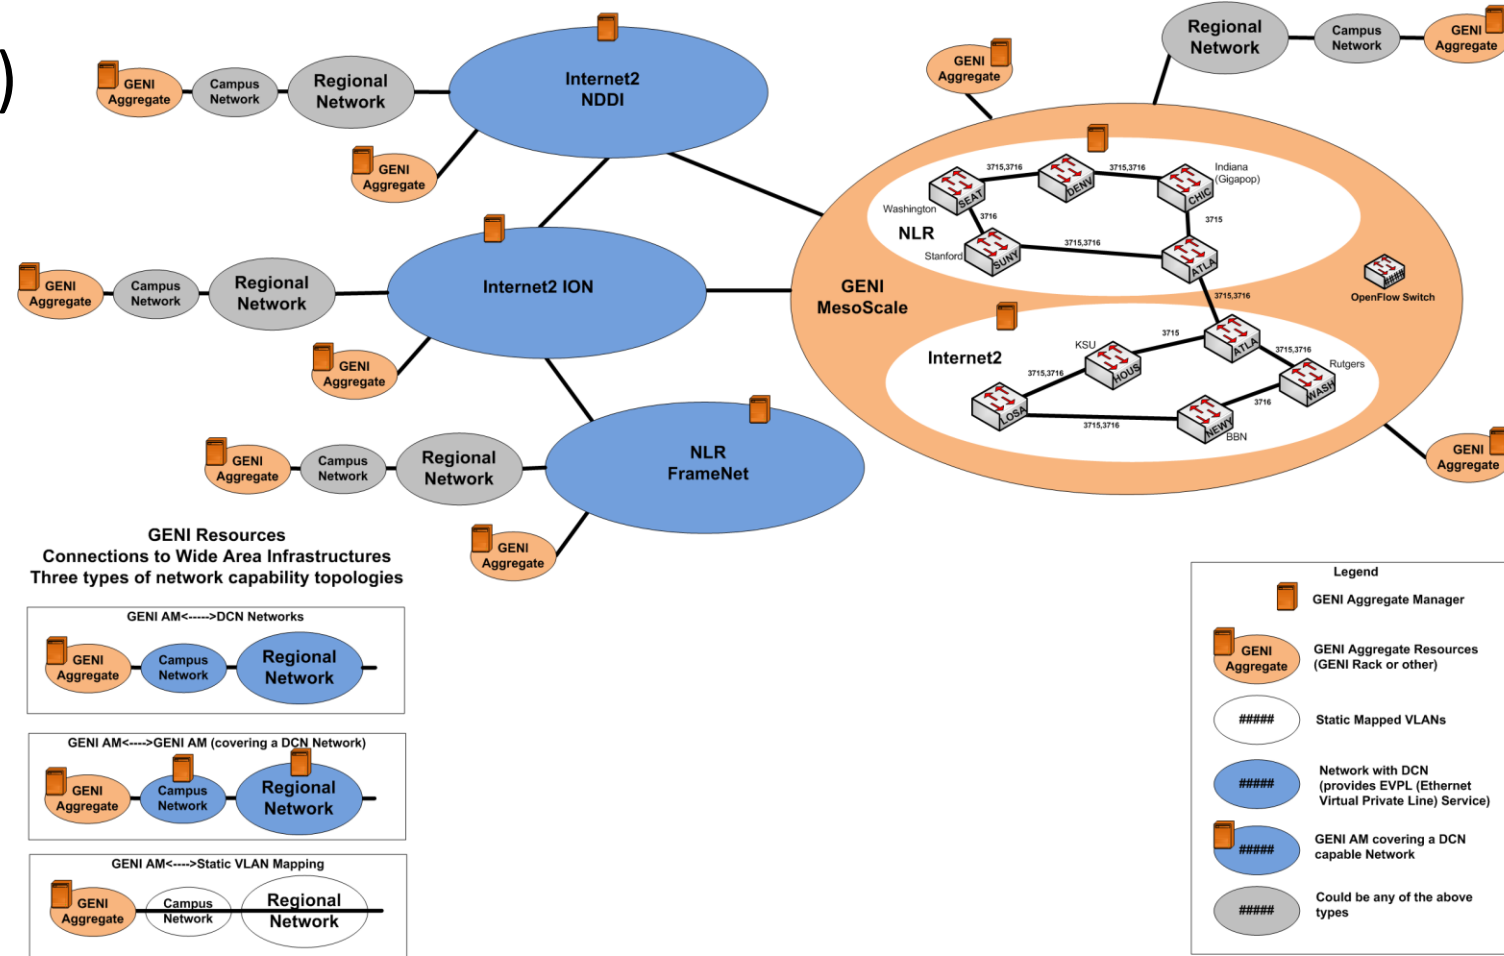


# History of SDN (1/3)



- **Background: NGN (Next Generation Network Initiative)**

- Internet2
- GENI (Global Environment for Network Innovations)



# History of SDN (2/3)



## 1. **Clean Slate Project** (Nick McKeown @Standord), 2005~2007

- Current network is bloated and convoluted by many things
- Martin Casado: “Ethane: Taking Control of the Enterprise”, [SIGCOMM 2007]
- Primary idea: central controller over multiple switches for a network

## 2. **Nicira startup** (Casado, McKeown, Scott Shenker@Berkely), 2007

- NOX controller
- OpenFlow interface  
(OpenFlow: Enabling Innovation in Campus Networks, [SIGCOMM 2008])
- Acquired by VMware in 2012 for **1.26 billion** dollars!

# History of SDN (3/3)



3. **OpenFlow 1.0** (Casado, Appenzeller@BigSwitch), 2009
4. ONF (Open Networking Foundation) founded, 2011
  - Deutsche Telecom, Facebook, Google, Microsoft, Verizon, Yahoo!
  - OpenFlow 1.1, 1.2, 1.3, 1.4
5. Many other organizations are founded, and many other SDN initiatives are launched since 2011...
6. Protocol-independent packet forwarding architecture is now being promoted by both academia (Stanford, Princeton, ...) and industry (Huawei)

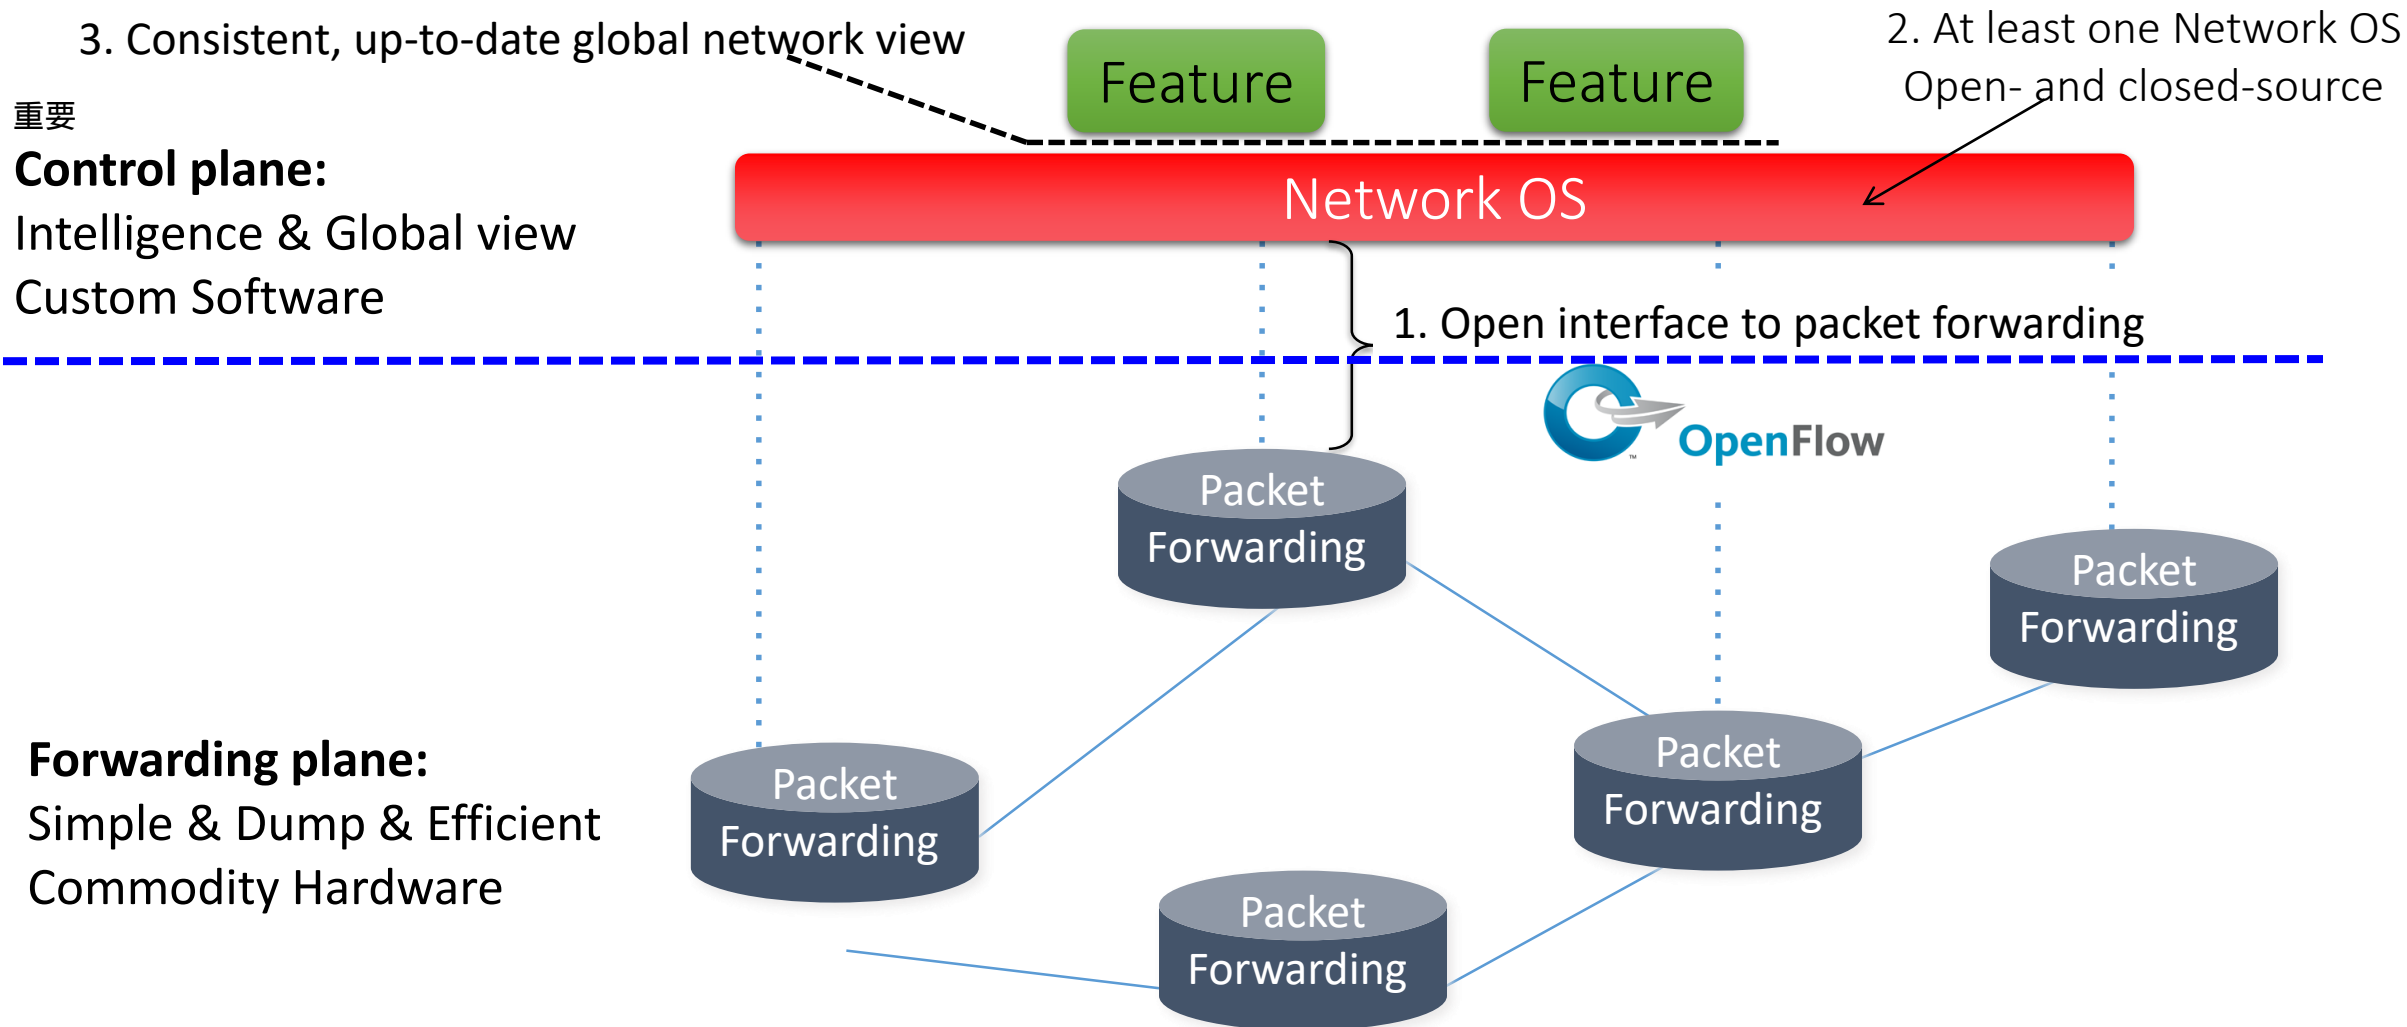
# Outline

- Motivation for SDN
- **SDN Overview**
- OpenFlow
- Case Study

# SDN Overview



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

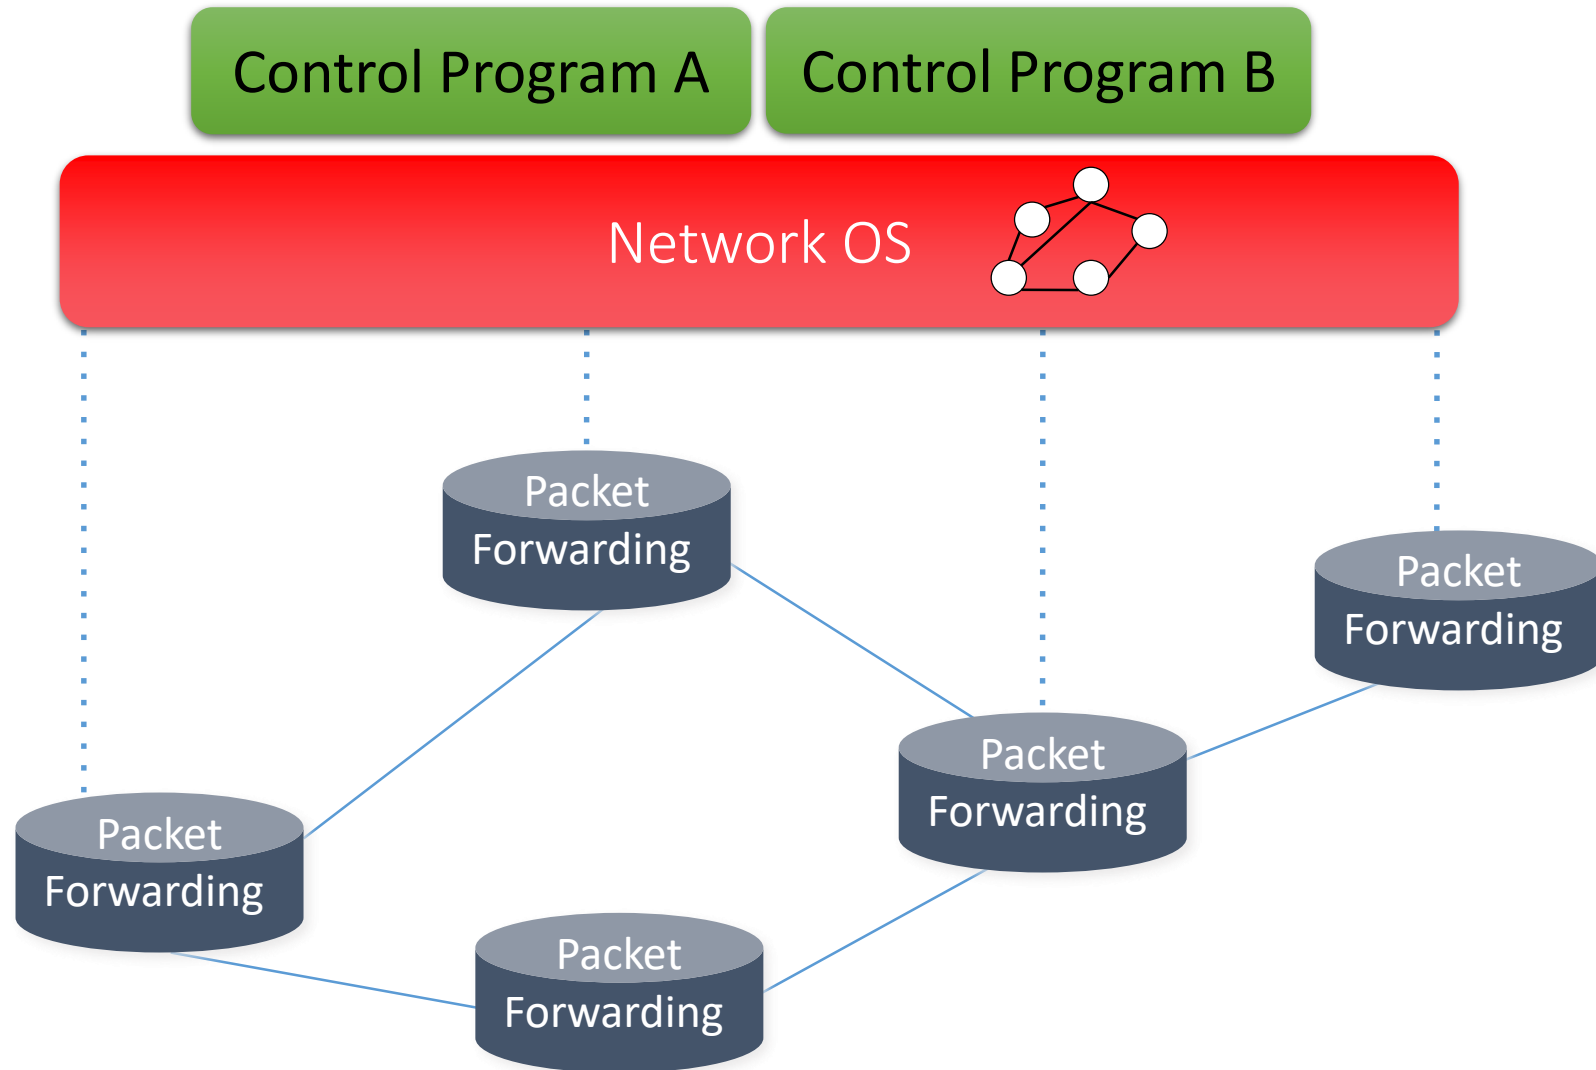


# Network OS



- A distributed system that creates a consistent, up-to-date network view
  - Runs on servers (controllers) in the network
  - Examples: Floodlight, POX, Pyretic, Nettle ONIX, Beacon, ...
- Uses forwarding abstraction to
  - Get state information **from** forwarding elements
  - Give control directives **to** forwarding elements

# Control Program

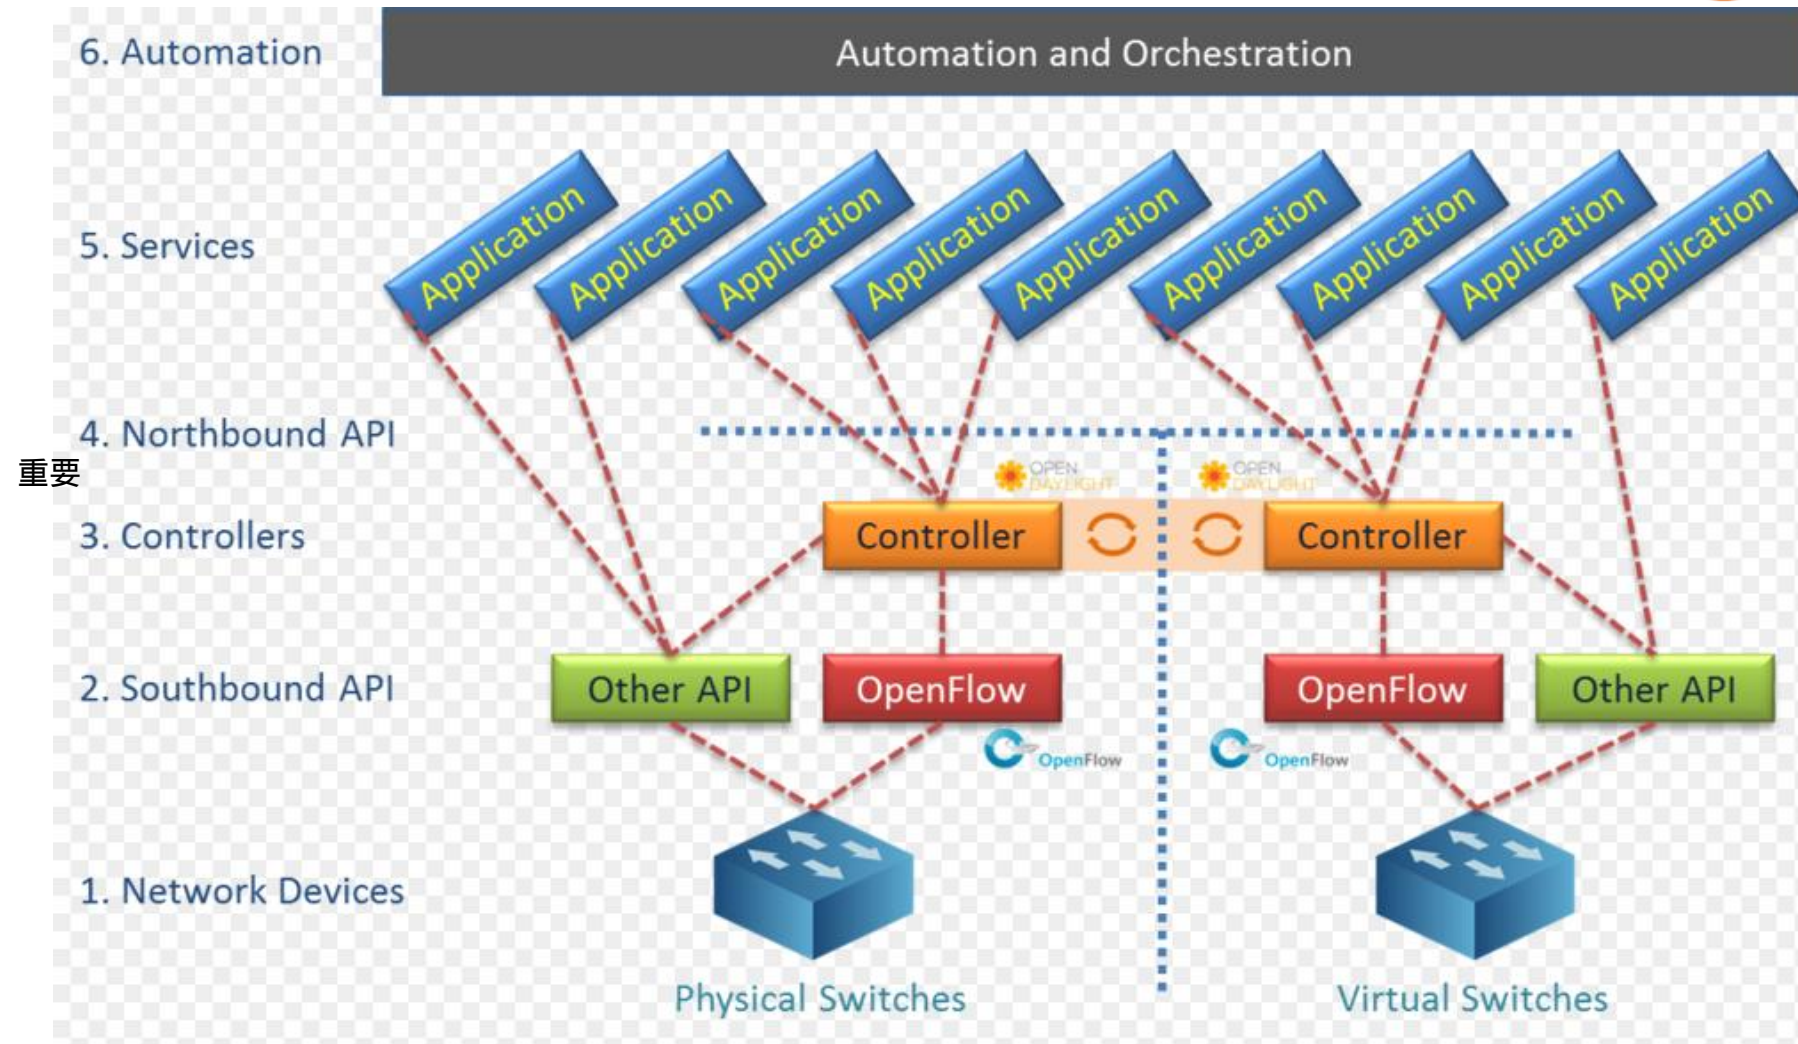


# Control Program

- Control program implements network features
- It operates on view of network
  - **Input:** global network view (graph/database)
  - **Output:** configuration of each network device
- Control program is not a distributed system
  - Abstraction hides details of distributed state



# A Complete SDN Hierarchy



# 1. Network Devices

- Abstracted as the forwarding plane or data plane, which forward packets based on flow tables
- Can be either hardware switches or software switches (OVS), can also be other devices like routers
- They receive instructions from the controller via the **Southbound Interface**, which will configure the flow table. They also report states and events to the controller through the **Southbound Interface**

## 2. Southbound Interface

重要

- The interface between the control plane and data plane
- Traditional Southbound interfaces are proprietary, and the differences of Southbound interfaces among different vendors prevent managing heterogenous network devices as a whole within or across a network
- Standardization of the Southbound interface is urgently needed, and OpenFlow is a big step forward

# 3. Controller

- SDN controller configures the switches with flow tables, which specifies the rules for packet forwarding
- An SDN controller can be any x86 Linux servers or Windows servers, or other computers running networking software
- An SDN controller may control multiple switches (and centralized control is promoted by SDN community), it is also possible that a switch is controlled by multiple controllers (one master, multiple slave)
- Controller decides the ecosystem, and is the kernel element in SDN network, therefore companies are competing in this area fiercely

# 4. Northbound Interface

重要

- Northbound Interface in traditional network refers to the interface between the switch control plane and the network management software (SNMP, TL1, ...)
- In SDN, Northbound Interface is the interface between the controller and the Apps/Services
- So far Northbound Interface is not standardized, and is much more complex to be standardized than the Southbound Interface

# 5. Apps/Services



重要

- With SDN, traditional network services can be better implemented, and new services can be practiced more easily
- **Load balancing:** global view and orchestration of the network simplifies load balancing
- **Security:** can be enforced with close collaboration between different switches coordinated by the centralized control plane
- **Monitoring and Performance Management:** can be implemented at the network level instead of at the switch level
- **LLDP:** topology discovery can be realized by the centralized controller easily and efficiently

# Outline

- Motivation for SDN
- SDN Overview
- **OpenFlow**
- Case Study

# Forwarding Abstraction

- Purpose: Abstract away forwarding hardware implementation details
- Flexible
  - Behavior specified by control plane
  - Built from basic set of forwarding primitives
- Minimal
  - Streamlined for speed and low-power
  - Control program not vendor-specific

**OpenFlow** is an example of such an abstraction




# OpenFlow



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- OpenFlow: An API between the control plane and dataplane

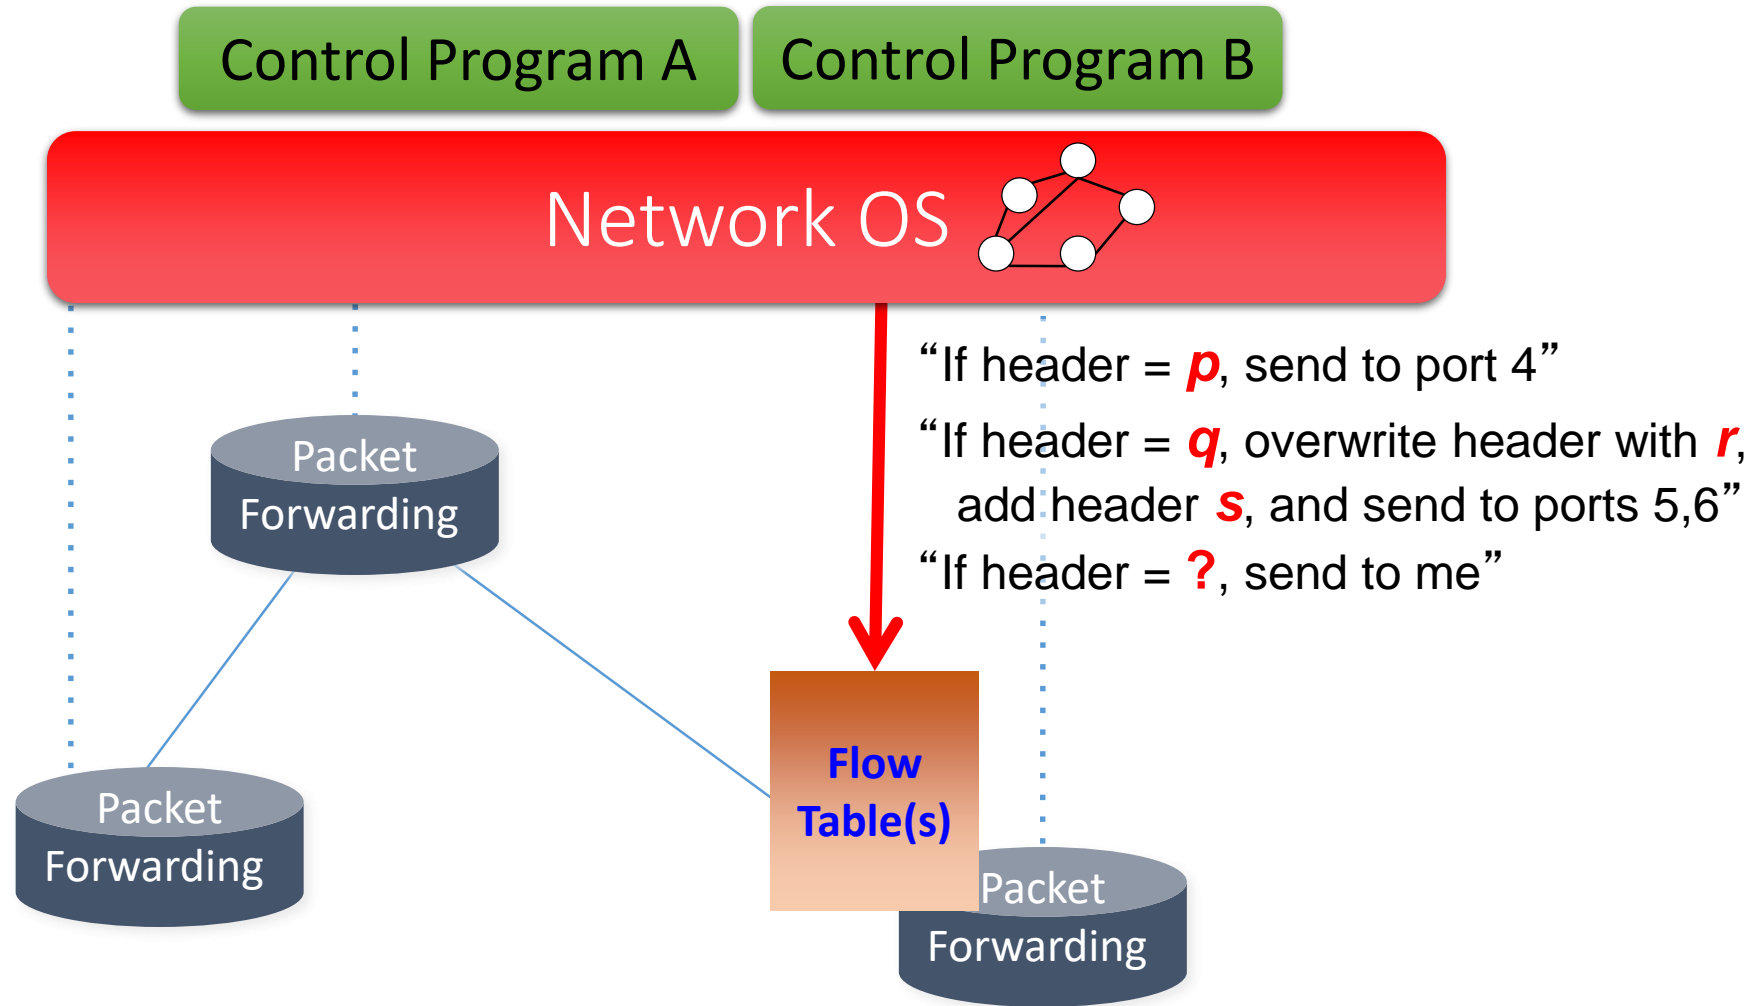


Like the **ISA** between  
SW & HW in a computer

# OpenFlow as An Interface

- OpenFlow is like an x86 instruction set for the network
- Provides open interface to “black box” networking node (ie. Routers, L2/L3 switch) to enable visibility and openness in network
- Separation of control plane and data plane by OpenFlow
  - The datapath of an OpenFlow Switch consists of a Flow Table, and an action associated with each flow entry
  - The control path consists of a controller which programs the flow entry in the flow table

# Flow Tables



# <Match, Action> Primitives



- **Match** arbitrary bits in headers:



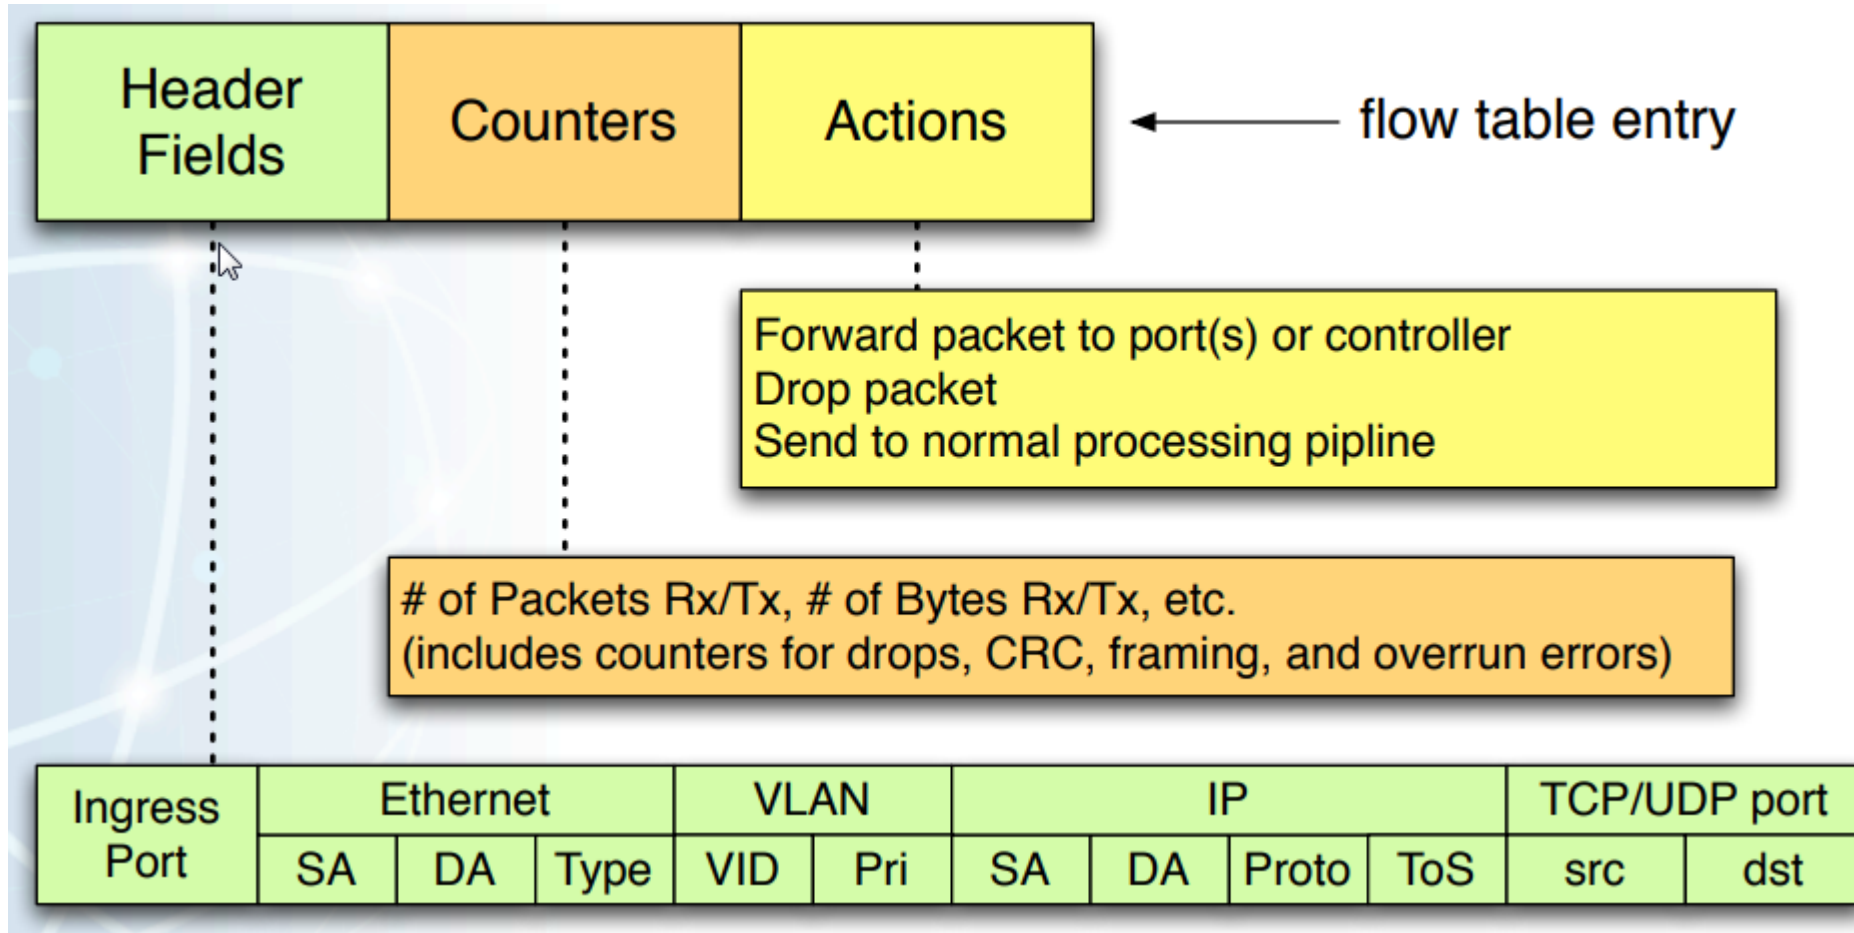
Match: 1000x01xx0101001x

- Match on any header, or new header
  - Allows any flow granularity
- 
- **Action**
    - Forward to port(s), drop, send to controller
    - Overwrite header with mask, push or pop
    - Forward at specific bit-rate

# Flow Table Entries



重要



# Examples



## Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f:...	*	*	*	*	*	*	*	port6

## Flow Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
port3	00:20..	00:1f..	0800	vlan1	1.2.3.4	5.6.7.8	4	17264	80	port6

## Firewall

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	*	*	*	22	drop

# Examples

## Routing

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	5.6.7.8	*	*	*	port6


## VLAN Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f..	*	vlan1	*	*	*	*	*	port6, port7, port9


# Centralized vs Distributed Control



## Centralized Control



## Distributed Control





# Flow Routing vs. Aggregation



## Flow-Based

- Every flow is individually set up by controller
- Exact-match flow entries
- Flow table contains one entry per flow
- Good for fine grain control, e.g. campus networks

## Aggregated

- One flow entry covers large groups of flows
- Wildcard flow entries
- Flow table contains one entry per category of flows
- Good for large number of flows, e.g. backbone

# Reactive vs. Proactive (pre-populated)



## Reactive

- First packet of flow triggers controller to insert flow entries
- Efficient use of flow table
- Every flow incurs small additional flow setup time
- If control connection lost, switch has limited utility

## Proactive

- Controller pre-populates flow table in switch
- Zero additional flow setup time
- Loss of control connection does not disrupt traffic
- Essentially requires aggregated (wildcard) rules

# Controller-Switch Interactions



quiz不考

- **Controller-to-Switch Messages**

- **Features:** query for functionalities supported by the switch
- **Configuration:** configure the switch or query its configuration parameters
- **Modify-State:** modify the flow table (add/delete/modify)
- **Read-State:** access various state information of the switch (e.g., counter)
- **Packet-out:** send out a packet matching a given flow entry
- **Barrier:** enforce order between messages
- **Role-request:** tell the switch its role when multiple controllers are connected to it
- ...

# Controller-Switch Interactions



- **Asynchronous Messages (Switch-to-Controller)**

- **Packet-in:** a message to controller when a packet matches a flow entry with action “output to Controller-Port”
- **Flow-removed:** when a flow entry is removed from the flow table (due to aging or instructed by the controller)
- **Port-status:** a message to the controller when a port status is changed (link down/up)
- **Error:** a message to the controller when the switch encounters errors
-

# Controller-Switch Interactions

- **Symmetric Messages**

- **Hello:** a message to inform the others when a controller or switch is booted up
- **Echo:** an echo/reply handshake to make sure the connection is OK, or to measure the delay of the connection
- **Experimenter:** a message for vendor-specific extensions
-

# Controller-Switch Workflow



## 1. Initialization

- Switch: a default flow table with a single default flow entry that either discards all packets or forwarding them to the controller for processing

## 2. Business-driven flow table update by controller

- According to business requirements, the controller adds new flow entries onto the flow table in the switch, updates existing entries, and deletes obsolete entries
- Flow table update can be carried out by network operators, or automated by network Apps/Services. The switch itself can also retire aging flow entries

## 3. Packet processing in switch

- An incoming packet searches the flow table for a match, then execute the corresponding actions

# SDN Standardization on OpenFlow



南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY


- Open Network Foundation (ONF)
  - Dec. 2009: OpenFlow 1.0
  - Feb. 2011: OpenFlow 1.1
  - Dec. 2011: OpenFlow 1.2
  - Jun. 2012: OpenFlow 1.3
  - Oct. 2013: OpenFlow 1.4

# Where It's Going


- OF v1.3
  - multiple tables: leverage additional tables
  - tags and tunnels
  - multipath forwarding
  - per flow meters
- OF v2+
  - generalized matching and actions: protocol independent forwarding



# Multiple Tables (OF v1.1+)



{a} Packets are matched against multiple tables in the pipeline



① Find highest - priority matching flow entry

② Apply instructions:

- Modify packet & update match fields (apply actions instruction)
- Update action set (clear actions and/or write actions instructions)
- Update metadata

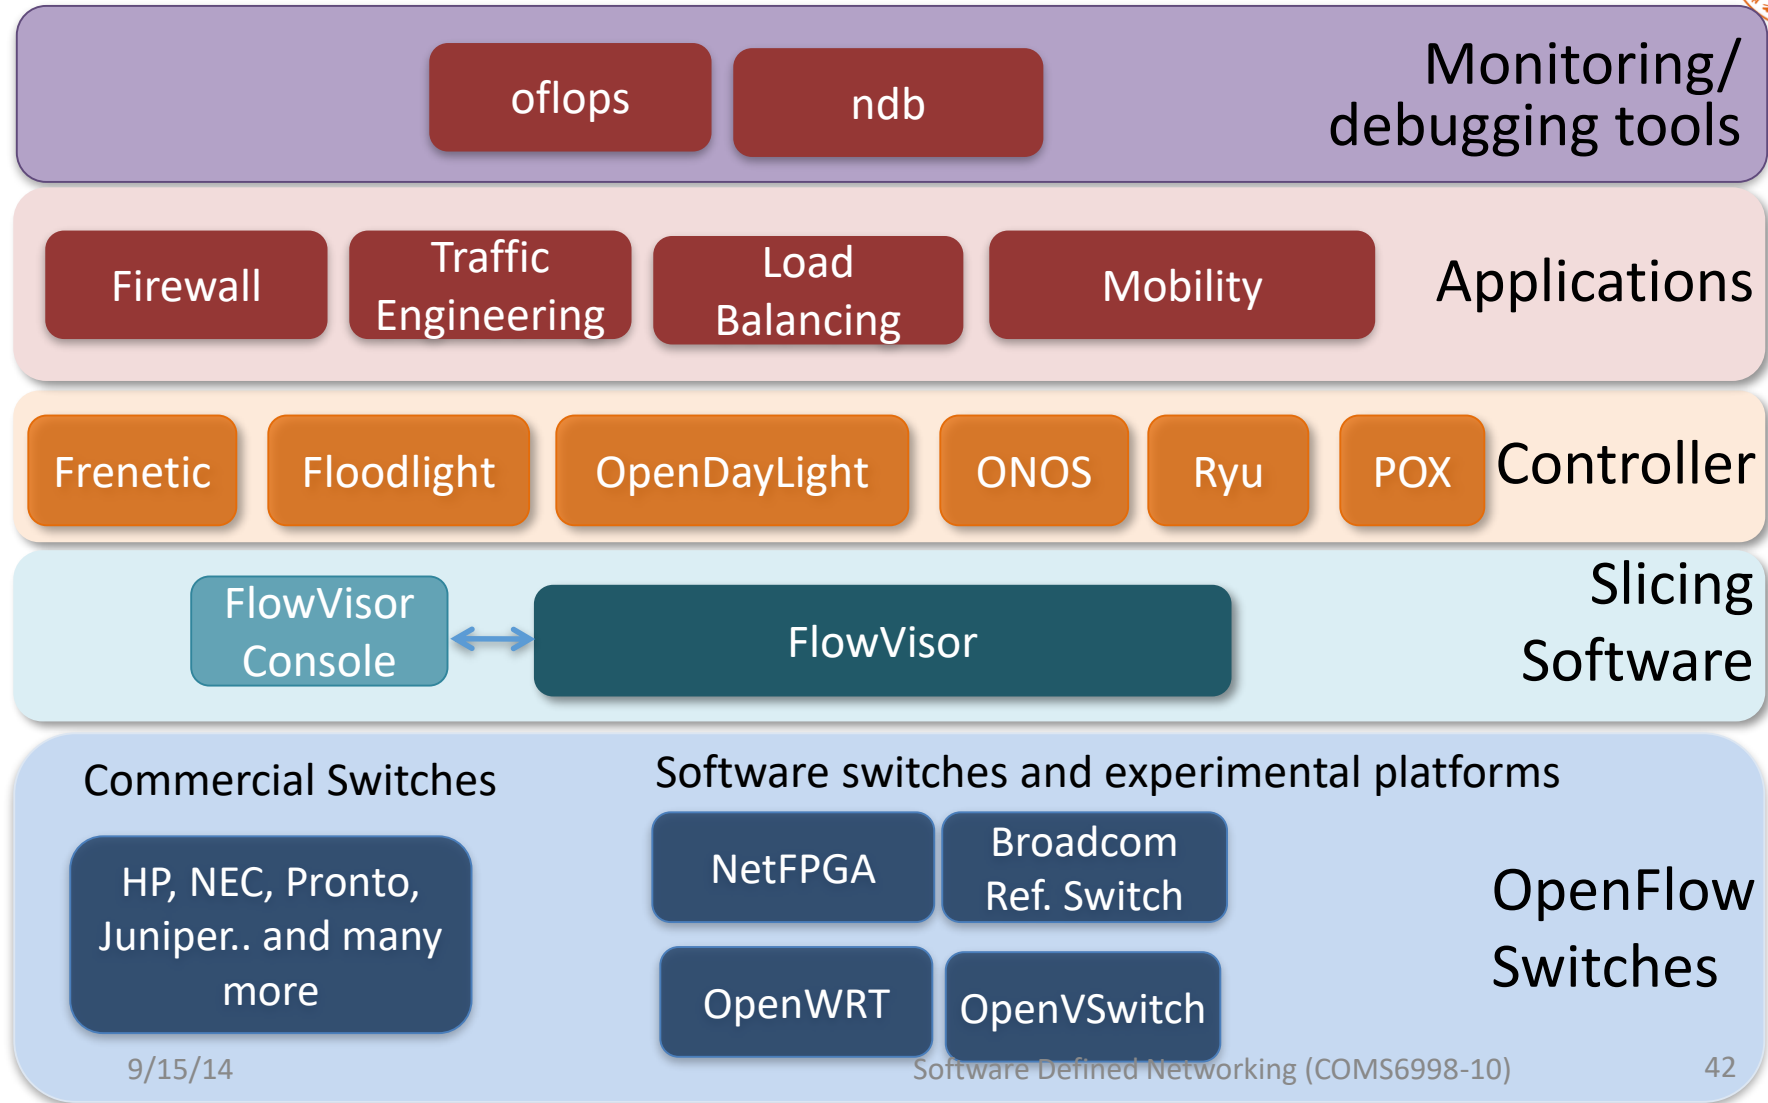
③ Send match data and action set to next table

{b} Per-table packet processing

# OpenFlow Building Blocks




南方科技大学  
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



# OpenFlow Configuration and Management Protocol: OF-Config




- Bootstrap OpenFlow network
  - Switch connects to controller
  - Controller(s) to connect to must be configured at switches
- Allocate resources within switches
  - Ports
  - Queues
  - . . .



# OpenFlow Configuration and Management Protocol: Reference Model



- Configuration Point
  - Source of switch configuration
- OpenFlow Capable Switch
  - Hosts one or more logical switches
- OpenFlow Controller
  - instance of an OpenFlow Switch



# Mininet


- Mininet is an SDN network emulator

```
# mn
# mn --topo tree,depth=3,fanout=3 --
link=tc,bw=10
mininet> xterm h1 h2
h1# wireshark &
h2# python -m SimpleHTTPServer 80 &
h1# firefox &
# mn --topo linear,100
# mn --custom custom.py --topo mytopo
```

# Mininet API Basics




```
net = Mininet()                                # net is a Mininet() object
h1 = net.addHost( 'h1' )                      # h1 is a Host() object
h2 = net.addHost( 'h2' )                      # h2 is a Host()
s1 = net.addSwitch( 's1' )                    # s1 is a Switch() object
c0 = net.addController( 'c0' )               # c0 is a Controller()
net.addLink( h1, s1 )                         # creates a Link() object
net.addLink( h2, s1 )
net.start()
h2.cmd( 'python -m SimpleHTTPServer 80 &' )
sleep( 2 )
h1.cmd( 'curl', h2.IP() )
CLI( net )
h2.cmd('kill %python')
net.stop()
```



# Performance Modeling






```
# Use performance-modeling link and host classes
net = Mininet(link=TCLink, host=CPULimitedHost)
# Limit link bandwidth and add delay
net.addLink(h2, s1, bw=10, delay='50ms')
# Limit CPU bandwidth
net.addHost('h1', cpu=.2)
```



# Switch Vendors



Model	Virtualize	Notes	
HP Procurve 5400zl or 6600	1 OF instance per VLAN	<ul style="list-style-type: none"><li>-LACP, VLAN and STP processing before OpenFlow</li><li>-Wildcard rules or non-IP pkts processed in s/w</li><li>-Header rewriting in s/w</li><li>-CPU protects mgmt during loop</li></ul>	
NEC IP8800	1 OF instance per VLAN	<ul style="list-style-type: none"><li>-OpenFlow takes precedence</li><li>-Most actions processed in hardware</li><li>-MAC header rewriting in h/w</li></ul>	
Pronto 3240 or 3290 with Pica8 or Indigo firmware	1 OF instance per switch	<ul style="list-style-type: none"><li>-No legacy protocols (like VLAN and STP)</li><li>-Most actions processed in hardware</li><li>-MAC header rewriting in h/w</li></ul>	



# Controller Vendors



Vendor	Notes
Nicira's NOX	<ul style="list-style-type: none"><li>•Open-source GPL</li><li>•C++ and Python</li><li>•Researcher friendly</li></ul>
Nicira's ONIX	<ul style="list-style-type: none"><li>•Closed-source</li><li>•Datacenter networks</li></ul>
Ryu	<ul style="list-style-type: none"><li>•Open-source GPL</li><li>•Python</li></ul>

Vendor	Notes
Stanford's Beacon	<ul style="list-style-type: none"><li>•Open-source</li><li>•Researcher friendly</li><li>•Java-based</li></ul>
BigSwitch controller	<ul style="list-style-type: none"><li>•Ha open source version</li><li>•Based on Beacon</li><li>•Enterprise network</li></ul>
OpenDayLight	<ul style="list-style-type: none"><li>•Open-source</li><li>•Based on Java</li></ul>
Frenetic	<ul style="list-style-type: none"><li>•Open-source</li><li>•Written in functional programming languages</li></ul>

# Outline

- Motivation for SDN
- SDN Overview
- OpenFlow
- **Case Study**

# Case Study: Google B4



- B4: a private WAN connecting Google's data centers across the planet



# Features and Goals



- **Features of B4**

- Massive bandwidth requirements deployed to a modest number of sites
- Elastic traffic demand that seeks to maximize average bandwidth
- Full control over the edge servers and network, which enables rate limiting and demand measurement at the edge

- **Goals**

- TE: Centralized traffic engineering service that can drive links to near 100% utilization
- Split application flows among multiple paths to balance capacity against application priority/demands

# B4 Architecture Overview

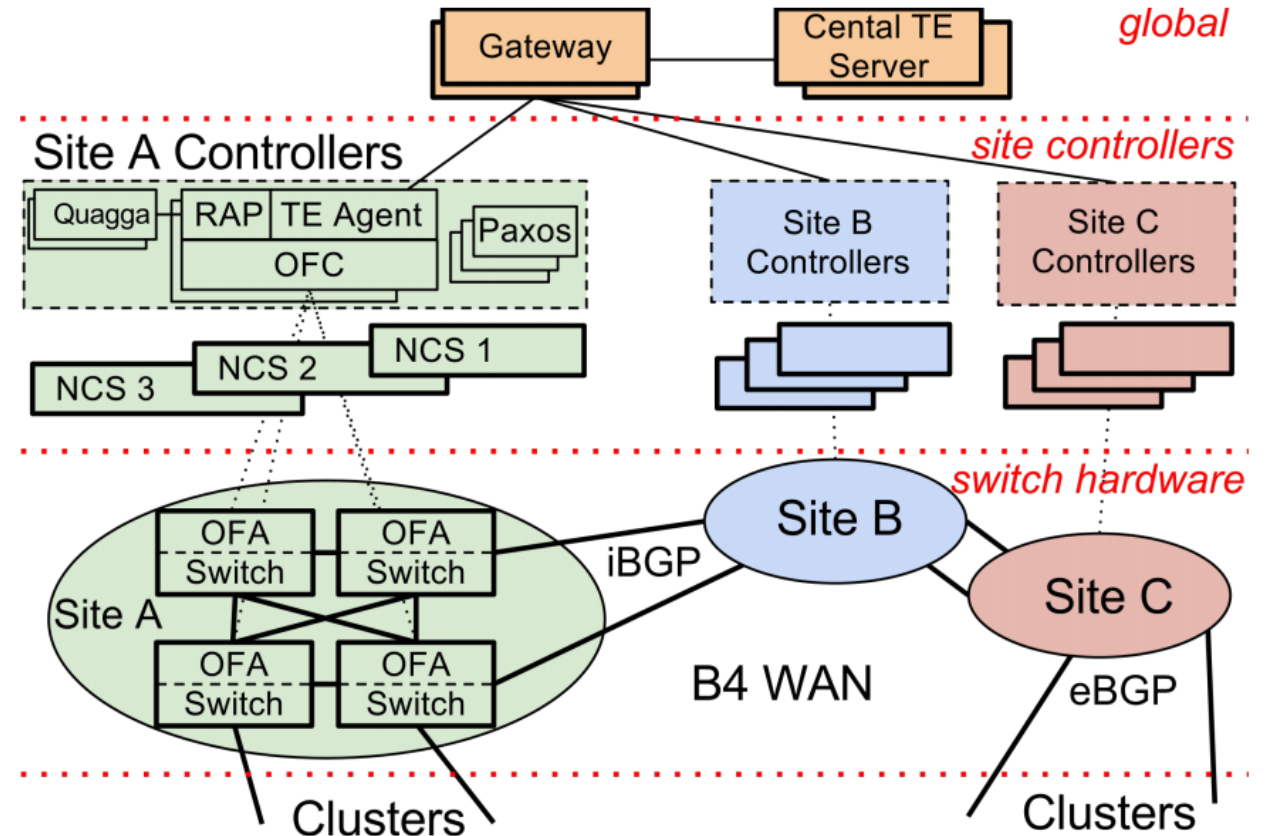


- **Switch layer**

- Primarily forwards traffic and does not run complex control software

- **Site controller layer**

- Consists of Network Control Servers (NCS) hosting both OpenFlow controllers (OFC) and Network Control Applications (NCAs)
- These servers enable distributed routing and central traffic engineering as a routing overlay




# B4 Architecture Overview



- **Global layer**



- Consists of logically centralized apps (e.g. an SDN Gateway and a central TE server) that enable the central control of the entire network via the site-level NCAs
- The SDN Gateway abstracts details of OpenFlow and switch hardware from the central TE server



# Switch Design



- Build B4 switches from multiple merchant silicon switch chips in a two-stage Clos topology with a copper backplane
  - OpenFlow Agent (OFA): a user-level process running on the switch hardware
  - OpenFlow Controller (OFC): OFA connects to a remote OFC, accepting OpenFlow (OF) commands and forwarding packets and link/switch events to the OFC





# Switch Design



## More details:

B4: Experience with a Globally-Deployed  
Software Defined WAN, [SIGCOMM 2013]

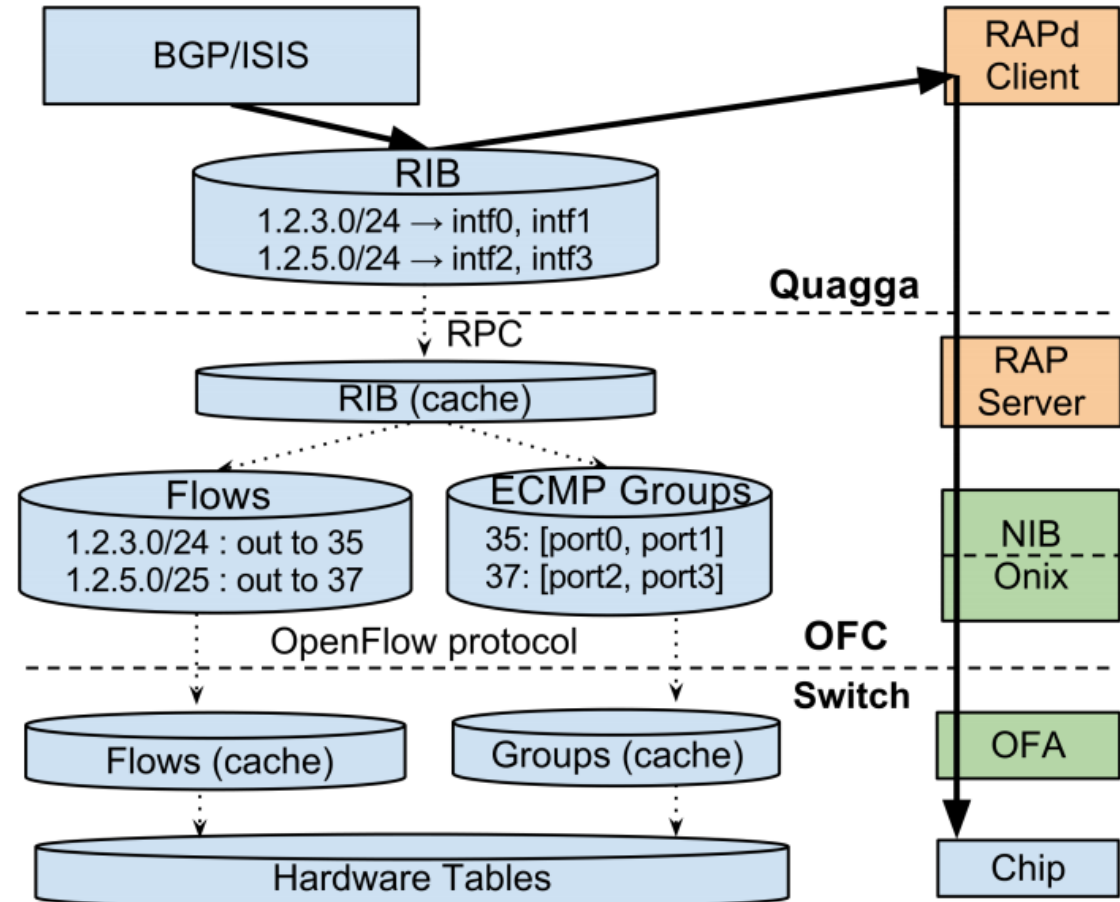


Figure 4: Integrating Routing with OpenFlow Control.