

Software Defined Networking (2/2)



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

Xuetao Wei

`weixt@sustech.edu.cn`

Outline

- Open vSwitch (OVS)
- Towards OpenFlow 2.0
- New Forwarding Plane Architectures
- Programming the Forwarding Plane

Where It's Going

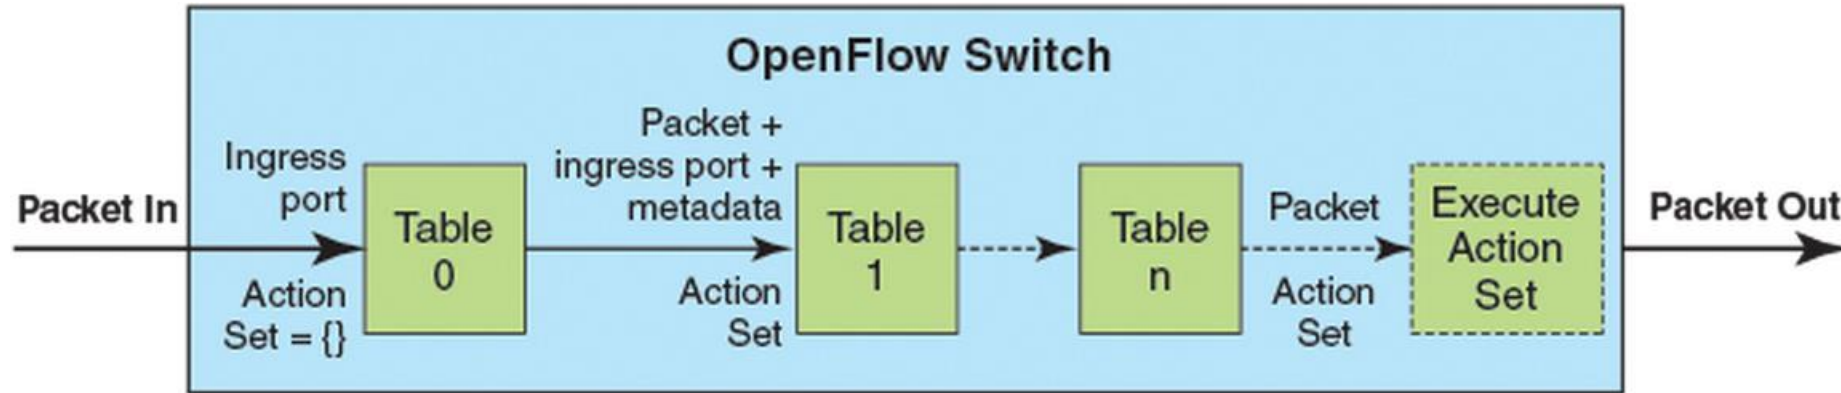
- **OF v1.x**

- multiple tables: leverage additional tables
- tags and tunnels
- multipath forwarding
- per flow meters


- **OF v2+ (yet to come)**

- generalized matching and actions: protocol independent forwarding

Multiple Tables (OF v1.x)



{a} Packets are matched against multiple tables in the pipeline



① Find highest - priority matching flow entry

② Apply instructions:

- i. Modify packet & update match fields (apply actions instruction)
- ii. Update action set (clear actions and/or write actions instructions)
- iii. Update metadata

③ Send match data and action set to next table

{b} Per-table packet processing

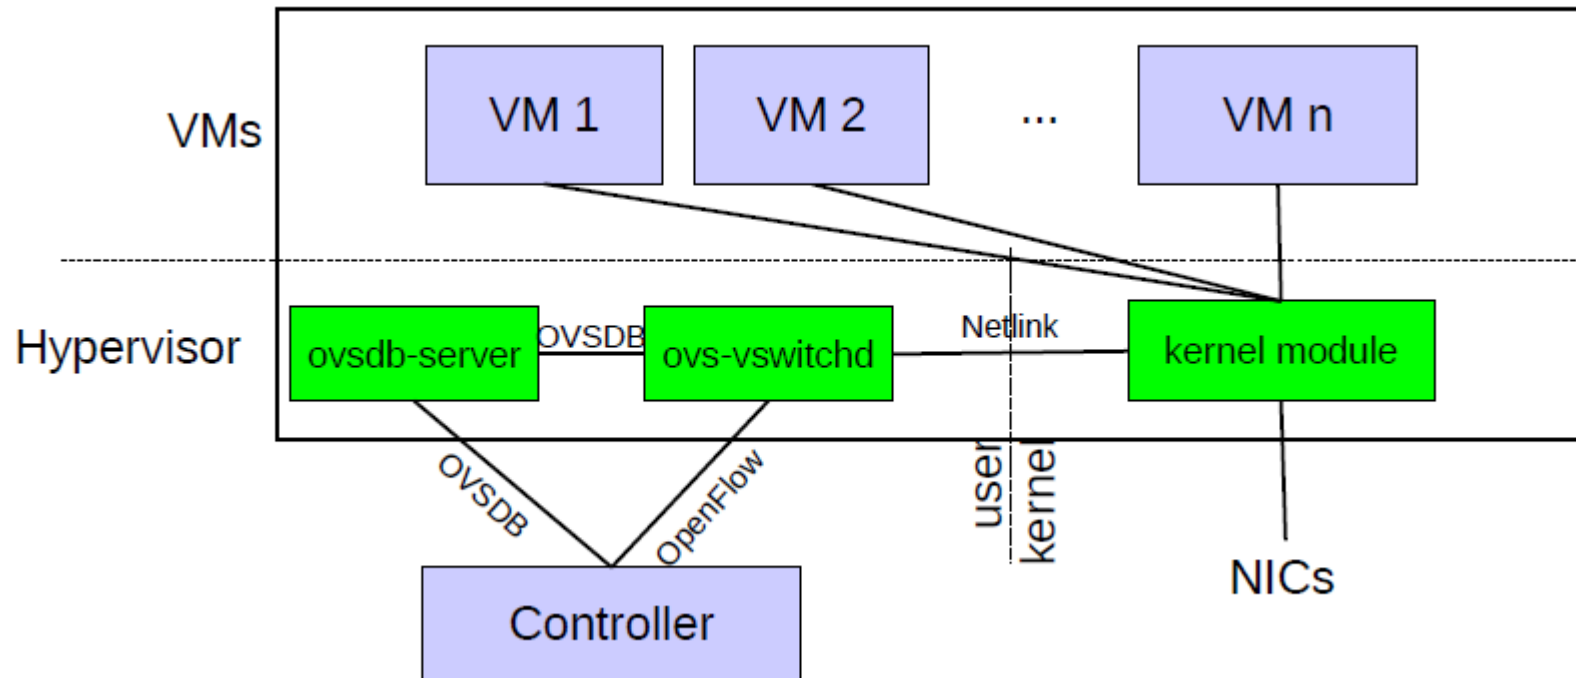
Open vSwitch (OVS)

- A virtual switch conforming to OpenFlow standard that is implemented in software
- Available from *openvswitch.org*
- Development code is available in git
- User-space (controller and tools) is under the Apache license
- Kernel (datapath) is under the GPLv2

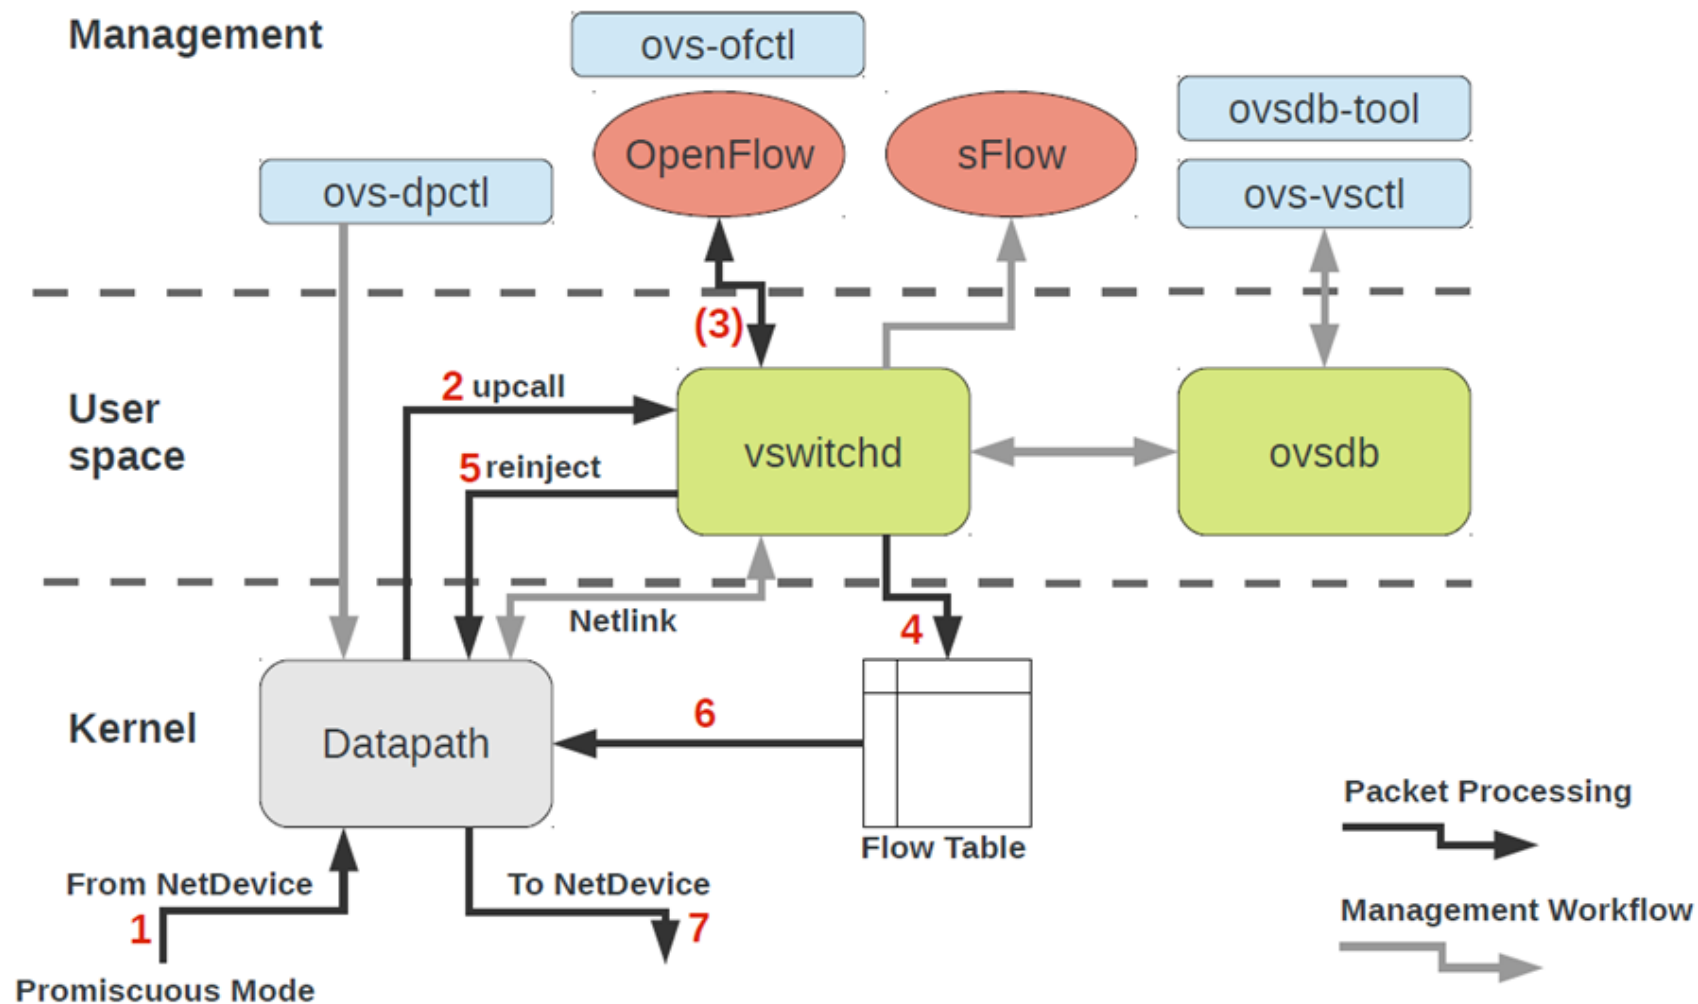
OVS Architecture



- **ovs-vswitchd**: core implementation of the switch
- **ovsdb-server**: manipulates the database of the vswitch configuration & flows



OVS Workflow



Usage (1/4)



- Check OpenFlow version supported by OVS

```
1 | $ ovs-ofctl --version
2 | ovs-ofctl (Open vSwitch) 1.11.1
3 | Compiled Oct 28 2013 14:17:17
4 | OpenFlow versions 0x1:0x4
```

- Create a new OVS switch

```
1 | $ ovs-vsctl add-br ovs-switch
```

- Create a new port p0 with port number 100

```
1 | $ ovs-vsctl add-port ovs-switch p0 -- set Interface p0 ofport_request=100
```


Usage (2/4)



- Check the switch created

```
1 $ ovs-ofctl show ovs-switch
2 OFPT_FEATURES_REPLY (xid=0x2): dpid:00001232a
3 n_tables:254, n_buffers:256
4 capabilities: FLOW_STATS TABLE_STATS PORT_STATS
5 actions: OUTPUT SET_VLAN_VID SET_VLAN_PCP STRIP
6 SET_NW_SRC SET_NW_DST SET_NW_TOS SET_TP_SRC SET_TP_DST
7 100(p0): addr:54:01:00:00:00:00
8     config:      PORT_DOWN
9     state:       LINK_DOWN
10    speed: 0 Mbps now, 0 Mbps max
11 101(p1): addr:54:01:00:00:00:00
12    config:      PORT_DOWN
13    state:       LINK_DOWN
14    speed: 0 Mbps now, 0 Mbps max
15 102(p2): addr:54:01:00:00:00:00
16    config:      PORT_DOWN
17    state:       LINK_DOWN
18    speed: 0 Mbps now, 0 Mbps max
19 LOCAL(ovs-switch): addr:12:32:a2:37:ea:45
20    config:      0
21    state:       0
22    speed: 0 Mbps now, 0 Mbps max
23 OFPT_GET_CONFIG_REPLY (xid=0x4): frags=normal
```

Usage (3/4)



- Deny all packets

```
1 | $ ovs-ofctl add-flow ovs-switch "table=0, dl_src=01:00:00:00:00:00/01:00:00:00:00:00, actions=drop"
```

- Dump all entries in flow tables

```
1 | ovs-ofctl dump-flows ovs-switch
```

- Change source address to 9.181.137.1 for all packets received on port p0

```
1 | $ ovs-ofctl add-flow ovs-switch "priority=1 idle_timeout=0,\  
2 |   in_port=100,actions=mod_nw_src:9.181.137.1,normal"
```

Usage (4/4)



- Sending testing packet from port p0 (192.168.1.100) to port p1 (192.168.1.101)

```
1 | $ ip netns exec ns0 ping 192.168.1.101
```

- Monitoring received packets at port p1, found out that its source address has been changed to 9.181.137.1

```
1 | $ ip netns exec ns3 tcpdump -i p2 icmp
2 | tcpdump: verbose output suppressed, use -v or -vv for full protocol decode
3 | listening on p2, link-type EN10MB (Ethernet), capture size 65535 bytes
4 | 16:07:35.677770 IP 192.168.1.100 > 192.168.1.101: ICMP echo request, id 23147, seq 25, length 64
5 | 16:07:36.685824 IP 192.168.1.100 > 192.168.1.101: ICMP echo request, id 23147, seq 26, length 64
```

Using Controller for OVS



- **Install floodlight**

- <http://www.projectfloodlight.org/getting-started/>
- `git clone git://github.com/floodlight/floodlight.git`
- `cd floodlight/`
- `ant`
- `nohup java -jar target/floodlight.jar > floodlight.log 2>&1 &`

- **Set Controller**


- `ovs-vsctl set-controller ubuntu_br tcp:192.168.100.1:6633`

Outline

- Open vSwitch (OVS)
- **Towards OpenFlow 2.0**
- New Forwarding Plane Architectures
- Programming the Forwarding Plane


Evolving of OpenFlow

- OpenFlow 1.0 (2008)
 - Single table with 12 fields



Version	Date	Header Fields
OF 1.0	Dec 2009	12 fields (Ethernet, TCP/IP _{v4})
OF 1.1	Feb 2011	15 fields (MPLS, inter-table metadata)
OF 1.2	Dec 2011	36 fields (ARP, ICMP, IPv6, etc.)
OF 1.3	Jun 2012	40 fields
OF 1.4	Oct 2013	41 fields

- OpenFlow 1.1 ~ 1.4 (2011 ~ 2013)
 - Multiple tables in pipelined processing
 - The number of fields increases from 12 to 40+



Evolving of OpenFlow

- **OpenFlow 2.0** (2013~)
 - Match & Action on any #fields, any combination of fields, and even any sequences of header bits
- **PIF**: Protocol Independent Forwarding (2013)
 - Stanford, Princeton
 - Coupled integrated with switch architecture and programming language
 - More towards ASIC implementations
- **POF**: Protocol Oblivious Forwarding (2013)
 - Huawei
 - Without specific target switch architecture in mind
 - More towards CPU/NPU implementations

POF: Protocol Oblivious Forwarding



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- Current OpenFlow-enabled Device
 - Decoupled control/data planes, “dumb” packet forwarding devices
 - Centralized control of network, intelligent controllers
 - Multiple flow tables


But...

- Still protocol conscious packet forwarding

Read:


Protocol-oblivious Forwarding: Unleash the Power of SDN Through a Future-proof Forwarding Plane [\[HotSDN'13\]](#)

Protocol Conscious Packet Forwarding



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

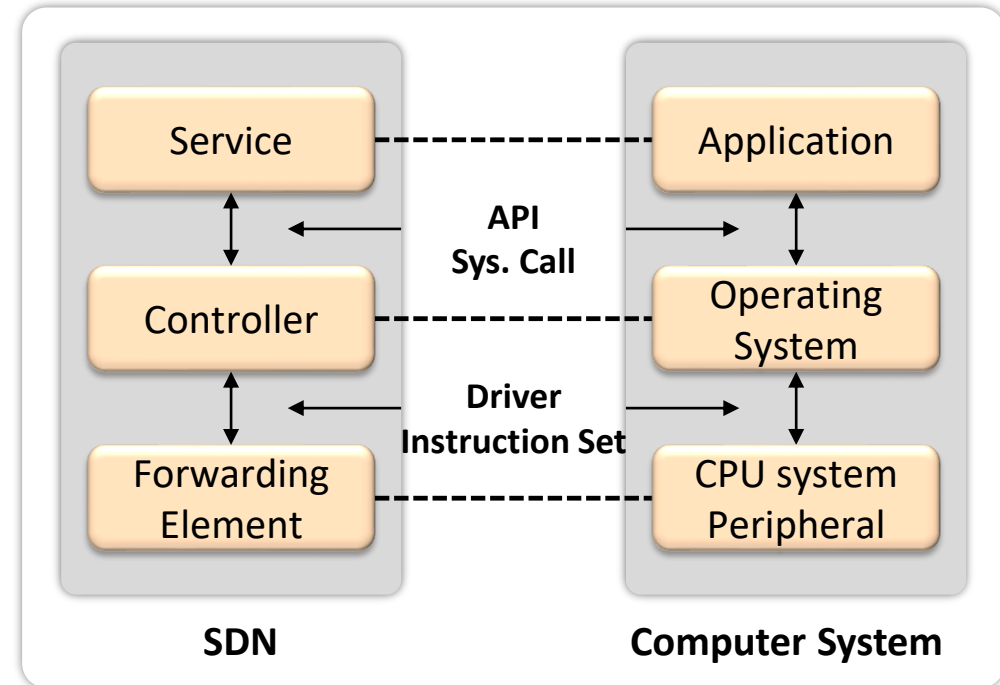
- Fixed parser extracts fixed set of header fields
- Pre-defined tables implement corresponding protocols or services
- Once deployed, switch functionalities cannot be programmed to support other protocols or services (only table content can be programmed)



Operate your Network Device like a PC



- Simple & generic instruction set
- Ultimate flexibility & extensibility
- Upgrade only on performance



Computer system components have been decoupled from the vertical integration model. SDN is on the track to mimic this transition. But, current OpenFlow still doesn't embrace this model to the full extent.

Core Concept of POF



- Table search keys are defined as {offset, length} tuples
- Instructions/Actions access packet data or metadata using {offset, length} tuples
- Include other math, logic, move, branching, and jump instructions

Match

~40 matching header fields defined yet still **many** uncovered protocols/headers

Action

OFPAT_COPY_TTL_OUT
OFPAT_COPY_TTL_IN
OFPAT_SET_MPLS_TTL
OFPAT_DEC_MPLS_TTL
OFPAT_PUSH_VLAN
OFPAT_POP_VLAN
OFPAT_PUSH_MPLS
OFPAT_POP_MPLS
OFPAT_SET_NW_TTL
OFPAT_DEC_NW_TTL
and on and on and on ...



{offset, length} covers **any** frame based formats

POFAT_SET_FIELD
POFAT_ADD_FIELD
POFAT_DELETE_FIELD
POFAT_MOD_FIELD
Period.

Current OpenFlow

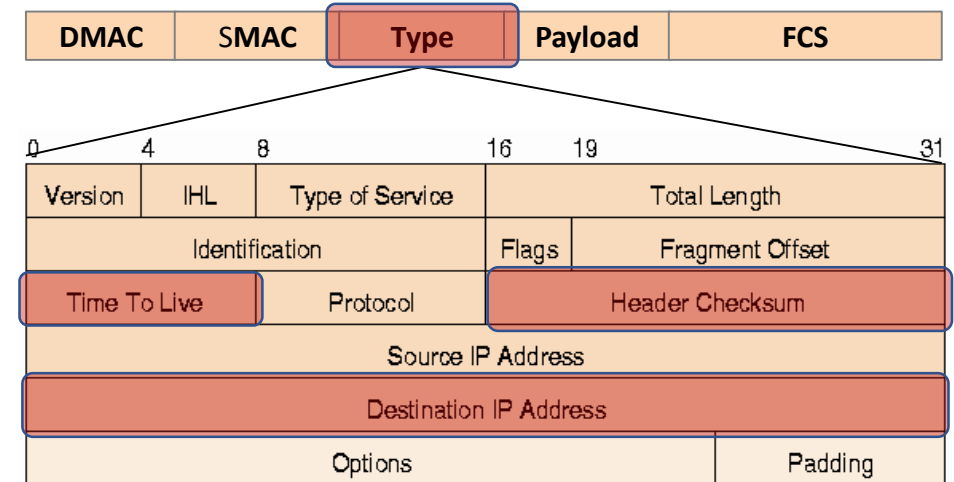
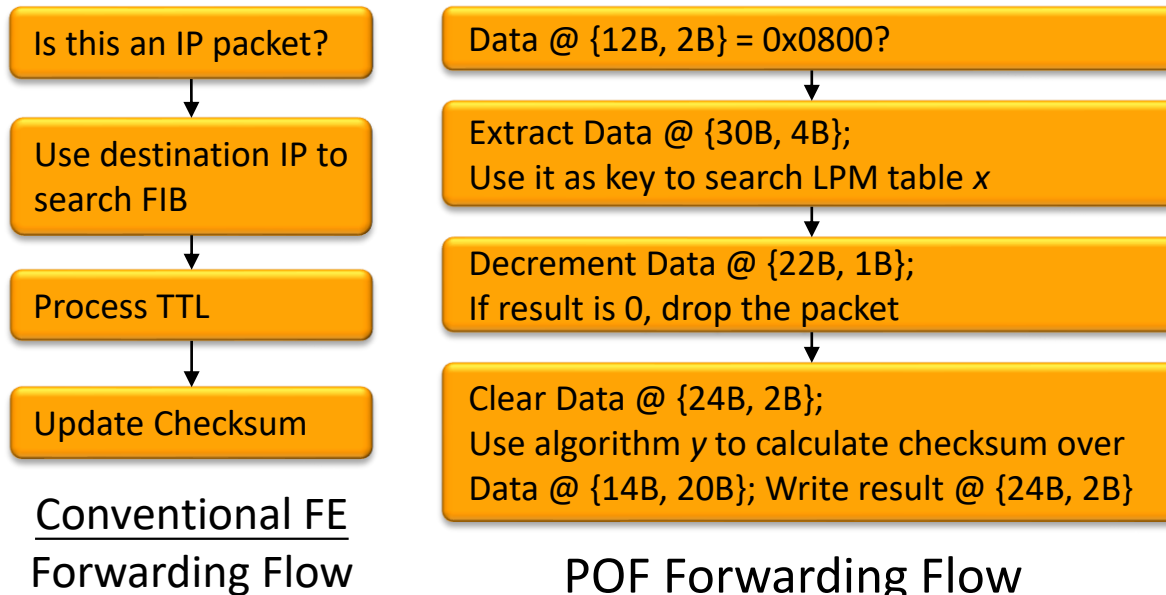
POF

Packet field parsing and handling are abstracted as generic instructions to enable flexible and future proof forwarding elements. This is simple yet has profound implications to SDN.

Ask a Dumb FE to Do Smart Things



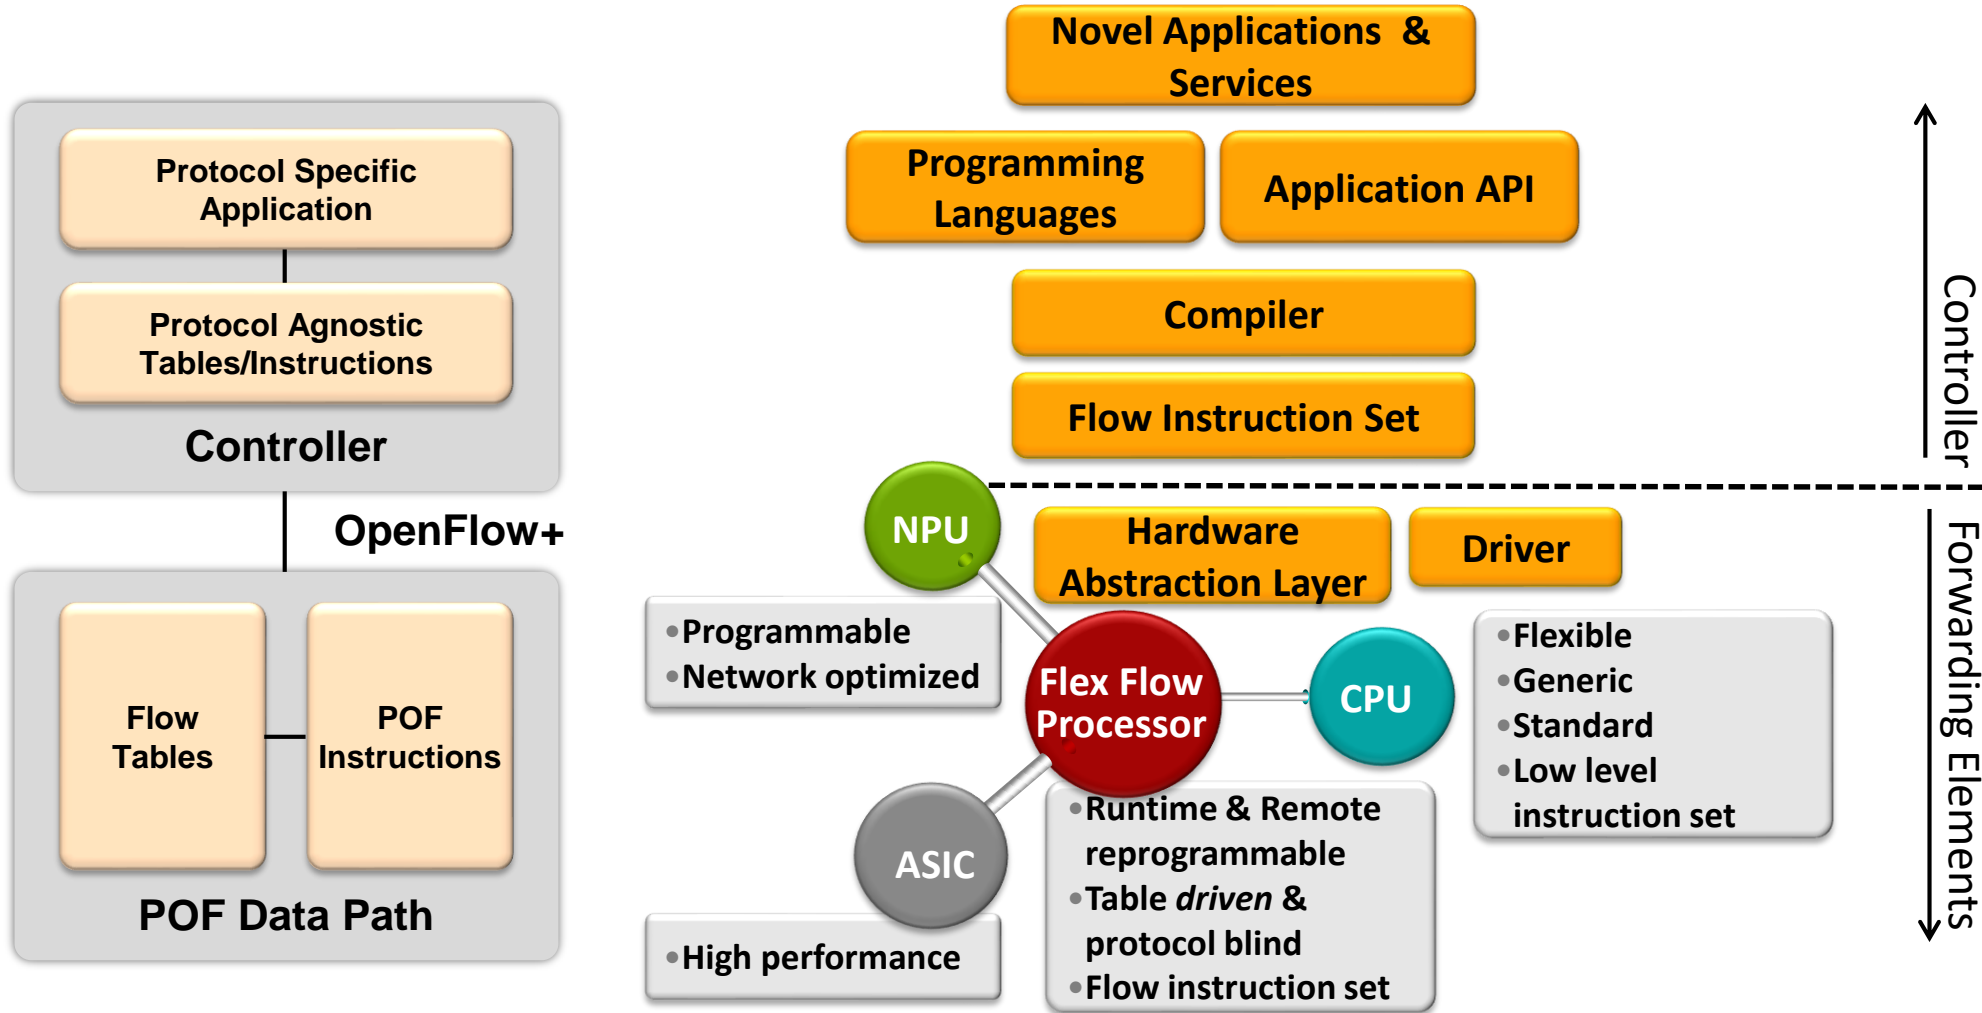
- The fine-grained bit-level manipulations used to be hardcoded or micro-coded in the FE are now explicitly described by controller




Ethernet/IPv4 Packet Format

OpenFlow's high level semantics ("what") is simple in communication but demands forwarding plane intelligence; POF's low level semantics ("how") moves all the intelligence up to the controller

POF-based SDN Architecture



PIF: Protocol Independent Forwarding



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY

- **Phase 0:** Initially, the switch does not know what a protocol is, or how to process packets (Protocol Independence)
- **Phase 1:** We tell the switch how we want it to process packets (Configuration)
- **Phase 2:** The switch runs (Run-time)

Three Goals

- **Protocol independence**

- Configure a packet parser to extract relevant header fields
- Define a set of typed <match, action> tables

- **Target independence**

- Program without knowledge of switch details
- Rely on compiler to configure the target switch

- **Reconfigurability**

- Change parsing and processing in the field

The Abstract Forwarding Model



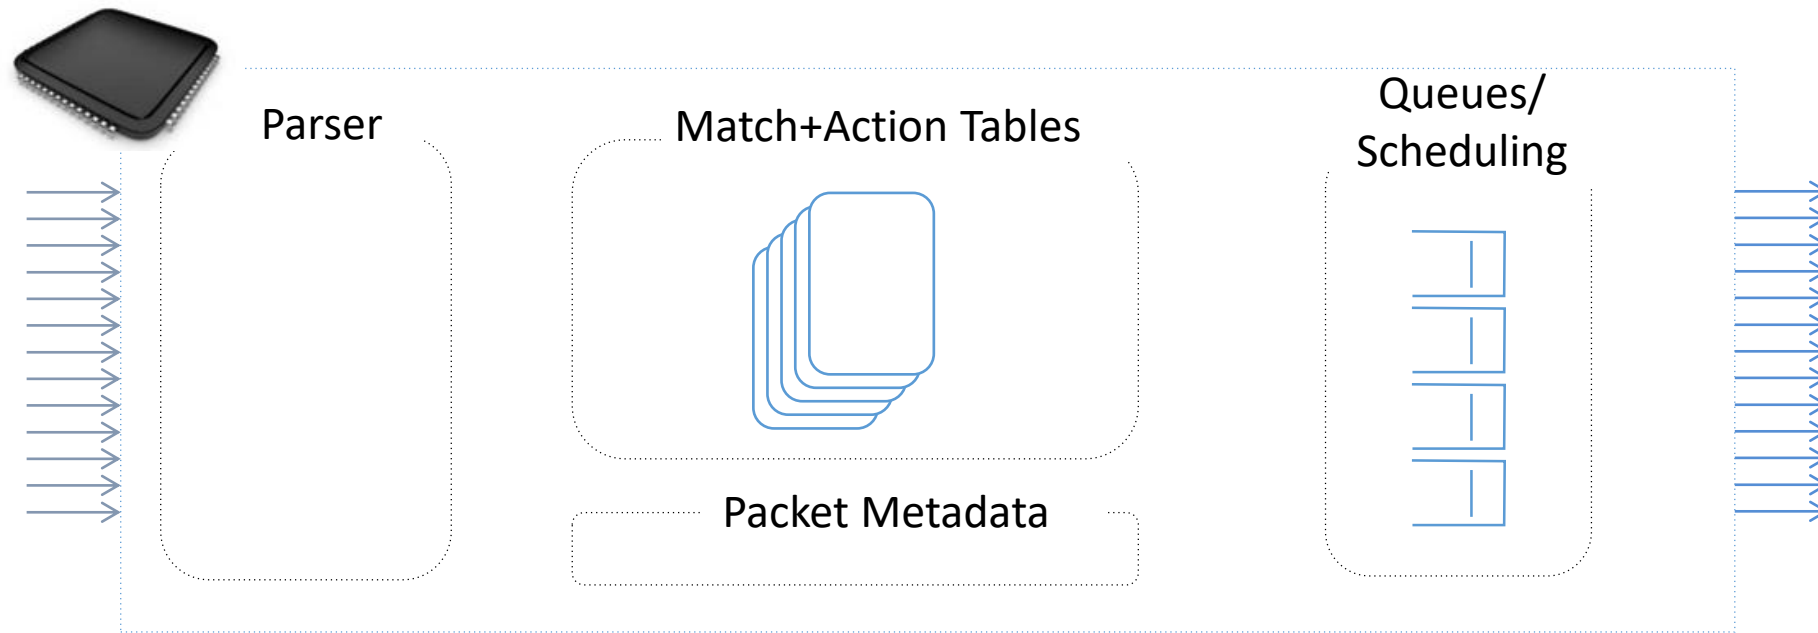
南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



The Abstract Forwarding Model



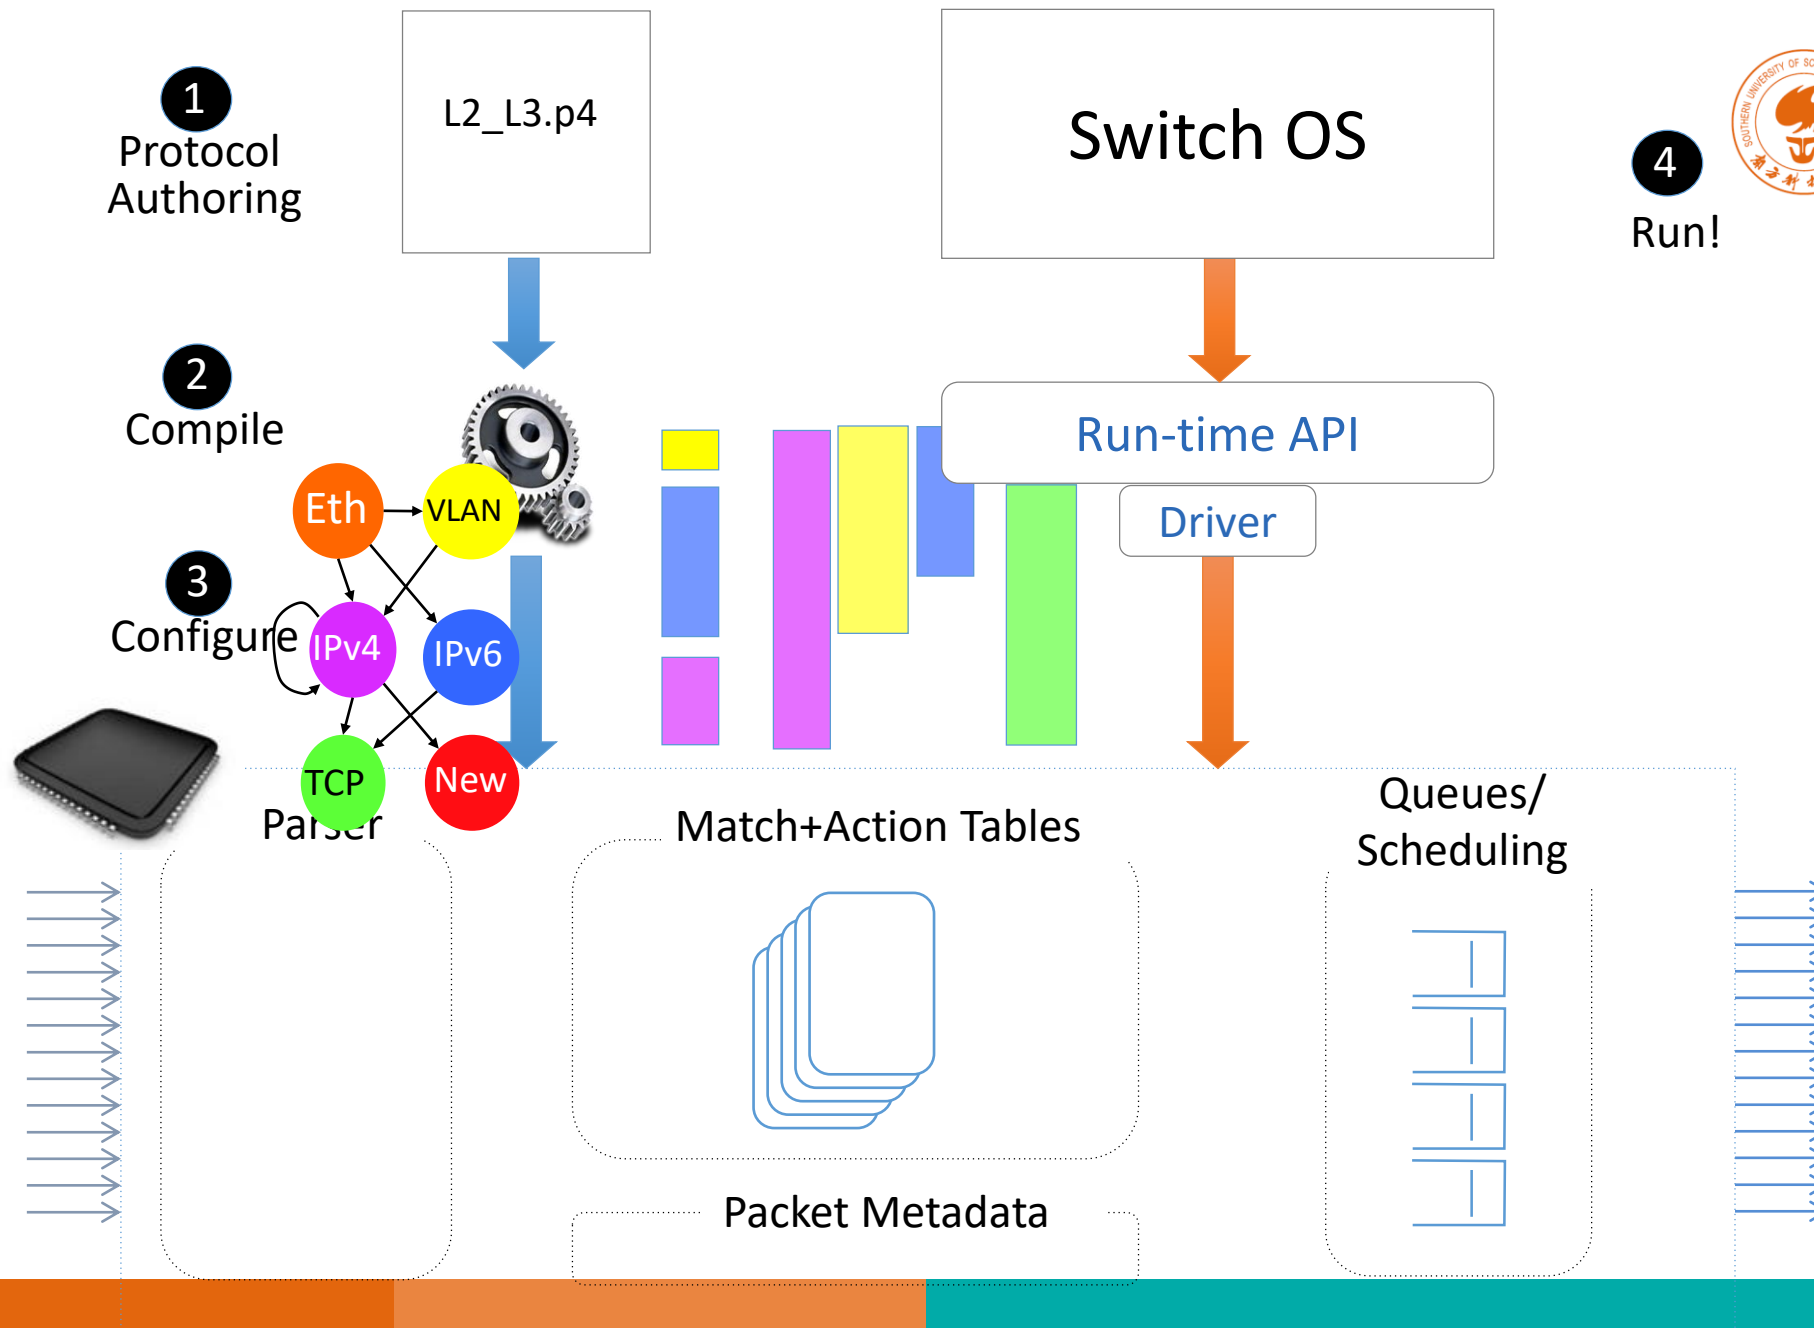
Initially, a switch is unprogrammed and does not know any protocols.

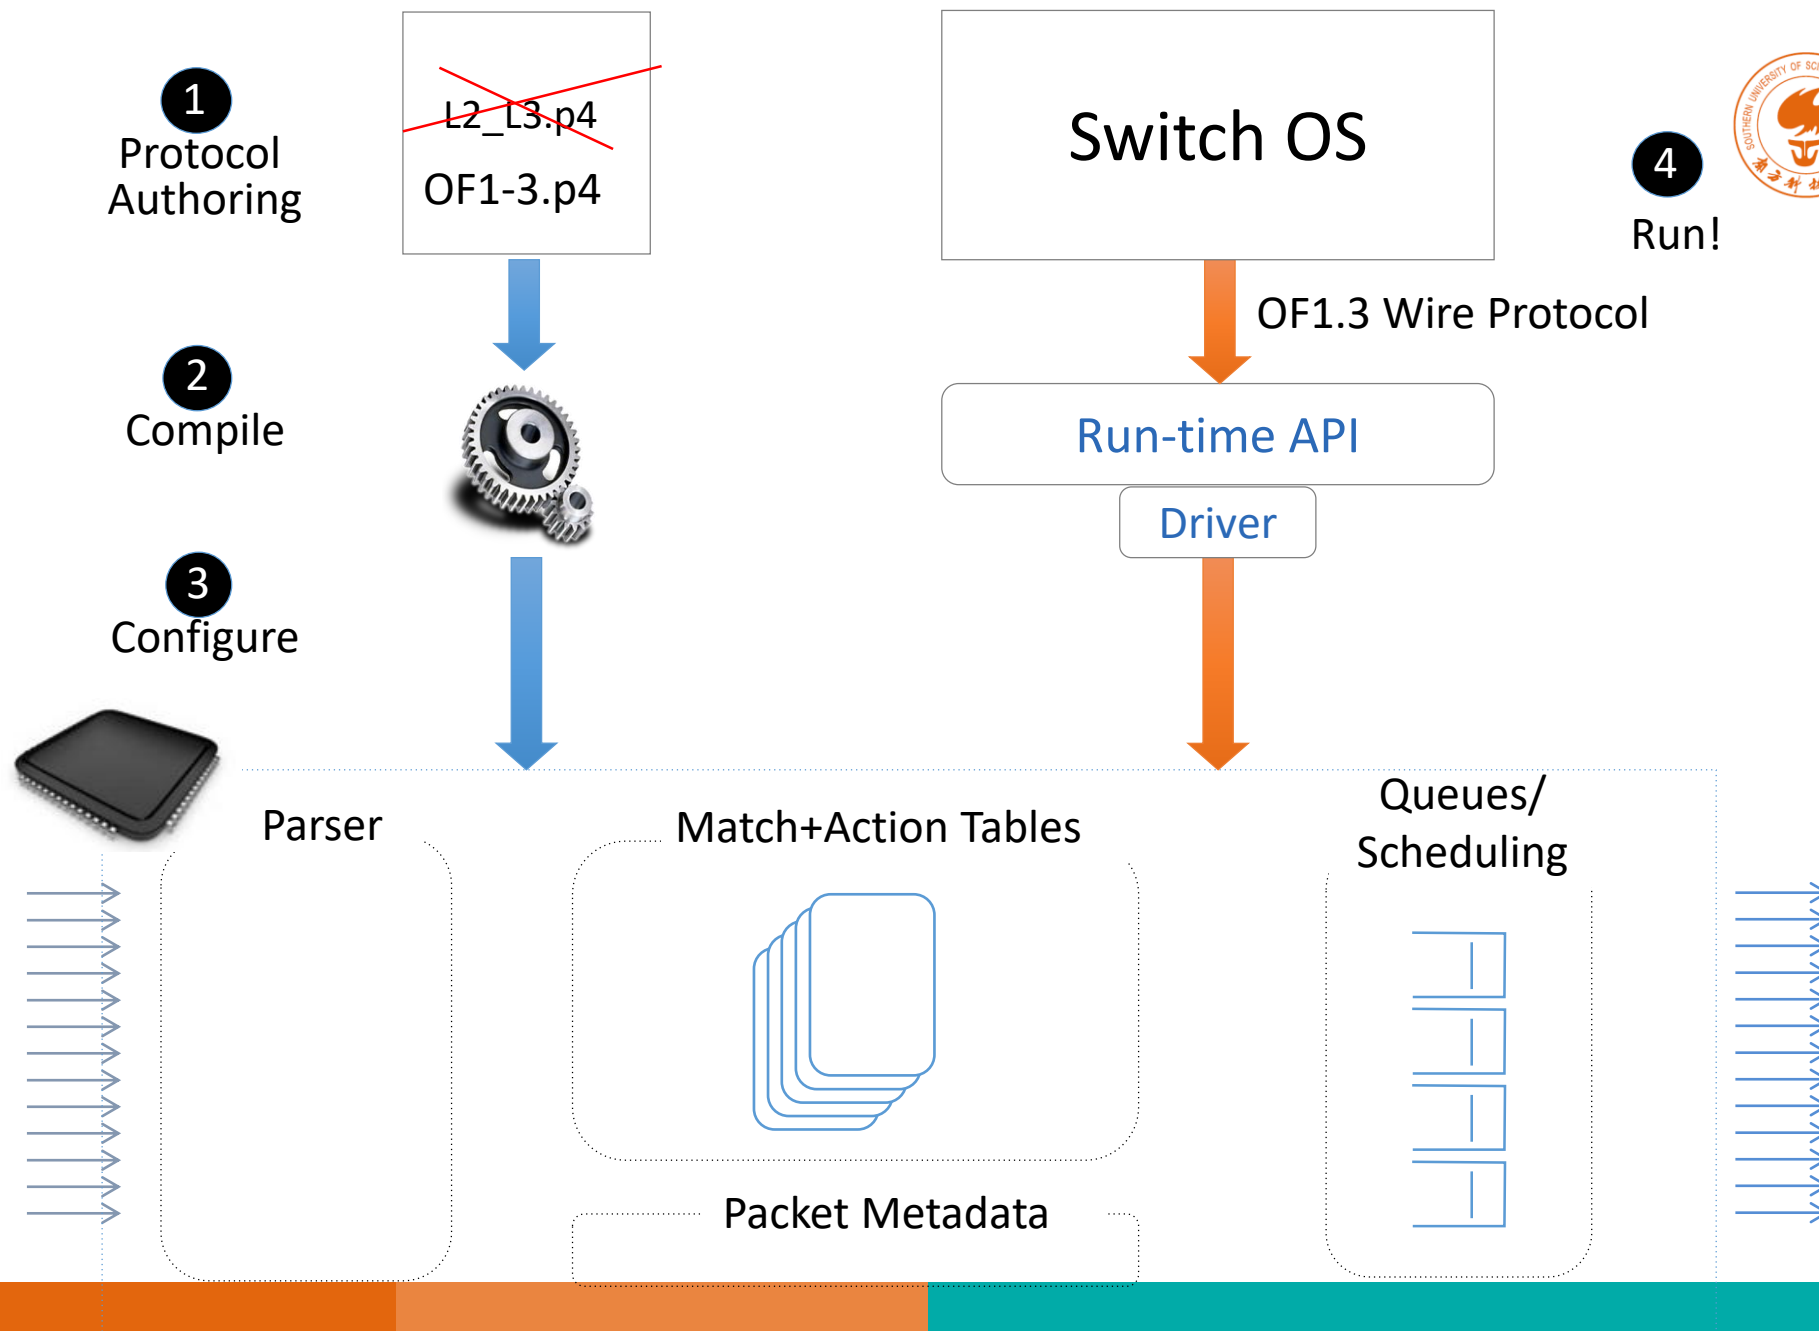




4

Run!





Outline

- Open vSwitch (OVS)
- Towards OpenFlow 2.0
- **New Forwarding Plane Architectures**
- Programming the Forwarding Plane

Match-Action Models

- **SMT (Single Match Table)**

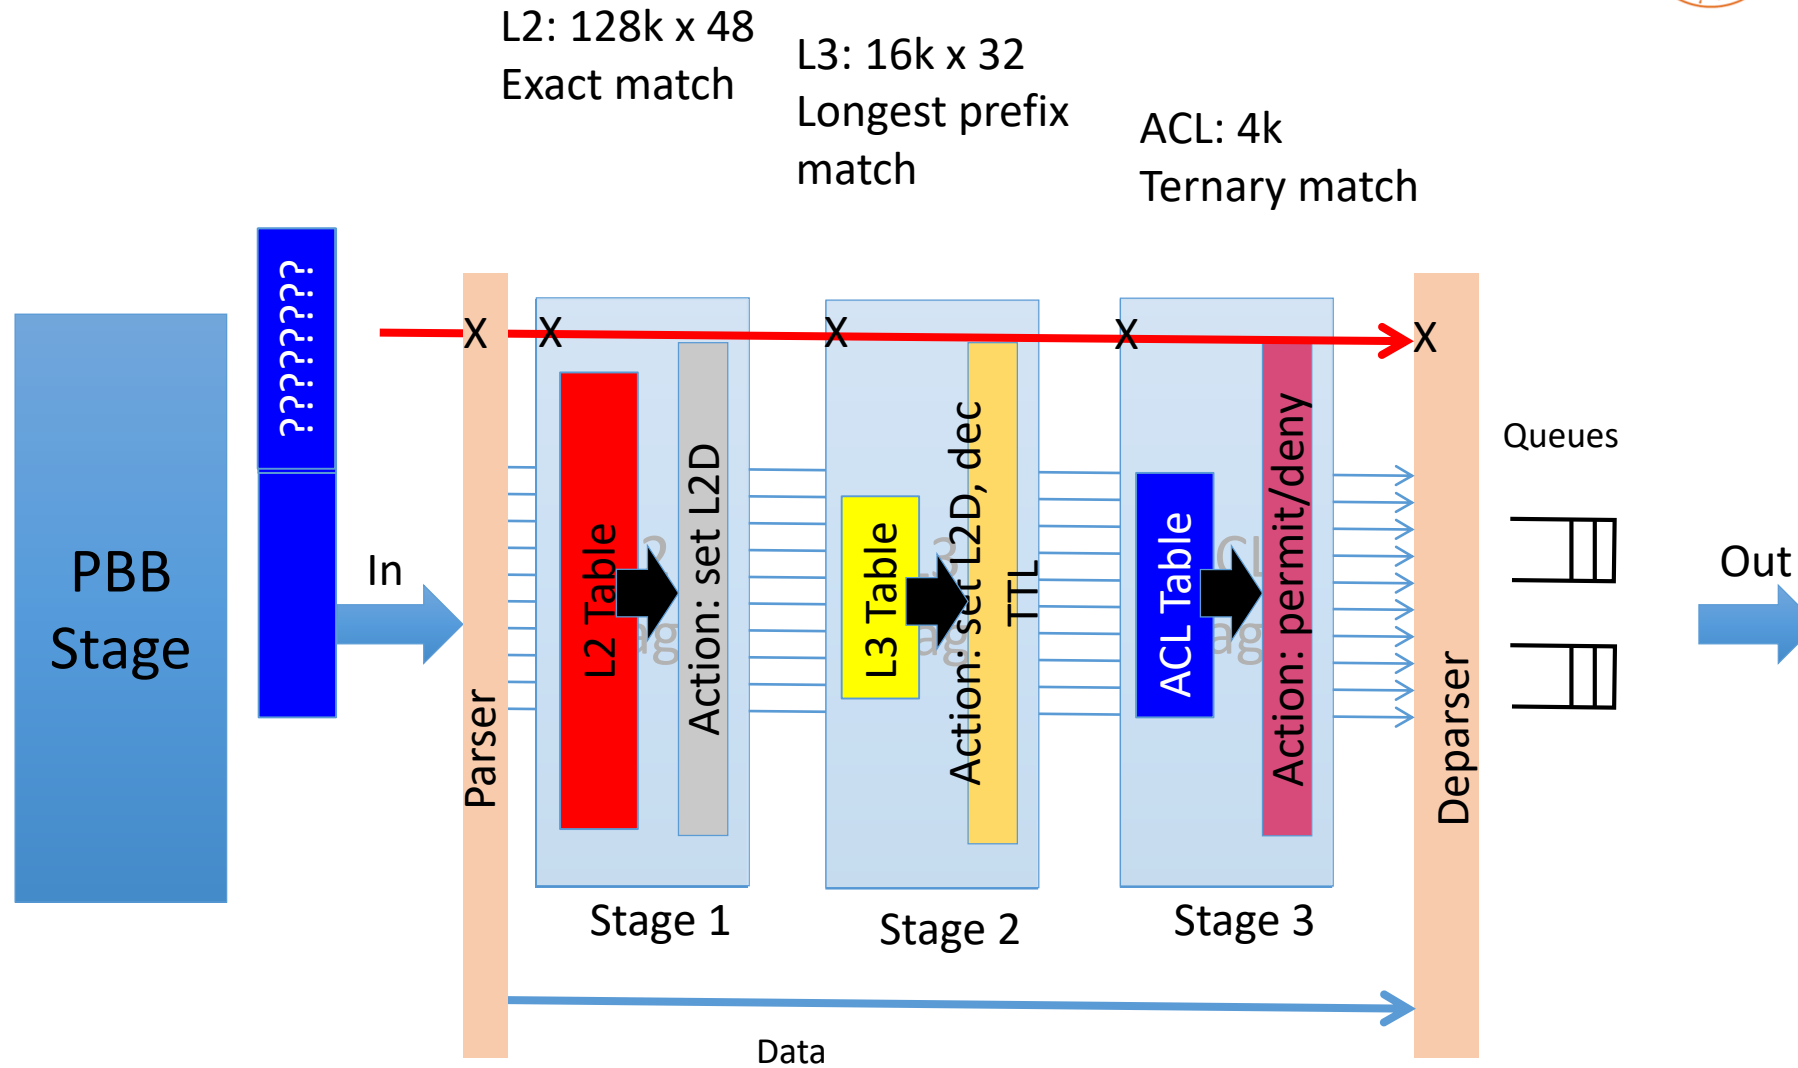
- The controller tells the switch to match any set of packet header fields against entries in a single match table
- It assumes that a parser extracts the correct header fields to match against the table
- SMT can be implemented using a wide Ternary Content Addressable Memory (TCAM)

- **MMT (Multiple Match Tables)**

- It allows multiple smaller match tables to be matched by a subset of packet fields
- The match tables are arranged into a pipeline of stages
- Existing switch chips implement a small (4-8) number of tables whose widths, depths, and execution order are set when the chip is fabricated, severely limiting their flexibility

- **RMT (Reconfigurable Match Tables)**

A Fixed Function Switch



The RMT Abstract Model

- RMT: Reconfigurable Match Tables
- Parse graph
- Table graph


Read:

Forwarding Metamorphosis: Fast Programmable Match-action Processing in Hardware for SDN [\[SIGCOMM'13\]](#)


Parse Graph



Packet:



- Describing arbitrary header fields to be extracted




Parse Graph



Packet:



- Describing arbitrary header fields to be extracted



Parse Graph



Packet:

Ethernet	RCP	IPV4	TCP
----------	-----	------	-----

- Describing arbitrary header fields to be extracted






Table Graph



Changes to Parse Graph and Table Graph



Parse Graph

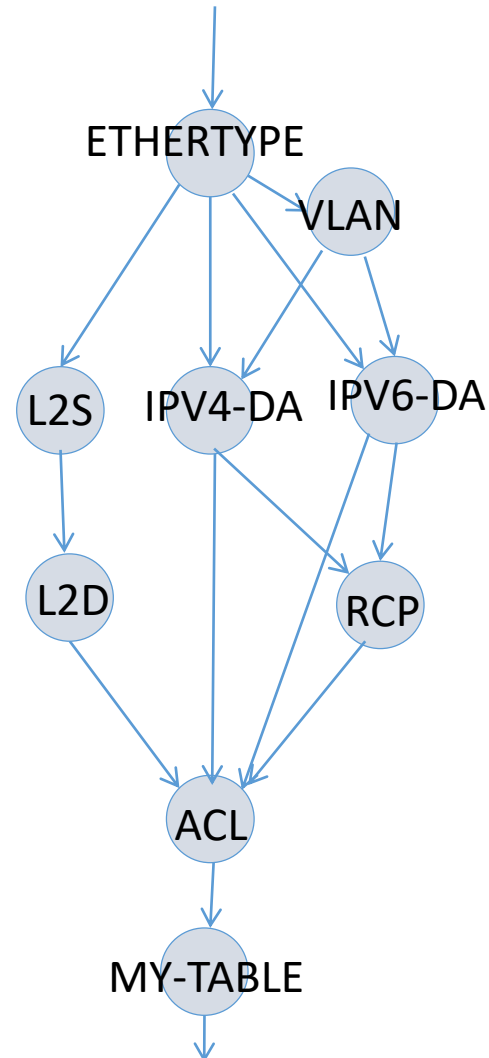
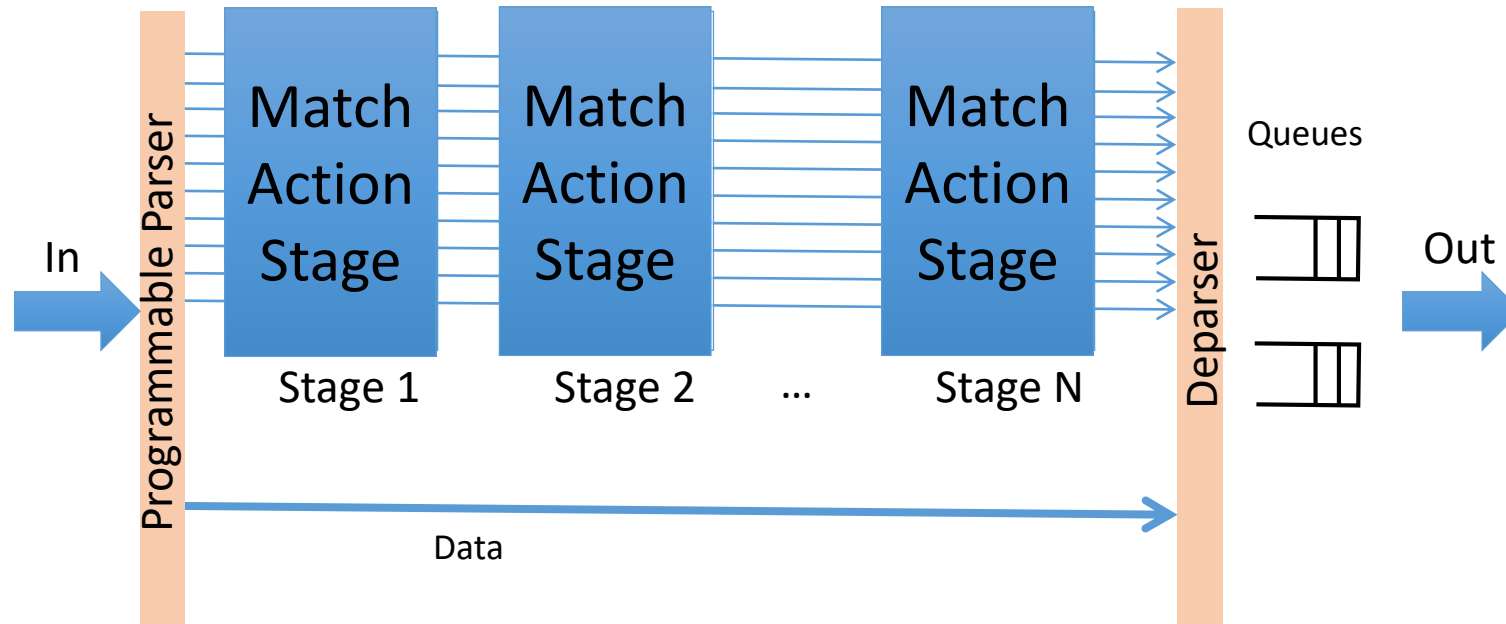


Table Graph





But the **Parse Graph** and **Table Graph**
don't show you how to build a switch




Match/Action Forwarding Model



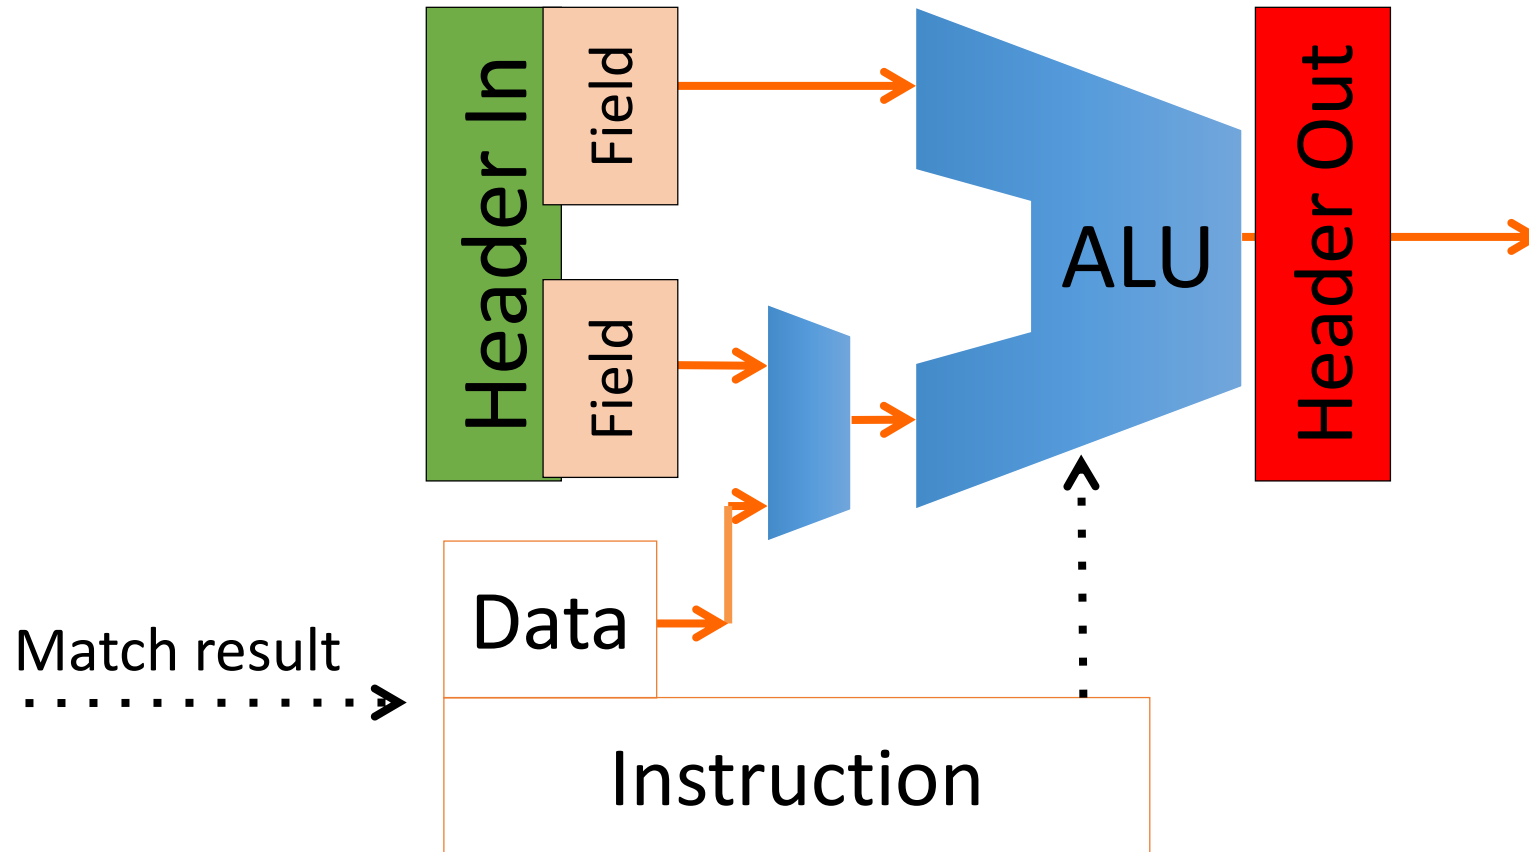
RMT Logical to Physical Table Mapping




RMT Logical to Physical Table Mapping




Action Processing Unit



Modeled as Multiple VLIW CPUs per Stage



南方科技大学
SOUTHERN UNIVERSITY OF SCIENCE AND TECHNOLOGY



RMT Chip Configuration



- 64 x 10Gb ports
 - 960M packets/second
 - 1GHz pipeline
- Programmable parser
- 32 Match/action stages
- Huge TCAM: 10x current chips
 - 64K TCAM words x 640b
- SRAM hash tables for exact matches
 - 128K words x 640b
- 224 action processors per stage
- All OpenFlow statistics counters

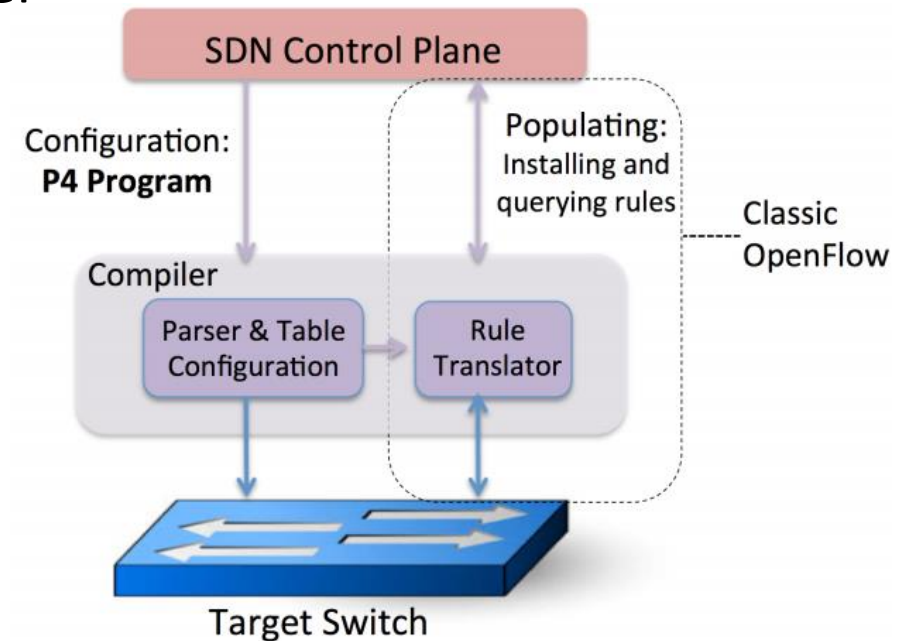
Outline

- Open vSwitch (OVS)
- Towards OpenFlow 2.0
- New Forwarding Plane Architectures
- **Programming the Forwarding Plane**

P4 Language



- P4: **P**rogramming **P**rotocol-Independent **P**acket **P**rocessors
 - An open source language allowing the specification of packet processing logic
 - Based on a Match+Action forwarding model
- Participated by
 - Stanford, Princeton
 - Google, Intel, Microsoft, Barefoot



Read:

P4: Programming Protocol-independent Packet Processors [\[SIGCOMM Review'14\]](#)

Programming RMT with P4

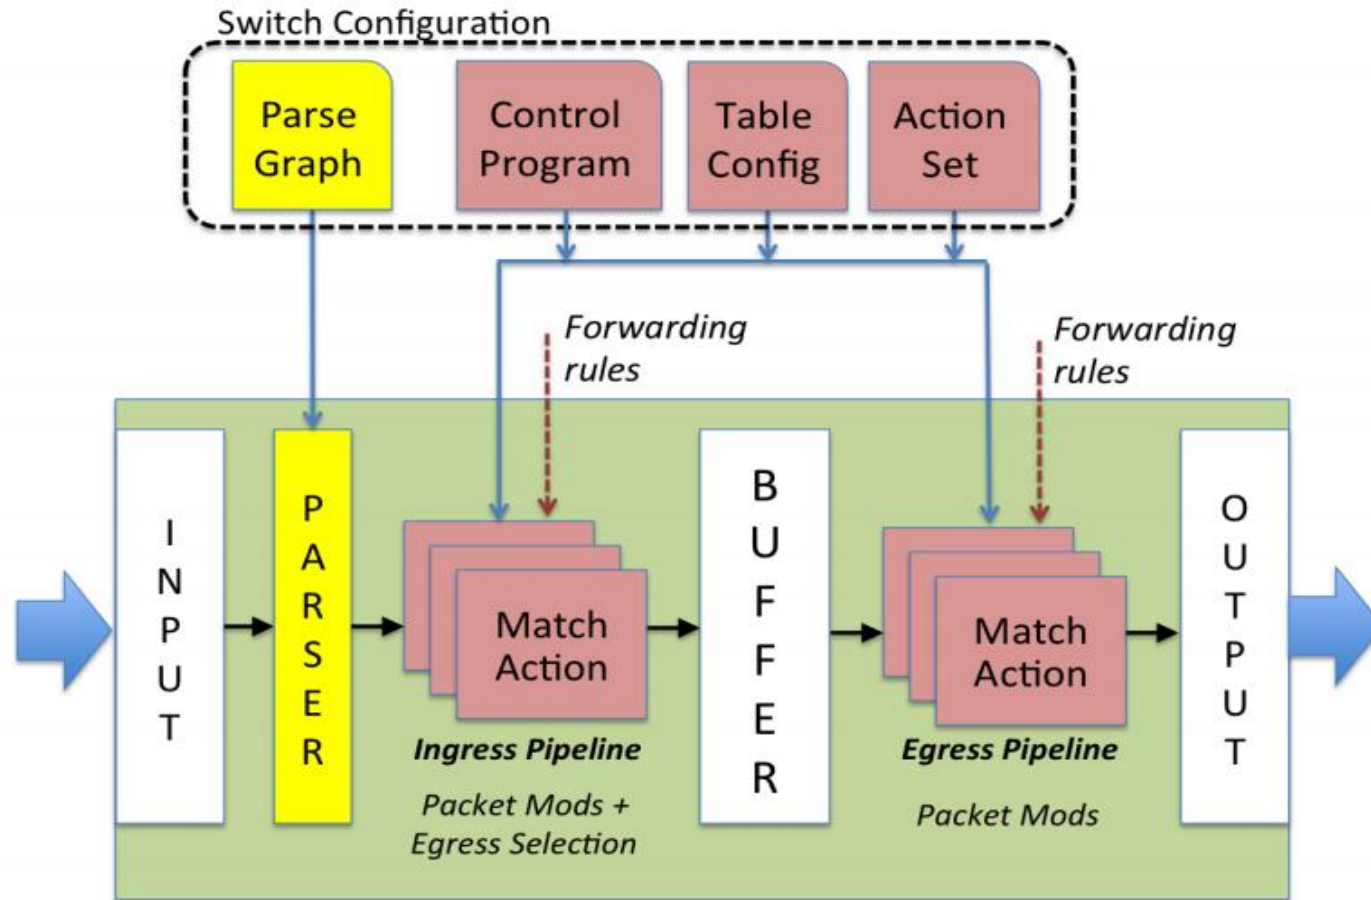
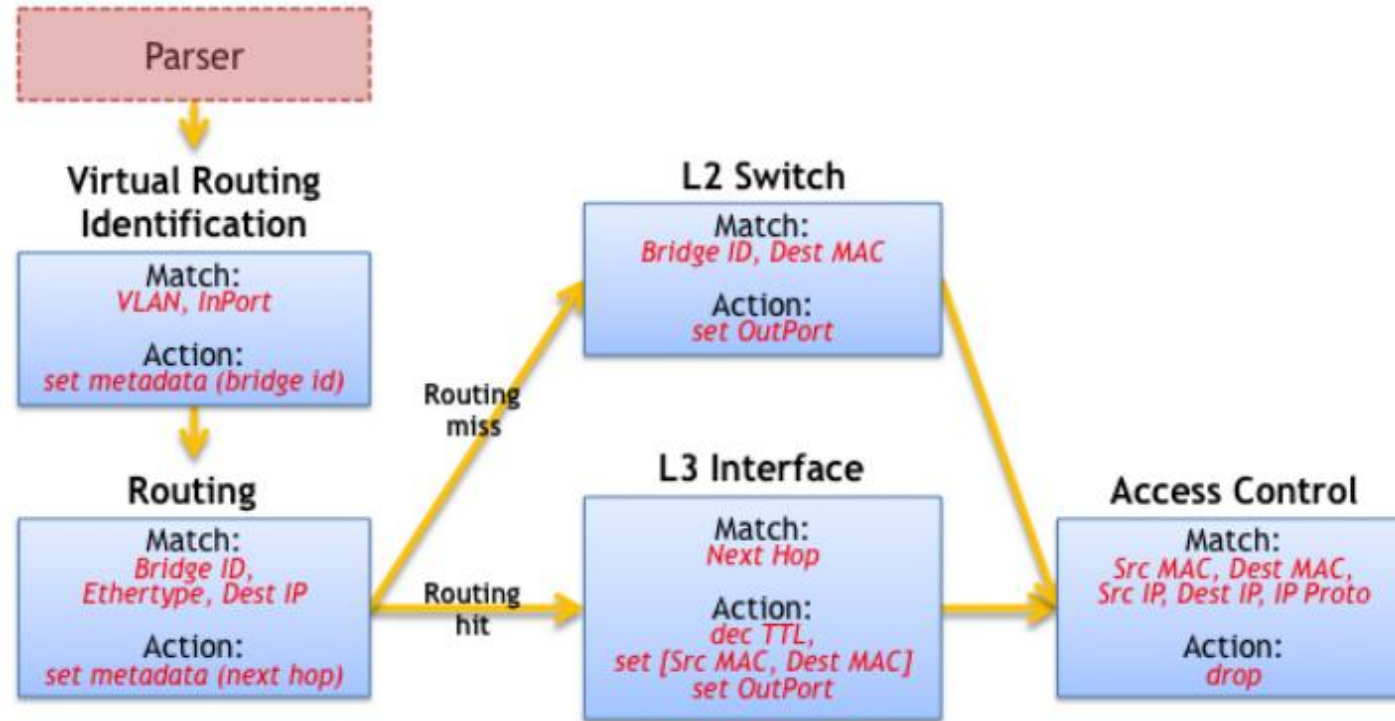


Table dependency graph (TDG) for an L2/L3 switch



P4 Language Elements

- **Headers**

- A header describes the sequence and structure of a series of fields

- **Parsers**

- A parser specifies how to identify headers within packets

- **Tables**

- A table defines the fields on which a packet may match and the actions may take

- **Actions**

- Complex actions can be constructed from simpler protocol-independent primitives

- **Control Programs**

- A control program gives the order of match+action tables applied to a packet

Headers and Fields

- **Fields** have width and other attributes
- **Headers** are collections of fields

```
header ethernet {  
    fields {  
        dst_addr : 48; // width in bits  
        src_addr : 48;  
        ethertype : 16;  
    }  
}
```

```
header vlan {  
    fields {  
        pcp : 3;  
        cfi : 1;  
        vid : 12;  
        ethertype : 16;  
    }  
}
```

Parser



- Extracts header fields from the packet

```
parser start {  
    ethernet;  
}
```

```
parser ethernet {  
    switch(ethertype) {  
        case 0x8100: vlan;  
        case 0x9100: vlan;  
        case 0x800: ipv4;  
        // Other cases  
    }  
}
```

```
parser vlan {  
    switch(ethertype) {  
        case 0xaaaa: mTag;  
        case 0x800: ipv4;  
        // Other cases  
    }  
}
```

```
parser mTag {  
    switch(ethertype) {  
        case 0x800: ipv4;  
        // Other cases  
    }  
}
```




Table Specification




```
table mTag_table {  
    reads {  
        ethernet.dst_addr : exact;  
        vlan.vid : exact;  
    }  
    actions {  
        // At runtime, entries are programmed with params  
        // for the mTag action.  See below.  
        add_mTag;  
    }  
    max_size : 20000;  
}
```

Action Specification



```
action add_mTag(up1, up2, down1, down2, egr_spec) {  
    add_header(mTag);  
    // Copy VLAN ethertype to mTag  
    copy_field(mTag.ethertype, vlan.ethertype);  
    // Set VLAN's ethertype to signal mTag  
    set_field(vlan.ethertype, 0xaaaa);  
    set_field(mTag.up1, up1);  
    set_field(mTag.up2, up2);  
    set_field(mTag.down1, down1);  
    set_field(mTag.down2, down2);  
  
    // Set the destination egress port as well  
    set_field(metadata.egress_spec, egr_spec);  
}
```

The Control Program



Flow chart for the mTag example

```
control main() {  
    // Verify mTag state and port are consistent  
    table(source_check);  
  
    // If no error from source_check, continue  
    if (!defined(metadata.ingress_error)) {  
        // Attempt to switch to end hosts  
        table(local_switching);  
  
        if (!defined(metadata.egress_spec)) {  
            // Not a known local host; try mtagging  
            table(mTag_table);  
        }  
  
        // Check for unknown egress state or  
        // bad retagging with mTag.  
        table(egress_check);  
    }  
}
```

Compiling a P4 Program

- The compiler translates the parser description into a parsing state machine

Current State	Lookup Value	Next State
vlan	0xaaaa	mTag
vlan	0x800	ipv4
vlan	*	stop
mTag	0x800	ipv4
mTag	*	stop

A partial state machine for the vlan and mTag section